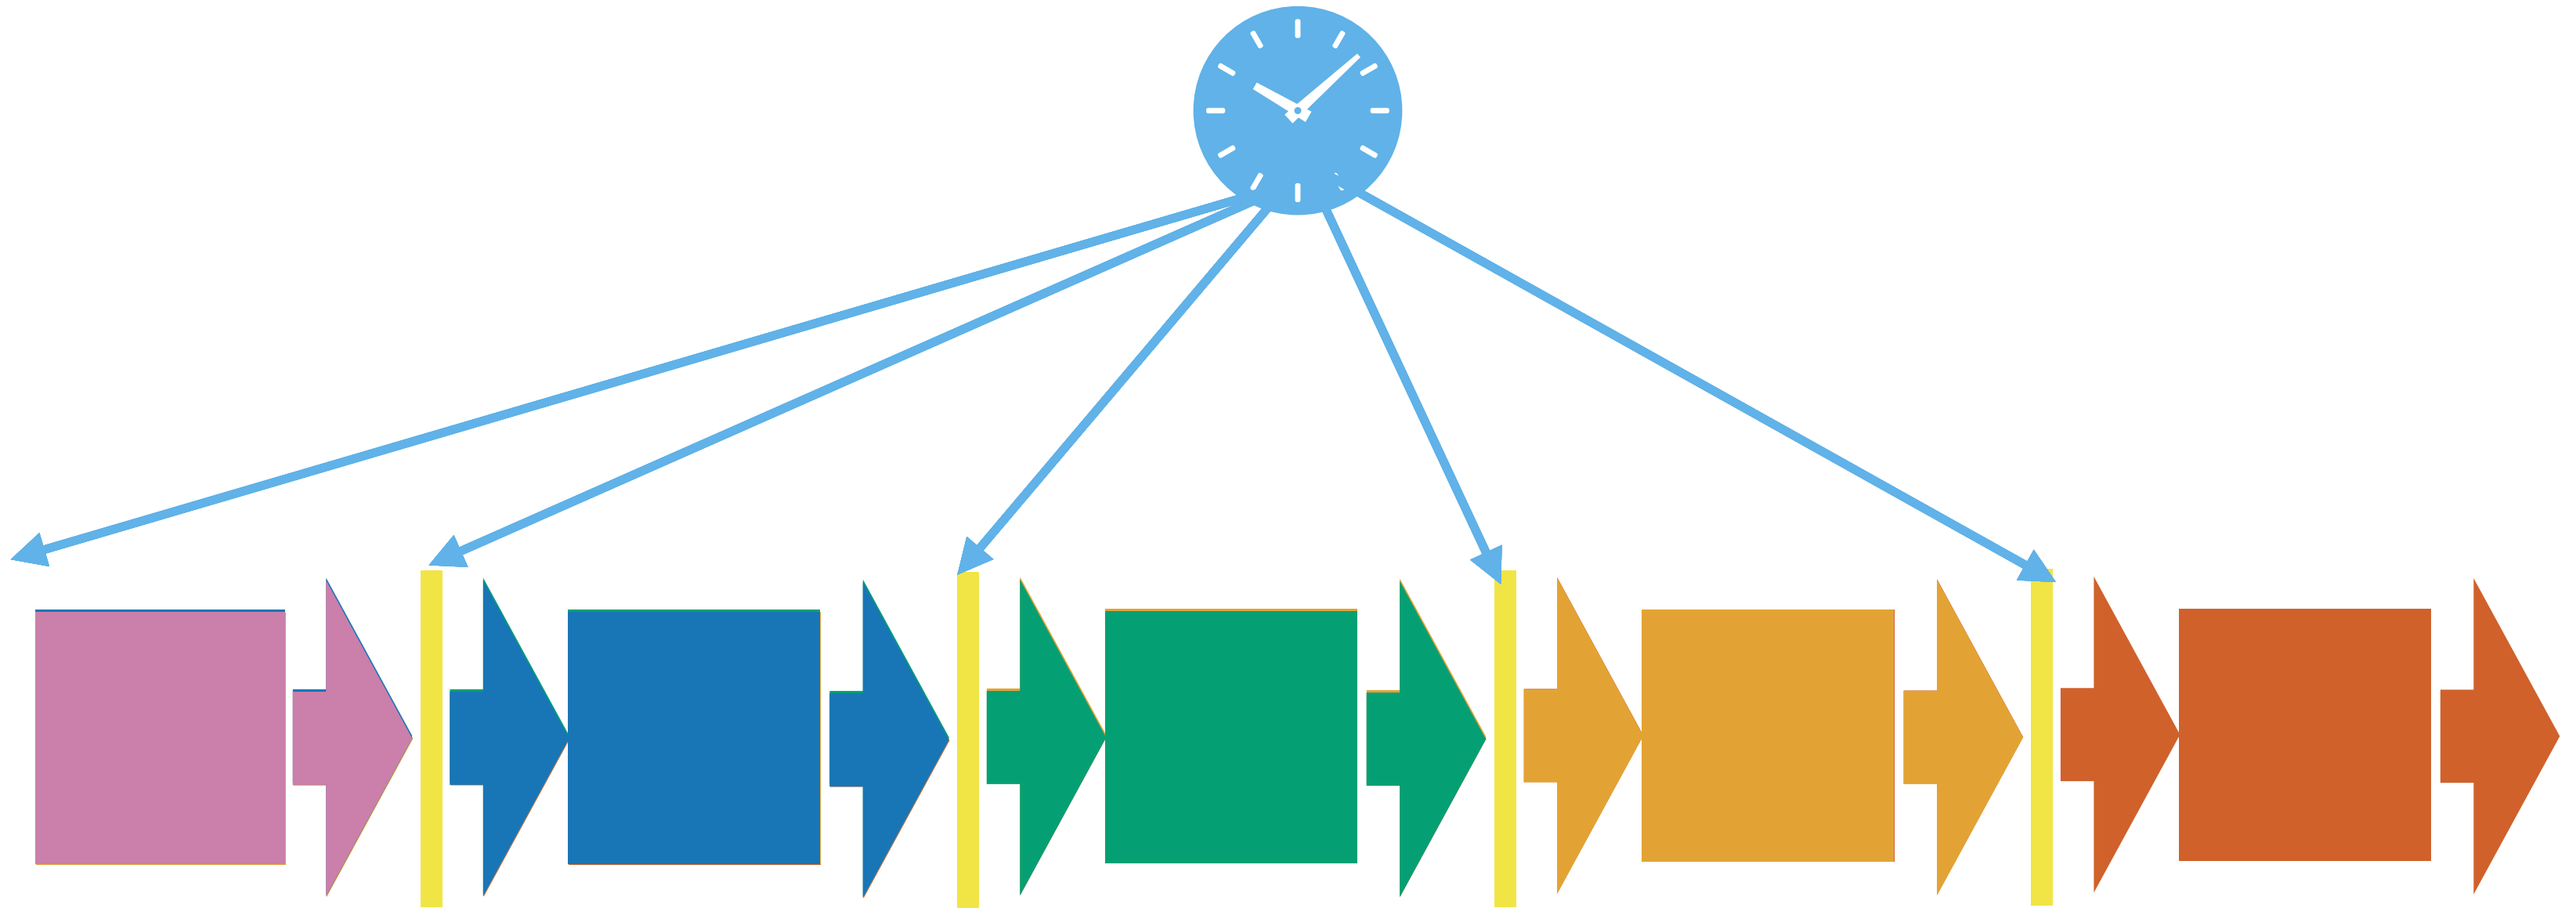


Data Hazards & Dynamic Instruction Scheduling

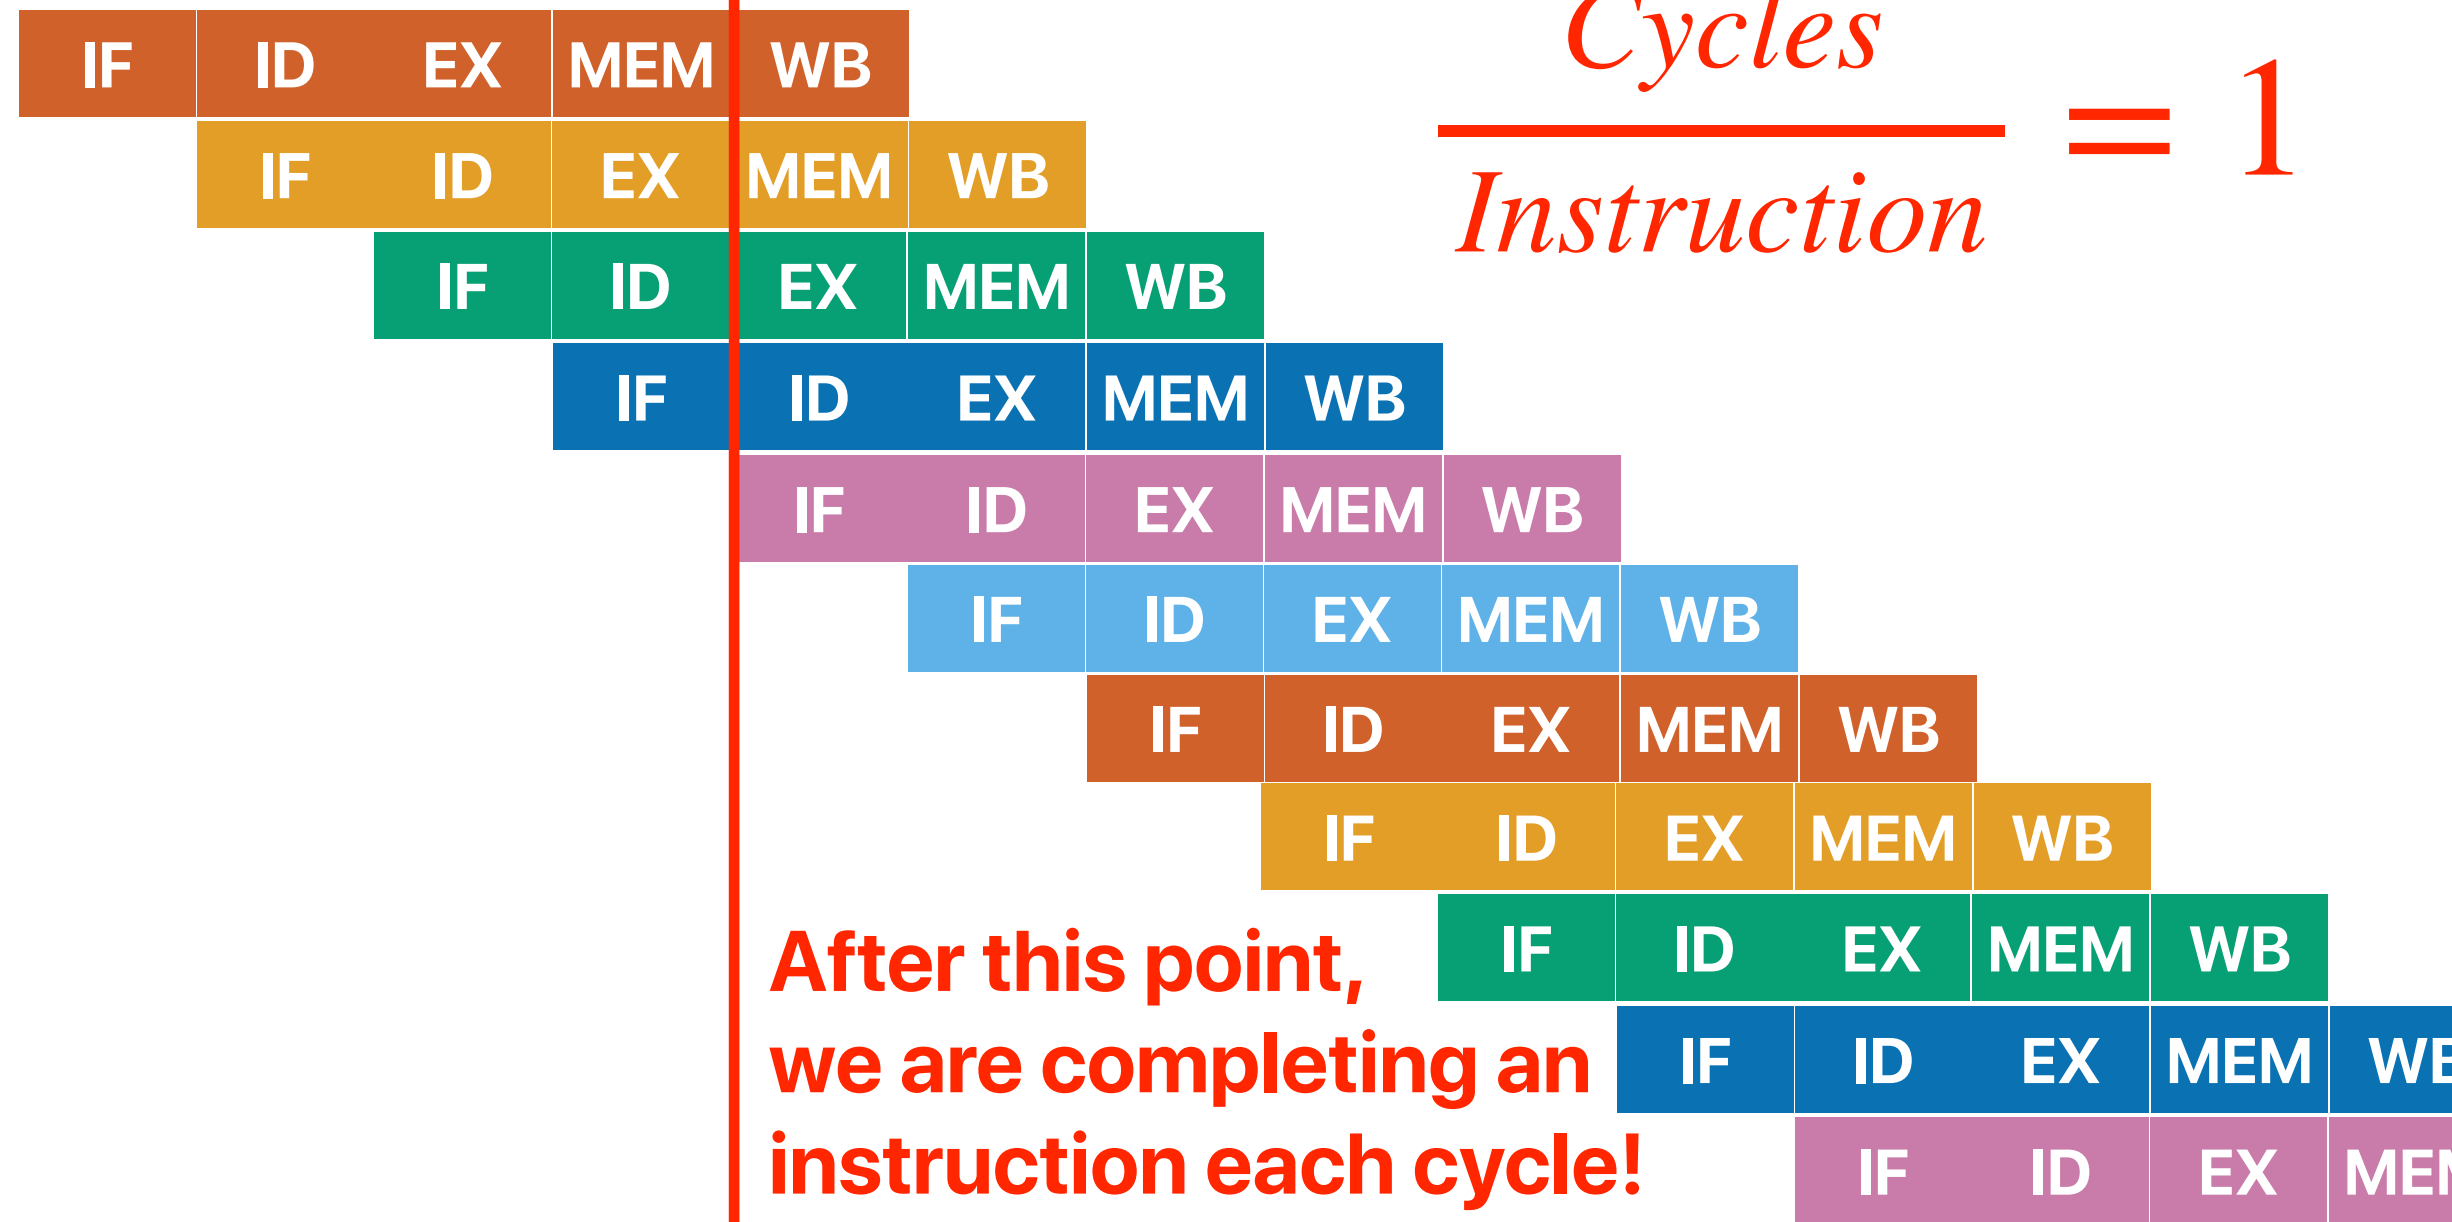
Hung-Wei Tseng

Recap: Pipelining



Recap: Pipelining

```
add x1, x2, x3
ld x4, 0(x5)
sub x6, x7, x8
sub x9, x10, x11
sd x1, 0(x12)
xor x13, x14, x15
and x16, x17, x18
add x19, x20, x21
sub x22, x23, x24
ld x25, 4(x26)
sd x27, 0(x28)
```



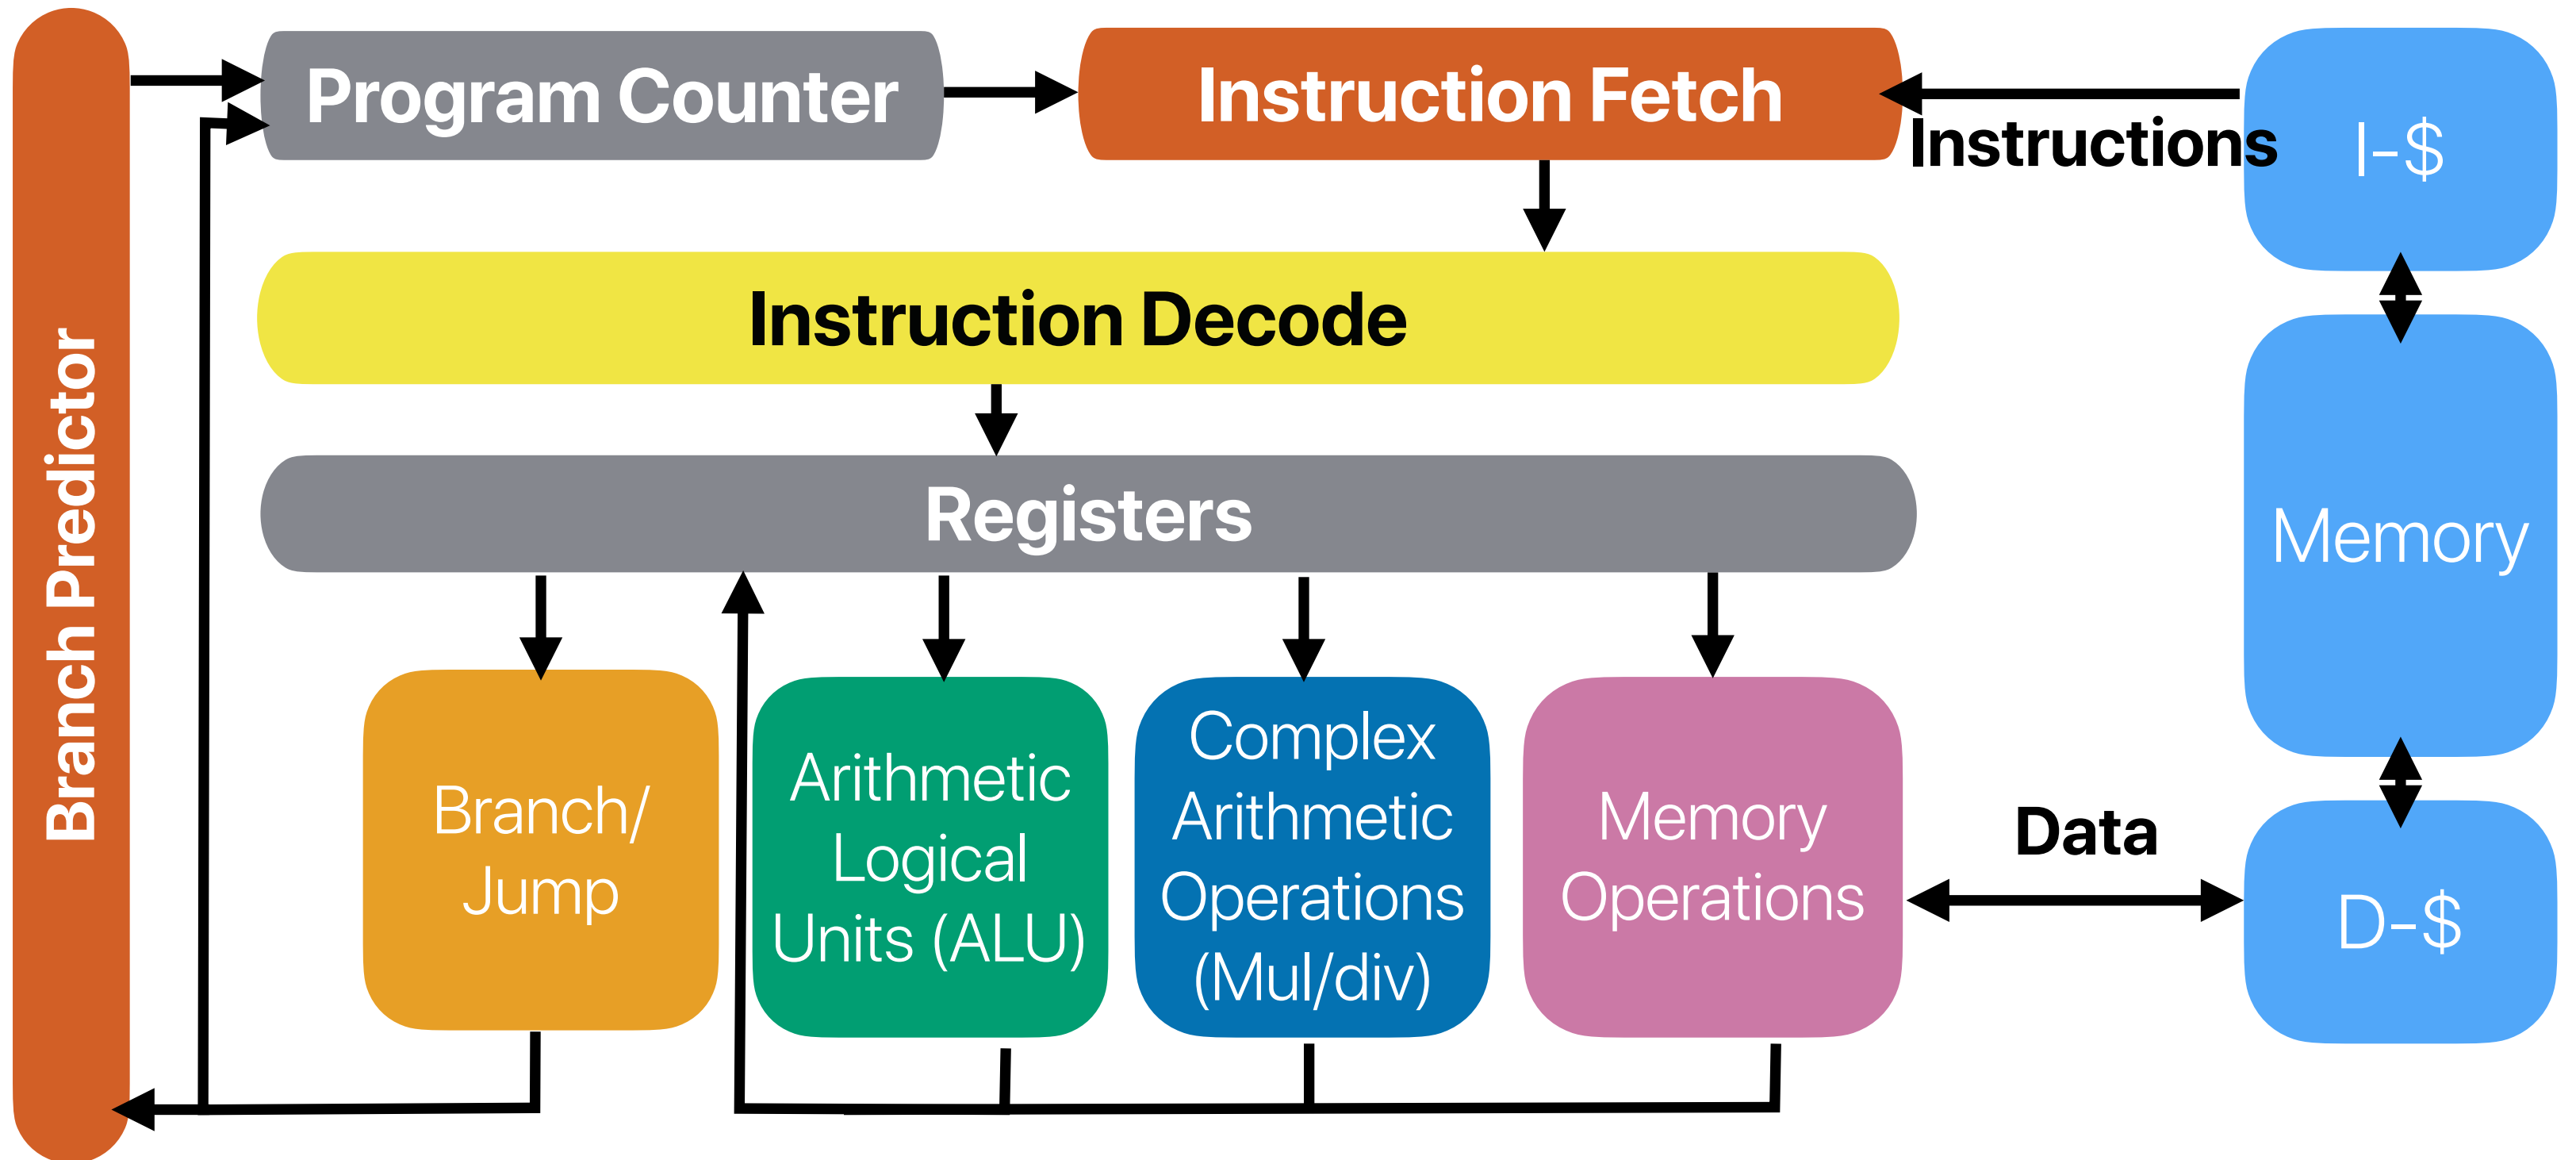
Recap: Three pipeline hazards

- Structural hazards — resource conflicts cannot support simultaneous execution of instructions in the pipeline
- Control hazards — the PC can be changed by an instruction in the pipeline
- Data hazards — an instruction depending on a the result that's not yet generated or propagated when the instruction needs that

Recap: addressing hazards

- Structural hazards
 - Stall
 - Modify hardware design
- Control hazards
 - Stall
 - Static prediction
 - Dynamic prediction — all “high-performance” processors nowadays have pretty decent branch predictors
 - Local bimodal
 - Global 2-level
 - Perceptron

The "current" pipeline



3:00

**What will you do if you're waiting for
someone else's response or
outcome?**

Ideas?

Outline

- Data hazards
 - Data forwarding
- SuperScalar
- Out-of-order, Dynamic instruction scheduling

Branch and programming

Demo revisited

```
if(option)
    std::sort(data, data + arraySize);

for (unsigned i = 0; i < 100000; ++i) {
    int threshold = std::rand();
    for (unsigned i = 0; i < arraySize; ++i) {
        if (data[i] >= threshold)
            sum ++;
    }
}
```

option = 1 is faster!!!

SELECT count(*) FROM TABLE WHERE val < A and val >= B;



Demo revisited

- Why the performance is better when option is not "0"
 - ① The amount of dynamic instructions needs to execute is a lot smaller
 - ② The amount of branch instructions to execute is smaller
 - ③ The amount of branch mis-predictions is smaller
 - ④ The amount of data accesses is smaller

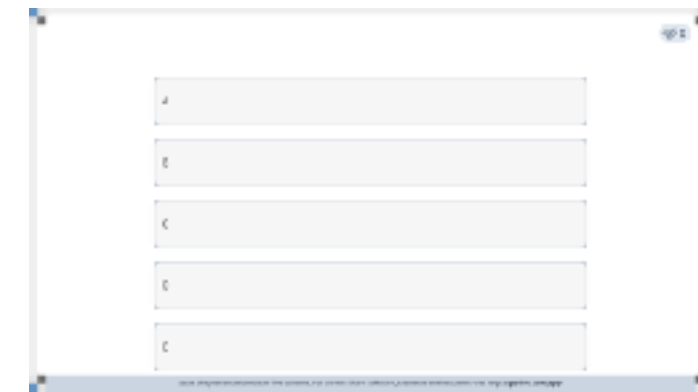
A. 0 `if(option)`
 `std::sort(data, data + arraySize);`

B. 1


C. 2 `for (unsigned i = 0; i < 100000; ++i) {`
 `int threshold = std::rand();`

D. 3 `for (unsigned i = 0; i < arraySize; ++i) {`
 `if (data[i] >= threshold)`
 `sum ++;`

E. 4 `}`
 `}`
 `}`




0%




A

0%




B

0%




C

0%



D

0%



E



Demo revisited

- Why the performance is better when option is not "0"
 - ① The amount of dynamic instructions needs to execute is a lot smaller
 - ② The amount of branch instructions to execute is smaller
 - ③ The amount of branch mis-predictions is smaller
 - ④ The amount of data accesses is smaller


A. 0 `if(option)`
 `std::sort(data, data + arraySize);`

B. 1


C. 2 `for (unsigned i = 0; i < 100000; ++i) {`
 `int threshold = std::rand();`

D. 3 `for (unsigned i = 0; i < arraySize; ++i) {`
 `if (data[i] >= threshold)`
 `sum ++;`

E. 4 `}`
 `}`
 `}`




0%




A

0%




B

0%




C

0%



D

0%



E

Demo revisited

- Why the performance is better when option is not "0"
 - ① The amount of dynamic instructions needs to execute is a lot smaller
 - ② The amount of branch instructions to execute is smaller
 - ✓③ The amount of branch mis-predictions is smaller
 - ④ The amount of data accesses is smaller

A. 0

B. 1

C. 2

D. 3

E. 4

```
if(option)
    std::sort(data, data + arraySize);

for (unsigned i = 0; i < 100000; ++i) {
    int threshold = std::rand();
    for (unsigned i = 0; i < arraySize;
        if (data[i] >= threshold)
            sum ++;
    }
}
```

branch X

	Without sorting	With sorting
The prediction accuracy of X before threshold	50%	100%
The prediction accuracy of X after threshold	50%	100%

Demo revisited: evaluating the cost of mis-predicted branches

- Compare the number of mis-predictions
- Calculate the difference of cycles
- We can get the “average CPI” of a mis-prediction!

34 cycles!!!

Data hazards

Data hazards

- An instruction currently in the pipeline cannot receive the “logically” correct value for execution
- Data dependencies
 - The output of an instruction is the input of a later instruction
 - May result in data hazard if the later instruction that consumes the result is still in the pipeline



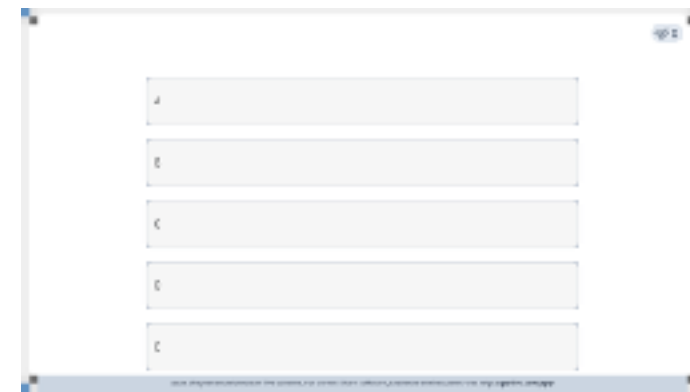
How many data dependencies do we have?

- How many pairs of data dependencies are there in the following x86 instructions?


```
movl    (%rdi), %eax
movl    (%rsi), %edx
movl    %edx, (%rdi)
movl    %eax, (%rsi)
```

```
int temp = *a;
*a = *b;
*b = temp;
```

- A. 1
- B. 2
- C. 3
- D. 4
- E. 5




0%




A

0%




B

0%




C

0%



D

0%



E



How many data dependencies do we have?


- How many pairs of data dependencies are there in the following x86 instructions?

```
movl    (%rdi), %eax  
movl    (%rsi), %edx  
movl    %edx, (%rdi)  
movl    %eax, (%rsi)
```

- A. 1
- B. 2
- C. 3
- D. 4
- E. 5


A screenshot of a poll interface. It shows five empty input boxes, each preceded by a letter (A, B, C, D, E) in a small font. The boxes are arranged vertically. The interface is part of a larger application, with some UI elements visible at the top and bottom of the poll area.

0%




A

0%




B

0%



C

0%



D

0%

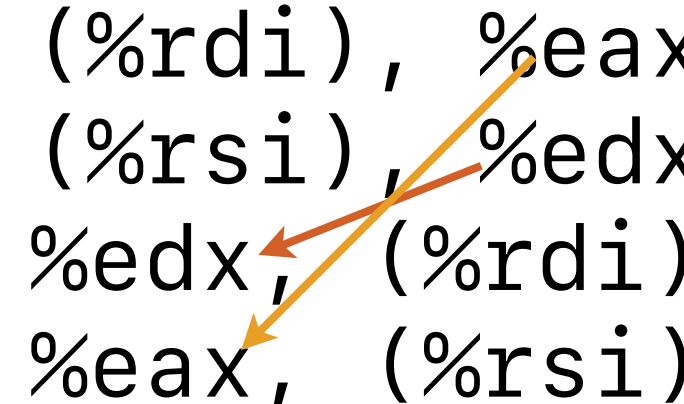


E

How many dependencies do we have?

- How many pairs of data dependences are there in the following x86 instructions?

```
movl    (%rdi), %eax
movl    (%rsi), %edx
movl    %edx, (%rdi)
movl    %eax, (%rsi)
```



```
int temp = *a;
*a = *b;
*b = temp;
```

A. 1

B. 2

C. 3

D. 4

E. 5



How many data dependencies do we have?

- How many pairs of data dependencies are there in the following x86 instructions?


```
movl    (%rdi), %eax
xorl    (%rsi), %eax
movl    %eax, (%rdi)
xorl    (%rsi), %eax
movl    %eax, (%rsi)
xorl    %eax, (%rdi)
```

```
*a ^= *b;
*b ^= *a;
*a ^= *b;
```

- A. 1
- B. 2
- C. 3
- D. 4
- E. 5




0%




A

0%




B

0%




C

0%



D

0%



E




How many data dependencies do we have?

- How many pairs of data dependencies are there in the following x86 instructions?


```
movl    (%rdi), %eax
xorl    (%rsi), %eax
movl    %eax, (%rdi)
xorl    (%rsi), %eax
movl    %eax, (%rsi)
xorl    %eax, (%rdi)
```

```
*a ^= *b;
*b ^= *a;
*a ^= *b;
```

- A. 1
- B. 2
- C. 3
- D. 4
- E. 5




0%




A

0%




B

0%




C

0%



D

0%



E

How many dependencies do we have?

- How many pairs of data dependencies are there in the following x86 instructions?

```
movl    (%rdi), %eax
xorl    (%rsi), %eax
movl    %eax, (%rdi)
xorl    (%rsi), %eax
movl    %eax, (%rsi)
xorl    %eax, (%rdi)
```

```
*a ^= *b;
*b ^= *a;
*a ^= *b;
```

- A. 1
- B. 2
- C. 3
- D. 4
- E. 5




Data hazards?


- How many pairs of data dependences in the following x86 instructions will result in data hazards if a memory operation (assume 100% cache hit rate) takes 4 cycles?

```
movl    (%rdi), %eax  
movl    (%rsi), %edx  
movl    %edx, (%rdi)  
movl    %eax, (%rsi)
```

- A. 0
- B. 1
- C. 2
- D. 3
- E. 4




0%




A

0%




B

0%




C

0%



D

0%



E




Data hazards?


- How many pairs of data dependences in the following x86 instructions will result in data hazards if a memory operation (assume 100% cache hit rate) takes 4 cycles?

```
movl    (%rdi), %eax  
movl    (%rsi), %edx  
movl    %edx, (%rdi)  
movl    %eax, (%rsi)
```

- A. 0
- B. 1
- C. 2
- D. 3
- E. 4




0%




A

0%




B

0%




C

0%



D


0%



E

Data hazards?

- How many pairs of data dependences in the following x86 instructions will result in data hazards if a memory operation (assume 100% cache hit rate) takes 4 cycles?



A. 0

B. 1

C. 2

D. 3

E. 4

Data hazards

① movl (%rdi), %eax

② movl (%rsi), %edx

③ movl %edx, (%rdi)

④ movl %eax, (%rsi)

%edx does not have our desired value

	IF	ID	ALU/BR/M1	M2	M3	M4	WB
1	(1)						
2	(2)	(1)					
3	(3)	(2)	(1)				
4	(4)	(3)	(2)	(1)			
5		(4)	(3)	(2)	(1)		
6			(4)	(3)	(2)	(1)	
7				(4)	(3)	(2)	(1)
8					(4)	(3)	(2)
9						(4)	(3)
10							(4)
11							
12							
13							
14							


%eax does not have our desired value

Solution 1: Let's try "stall" again

- Whenever the input is not ready when the consumer is decoding, just stall — the consumer stays at ID.

we have the value for %edx already! Why another cycle?

```
movl    (%rdi), %eax
movl    (%rsi), %edx
movl    %edx, (%rdi)
movl    %eax, (%rsi)
```



3 additional cycles

Solution 1: Let's try "stall" again

- Whenever the input is not ready when the consumer is decoding, just stall — the consumer stays at ID.

① `movl (%rdi), %eax`
② `movl (%rsi), %edx`
③ `movl %edx, (%rdi)`
④ `movl %eax, (%rsi)`

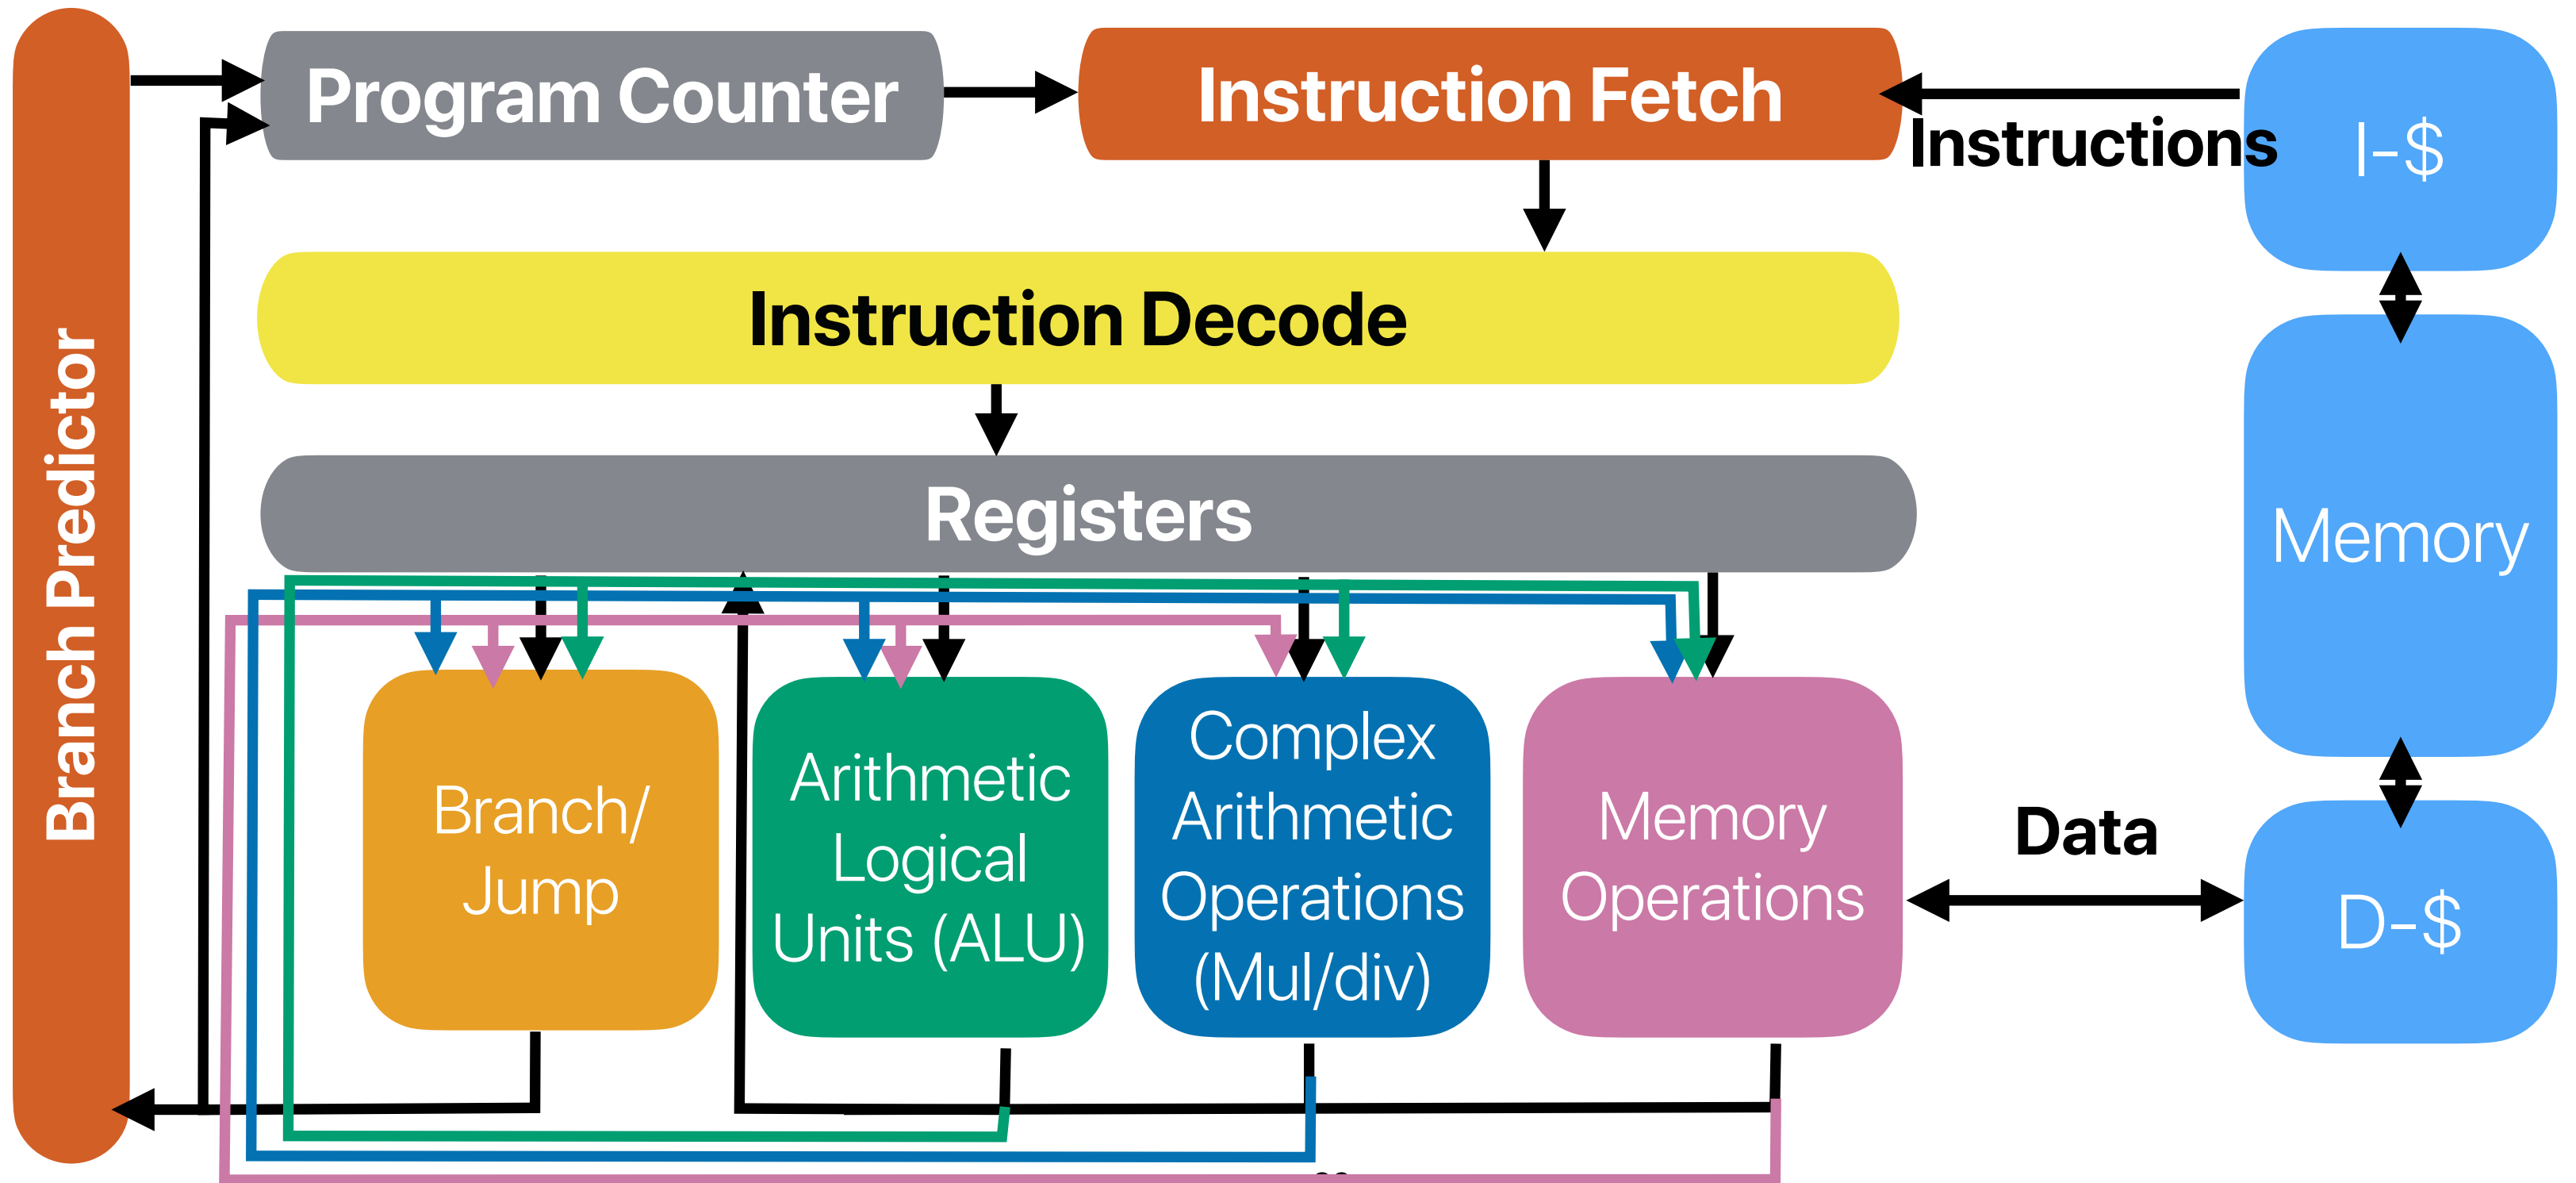
we have the value for %edx already!
Why another cycle?

	IF	ID	ALU/BR/M1	M2	M3	M4	WB
1	(1)						
2	(2)	(1)					
3	(3)	(2)	(1)				
4	(4)	(3)	(2)	(1)			
5	(4)	(3)		(2)	(1)		
6	(4)	(3)			(2)	(1)	
7	(4)	(3)				(2)	(1)
8	(4)	(3)					(2)
9		(4)	(3)				
10			(4)	(3)			
11				(4)	(3)		
12					(4)	(3)	
13						(4)	(3)
14							(4)

Solution 2: Data forwarding

- Add logics/wires to forward the desired values to the demanding instructions

Data "forwarding"






How many data dependencies are still problematic?


- How many pairs of data dependences in the following x86 instructions are still problematic with data forwarding if a memory operation (assume 100% cache hit rate) takes 4 cycles?

```
movl    (%rdi), %eax
movl    (%rsi), %edx
movl    %edx, (%rdi)
movl    %eax, (%rsi)
```

- A. 0
- B. 1
- C. 2
- D. 3
- E. 4




0%




A

0%




B

0%




C

0%



D

0%



E



How many data dependencies are still problematic?


- How many pairs of data dependences in the following x86 instructions are still problematic with data forwarding if a memory operation (assume 100% cache hit rate) takes 4 cycles?

```
movl    (%rdi), %eax
movl    (%rsi), %edx
movl    %edx, (%rdi)
movl    %eax, (%rsi)
```

- A. 0
- B. 1
- C. 2
- D. 3
- E. 4


A screenshot of a poll interface. It shows five empty input boxes, each preceded by a letter (A, B, C, D, E) in a small font. The boxes are arranged vertically. At the bottom, there is a small text line that reads "You are not logged in. Please log in to participate in the poll." and a "Log in" button.

0%




A

0%




B

0%




C

0%



D

0%

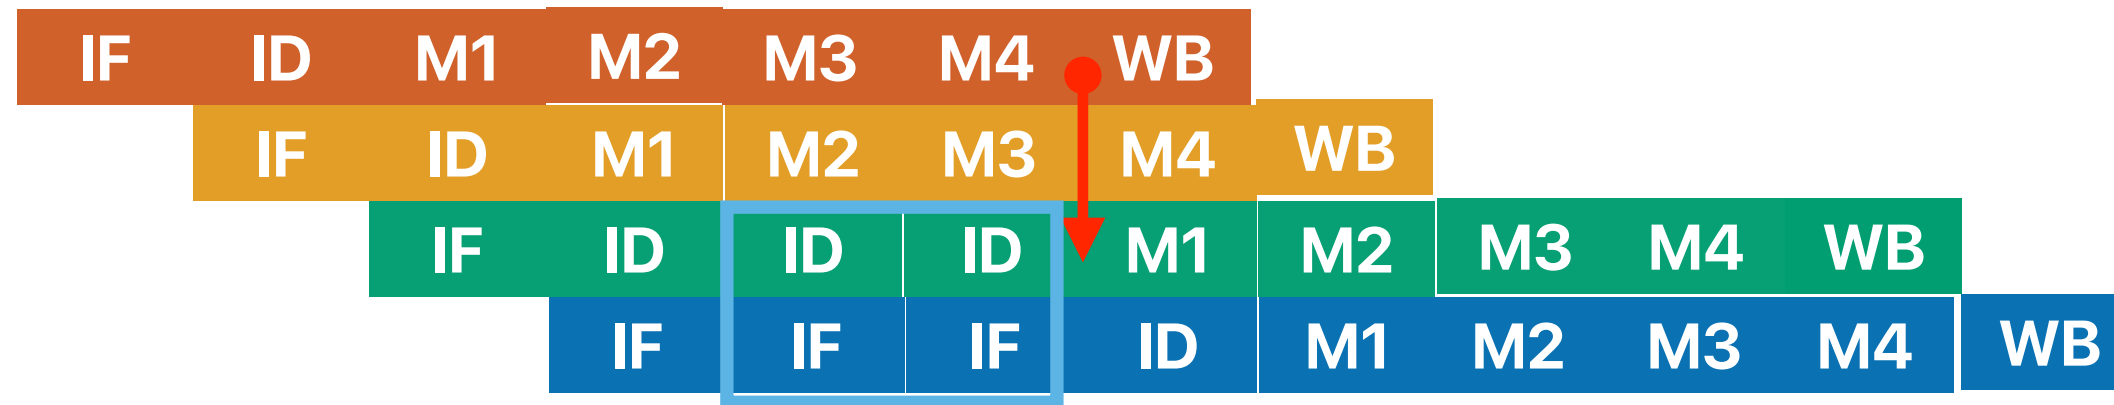


E

How many data dependencies are still problematic?

- How many pairs of data dependencies in the following x86 instructions are still problematic with data forwarding if a memory operation (assume 100% cache hit rate) takes 4 cycles?

```
movl    (%rdi), %eax
movl    (%rsi), %edx
movl    %edx, (%rdi)
movl    %eax, (%rsi)
```



2 additional cycles

A. 0

B. 1

C. 2

D. 3

E. 4

```
int temp = *a;
*a = *b;
*b = temp;
```

Solution 2: Data forwarding

- Whenever the input is not ready when the consumer is decoding, just stall — the consumer stays at ID.

① `movl (%rdi), %eax`
② `movl (%rsi), %edx`
③ `movl %edx, (%rdi)`
④ `movl %eax, (%rsi)`

	IF	ID	ALU/BR/M1	M2	M3	M4	WB
1	(1)						
2	(2)	(1)					
3	(3)	(2)	(1)				
4	(4)	(3)	(2)	(1)			
5	(4)	(3)		(2)	(1)		
6	(4)	(3)			(2)	(1)	
7	(4)	(3)	data forwarding			(2)	(1)
8		(4)					(2)
9			(4)	(3)			
10				(4)	(3)		
11					(4)	(3)	
12						(4)	(3)
13							(4)
14							



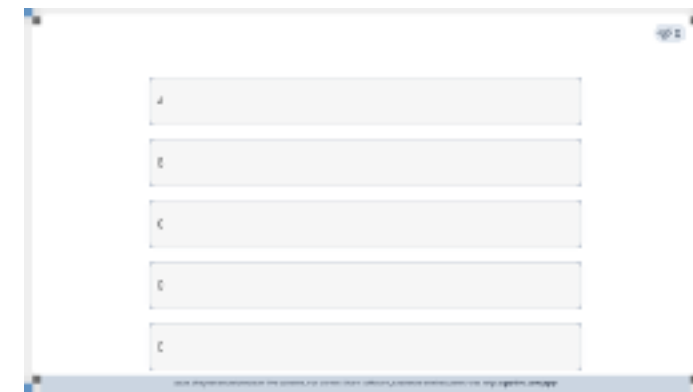
How many of data hazards w/ Data Forwarding?

- How many pairs of data dependences in the following x86 instructions are still problematic with data forwarding if both a memory operation and an xorl take 4 cycles?


```
① movl    (%rdi), %eax
② xorl    (%rsi), %eax
③ movl    %eax, (%rdi)
④ xorl    (%rsi), %eax
⑤ movl    %eax, (%rsi)
⑥ xorl    %eax, (%rdi)
```

```
*a ^= *b;
*b ^= *a;
*a ^= *b;
```

- A. 0
- B. 1
- C. 2
- D. 3
- E. 4




0%




A

0%




B

0%




C

0%



D

0%



E




How many of data hazards w/ Data Forwarding?

- How many pairs of data dependences in the following x86 instructions are still problematic with data forwarding if both a memory operation and an xorl take 4 cycles?


```
① movl    (%rdi), %eax
② xorl    (%rsi), %eax
③ movl    %eax, (%rdi)
④ xorl    (%rsi), %eax
⑤ movl    %eax, (%rsi)
⑥ xorl    %eax, (%rdi)
```

```
*a ^= *b;
*b ^= *a;
*a ^= *b;
```

- A. 0
- B. 1
- C. 2
- D. 3
- E. 4




0%




A

0%




B

0%




C

0%



D

0%



E

How many of data hazards w/ Data Forwarding?

- How many pairs of data dependences in the following x86 instructions are still problematic with data forwarding if both a memory operation and an xorl take 4 cycles?

① movl (%rdi), %eax
② xorl (%rsi), %eax
③ movl %eax, (%rdi)
④ xorl (%rsi), %eax
⑤ movl %eax, (%rsi)
⑥ xorl %eax, (%rdi)

A. 0

B. 1

C. 2

D. 3

E. 4

	IF	ID	ALU/BR/M1	M2	M3	M4/XORL	WB
1	(1)						
2	(2)	(1)					
3	(3)	(2)	(1)				
4	(3)	(2)		(1)			
5	(3)	(2)			(1)		
6	(3)	(2)				(1)	
7	(4)	(3)	(2)				(1)
8	(4)	(3)		(2)			
9	(4)	(3)			(2)		
10	(4)	(3)				(2)	
11	(5)	(4)	(3)				(2)
12	(6)	(5)	(4)	(3)			
13	(6)	(5)		(4)	(3)		
14	(6)	(5)			(4)	(3)	
15	(6)	(5)				(4)	(3)
16		(6)	(5)				(4)
17			(6)	(5)			
18				(6)	(5)		
19					(6)	(5)	
20						(6)	(5)
21							(6)
22							

Another code example

```
for(i = 0; i < count; i++) {  
    s += a[i];  
}
```

```
.L3:  
①      movl    (%rdi), %ecx  
②      addl    %ecx, %eax  
③      addq    $4, %rdi  
④      cmpq    %rdx, %rdi  
⑤      jne     .L3  
⑥      ret
```

	IF	ID	ALU/BR/M1	M2	M3	M4/XORL	WB
1	(1)						
2	(2)	(1)					
3	(3)	(2)	(1)				
4	(3)	(2)		(1)			
5	(3)	(2)			(1)		
6	(3)	(2)				(1)	
7	(4)	(3)	(2)				(1)
8	(5)	(4)	(3)	(2)			
9		(5)	(4)	(3)	(2)		
10			(5)	(4)	(3)	(2)	
11				(5)	(4)	(3)	(2)
12					(5)	(4)	(3)
13						(5)	(4)



The effect of code optimization

- By reordering which pair of the following instruction stream can we eliminate all stalls without affecting the correctness of the code?

① `movl (%rdi), %ecx`

② `addl %ecx, %eax`

③ `addq $4, %rdi`

④ `cmpq %rdx, %rdi`

⑤ `jne .L3`

⑥ `ret`

A. (1) & (2)

B. (2) & (3)


C. (3) & (4)

D. (4) & (5)

E. None of the pairs can be reordered




0%




A

0%




B

0%




C

0%



D

0%



E




The effect of code optimization


- By reordering which pair of the following instruction stream can we eliminate all stalls without affecting the correctness of the code?

① `movl (%rdi), %ecx`
② `addl %ecx, %eax`
③ `addq $4, %rdi`
④ `cmpq %rdx, %rdi`
⑤ `jne .L3`
⑥ `ret`

- A. (1) & (2)
B. (2) & (3)
C. (3) & (4)
D. (4) & (5)
E. None of the pairs can be reordered




0%




A

0%




B

0%




C

0%



D

0%



E

The effect of code optimization

- By reordering which pair of the following instruction stream can we eliminate all stalls without affecting the correctness of the code?

① `movl (%rdi), %ecx`
② `addl %ecx, %eax`
③ `addq $4, %rdi`
④ `cmpq %rdx, %rdi`
⑤ `jne .L3`
⑥ `ret`

A. (1) & (2)

B. (2) & (3)

C. (3) & (4)

D. (4) & (5)

E. None of the pairs can be reordered

Compiler optimization

```
for(i = 0; i < count; i++) {  
    s += a[i];  
}
```

```
.L3:  
①    movl    (%rdi), %ecx  
②    addq    $4, %rdi  
③    addl    %ecx, %eax  
④    cmpq    %rdx, %rdi  
⑤    jne     .L3  
⑥    ret
```

	IF	ID	ALU/BR/M1	M2	M3	M4/XORL	WB
1	(1)						
2	(2)	(1)					
3	(3)	(2)	(1)				
4	(3)	(2)	(2)	(1)			
5	(3)	(2)		(2)	(1)		
6	(4)	(2)			(2)	(1)	
7	(5)	(4)	(3)			(2)	(1)
8		(5)	(4)	(3)			(2)
9			(5)	(4)	(3)		
10				(5)	(4)	(3)	
11					(5)	(4)	(3)
12						(5)	(4)
13							(5)

Compiler optimization

```
for(i = 0; i < count; i++) {  
    s += a[i];  
}
```


```
.L3:  
①    movl    (%rdi), %ecx  
②    addq    $4, %rdi  
③    addl    %ecx, %eax  
④    cmpq    %rdx, %rdi  
⑤    jne     .L3  
⑥    ret
```

	IF	ID	ALU/BR/M1	M2	M3	M4/XORL	WB
1	(1)						
2	(2)	(1)					
3	(3)	(2)	(1)				
4	(3)	(2)	(2)	(1)			
5	(3)	(2)		(2)	(1)		
6	(4)	(2)			(2)	(1)	
7	(5)	(4)	(3)			(2)	(1)
8		(5)	(4)	(3)			(2)
9			(5)	(4)	(3)		
10				(5)	(4)	(3)	
11					(5)	(4)	(3)
12						(5)	(4)
13							(5)


addq is not depending on movl and
ALU is free! can we execute them
together?

If $CPI == 1$ the limitation?

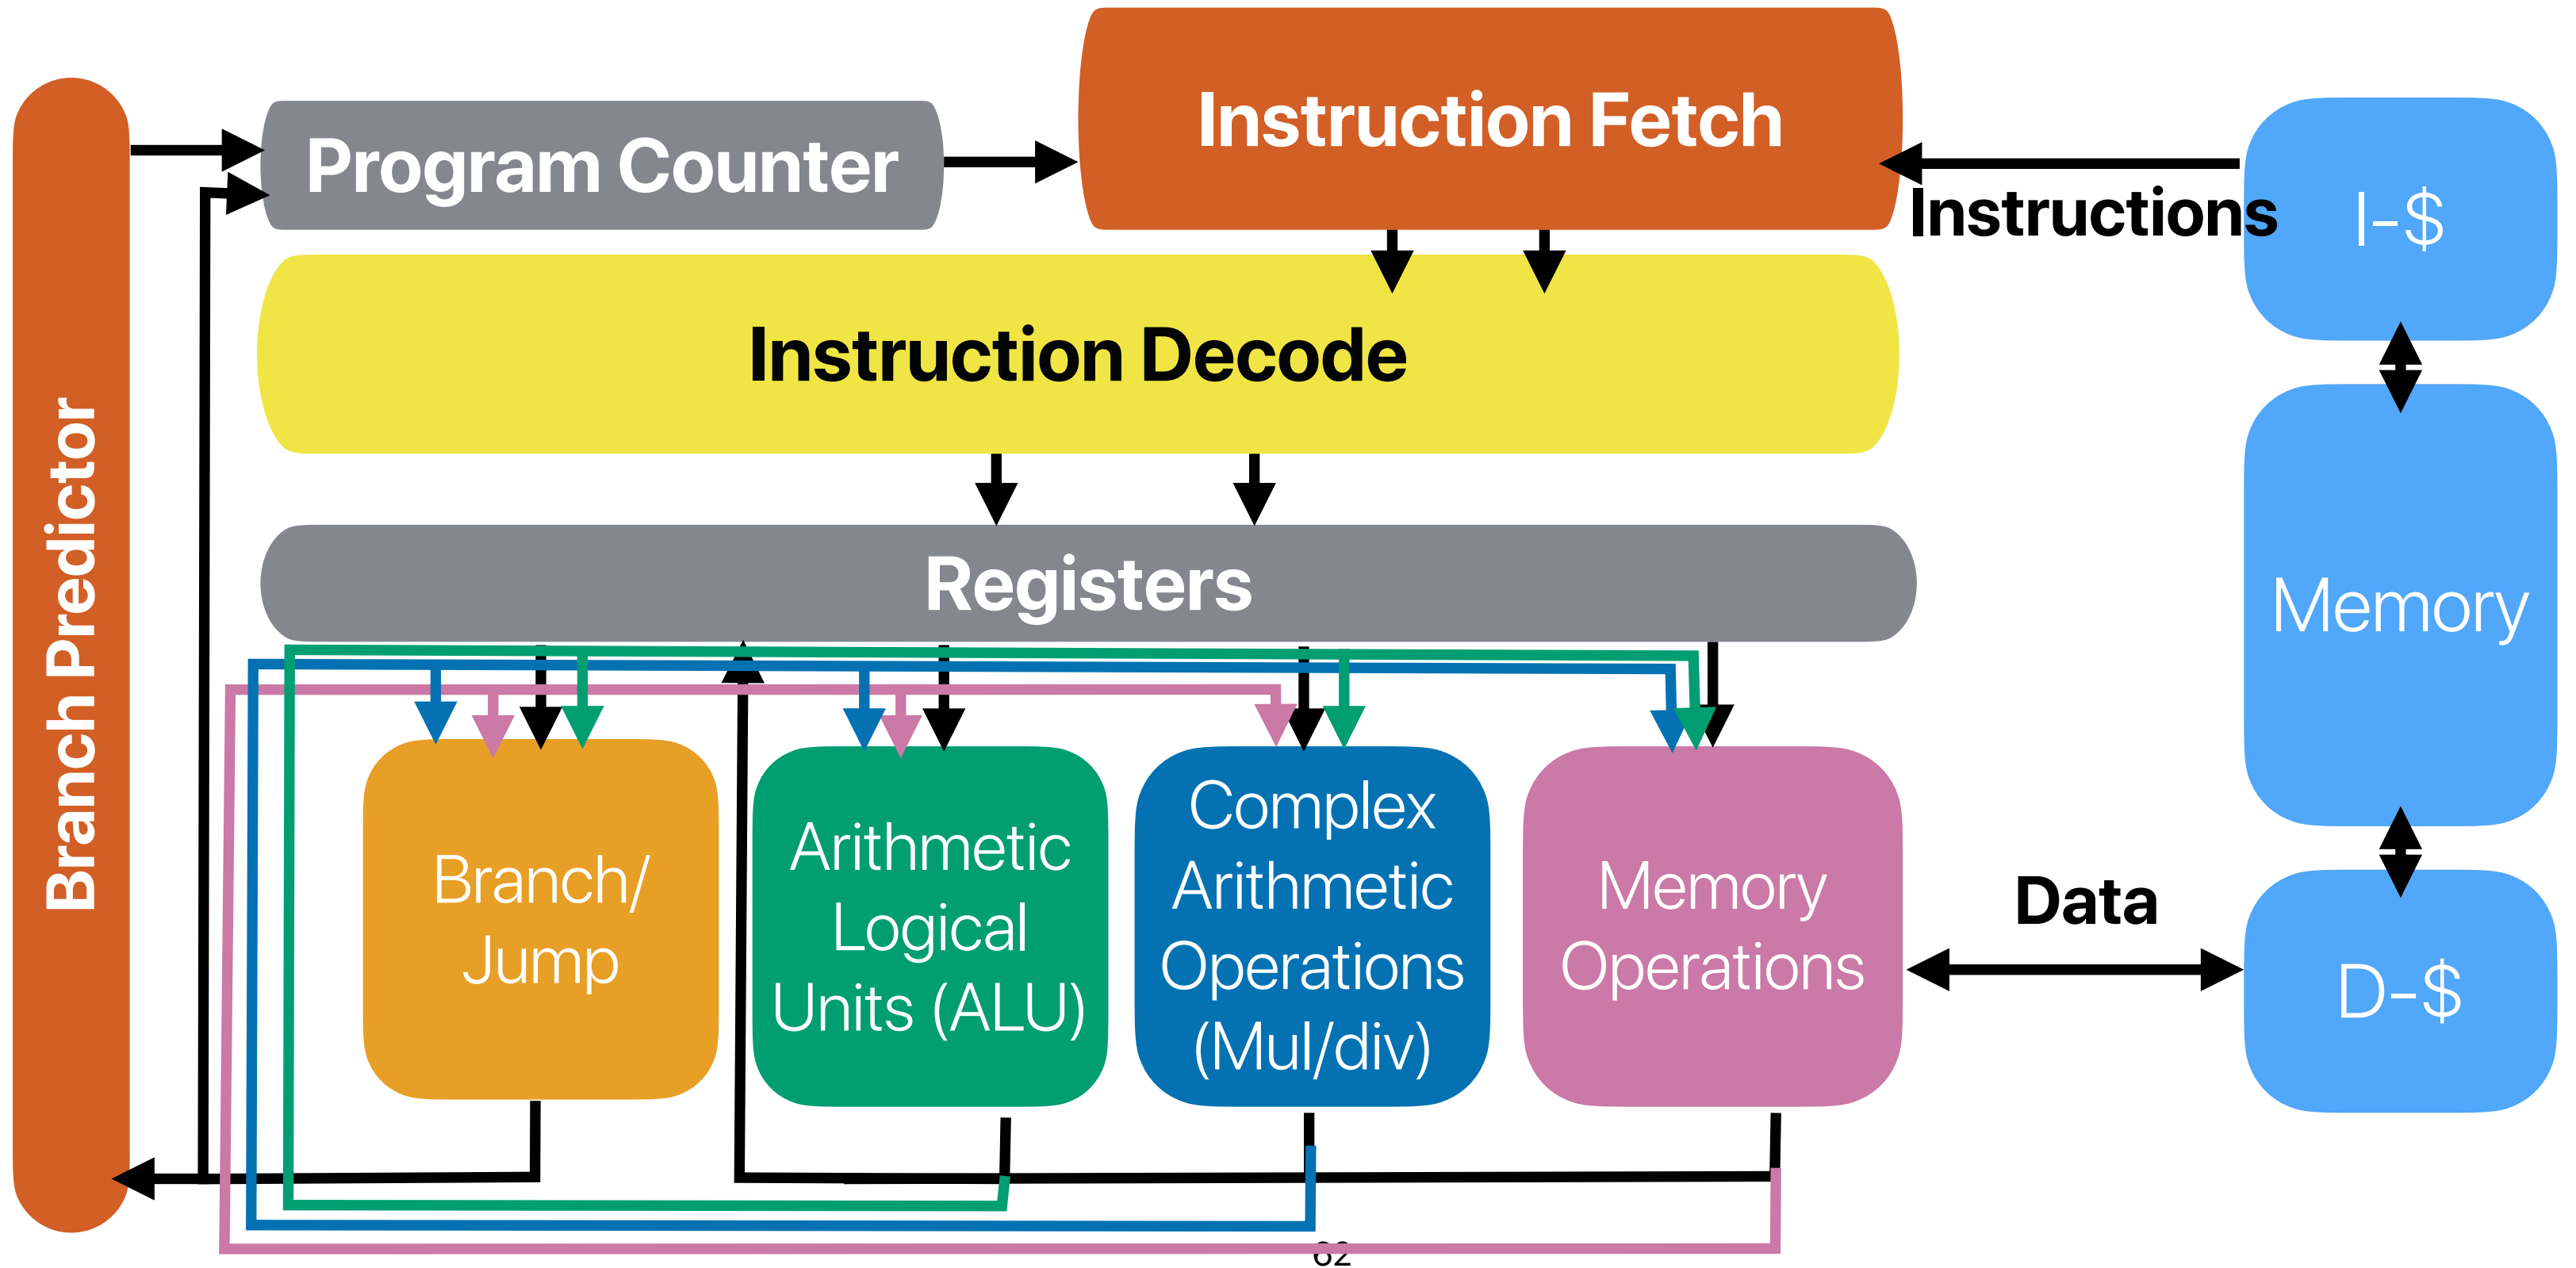
Data "forwarding"



Data "forwarding"



Super Scalar



Super Scalar

Superscalar

- Since we have many functional units now, we should fetch/decode more instructions each cycle so that we can have more instructions to issue!
- Super-scalar: fetch/decode/issue more than one instruction each cycle
 - **Fetch width:** how many instructions can the processor fetch/decode each cycle
 - **Issue width:** how many instructions can the processor issue each cycle
- The theoretical CPI should now be

1

$\min(\text{issue width}, \text{fetch width}, \text{decode width})$

Superscalar: fetch/issue width == 2, theoretical CPI = 0.5

```
for(i = 0; i < count; i++) {  
    s += a[i];  
}
```

.L3:

①

movl

(%rdi), %ecx

②

addq

\$4, %rdi

③

addl

%ecx, %eax

④

cmpq

%rdx, %rdi

⑤

jne

.L3

⑥

ret

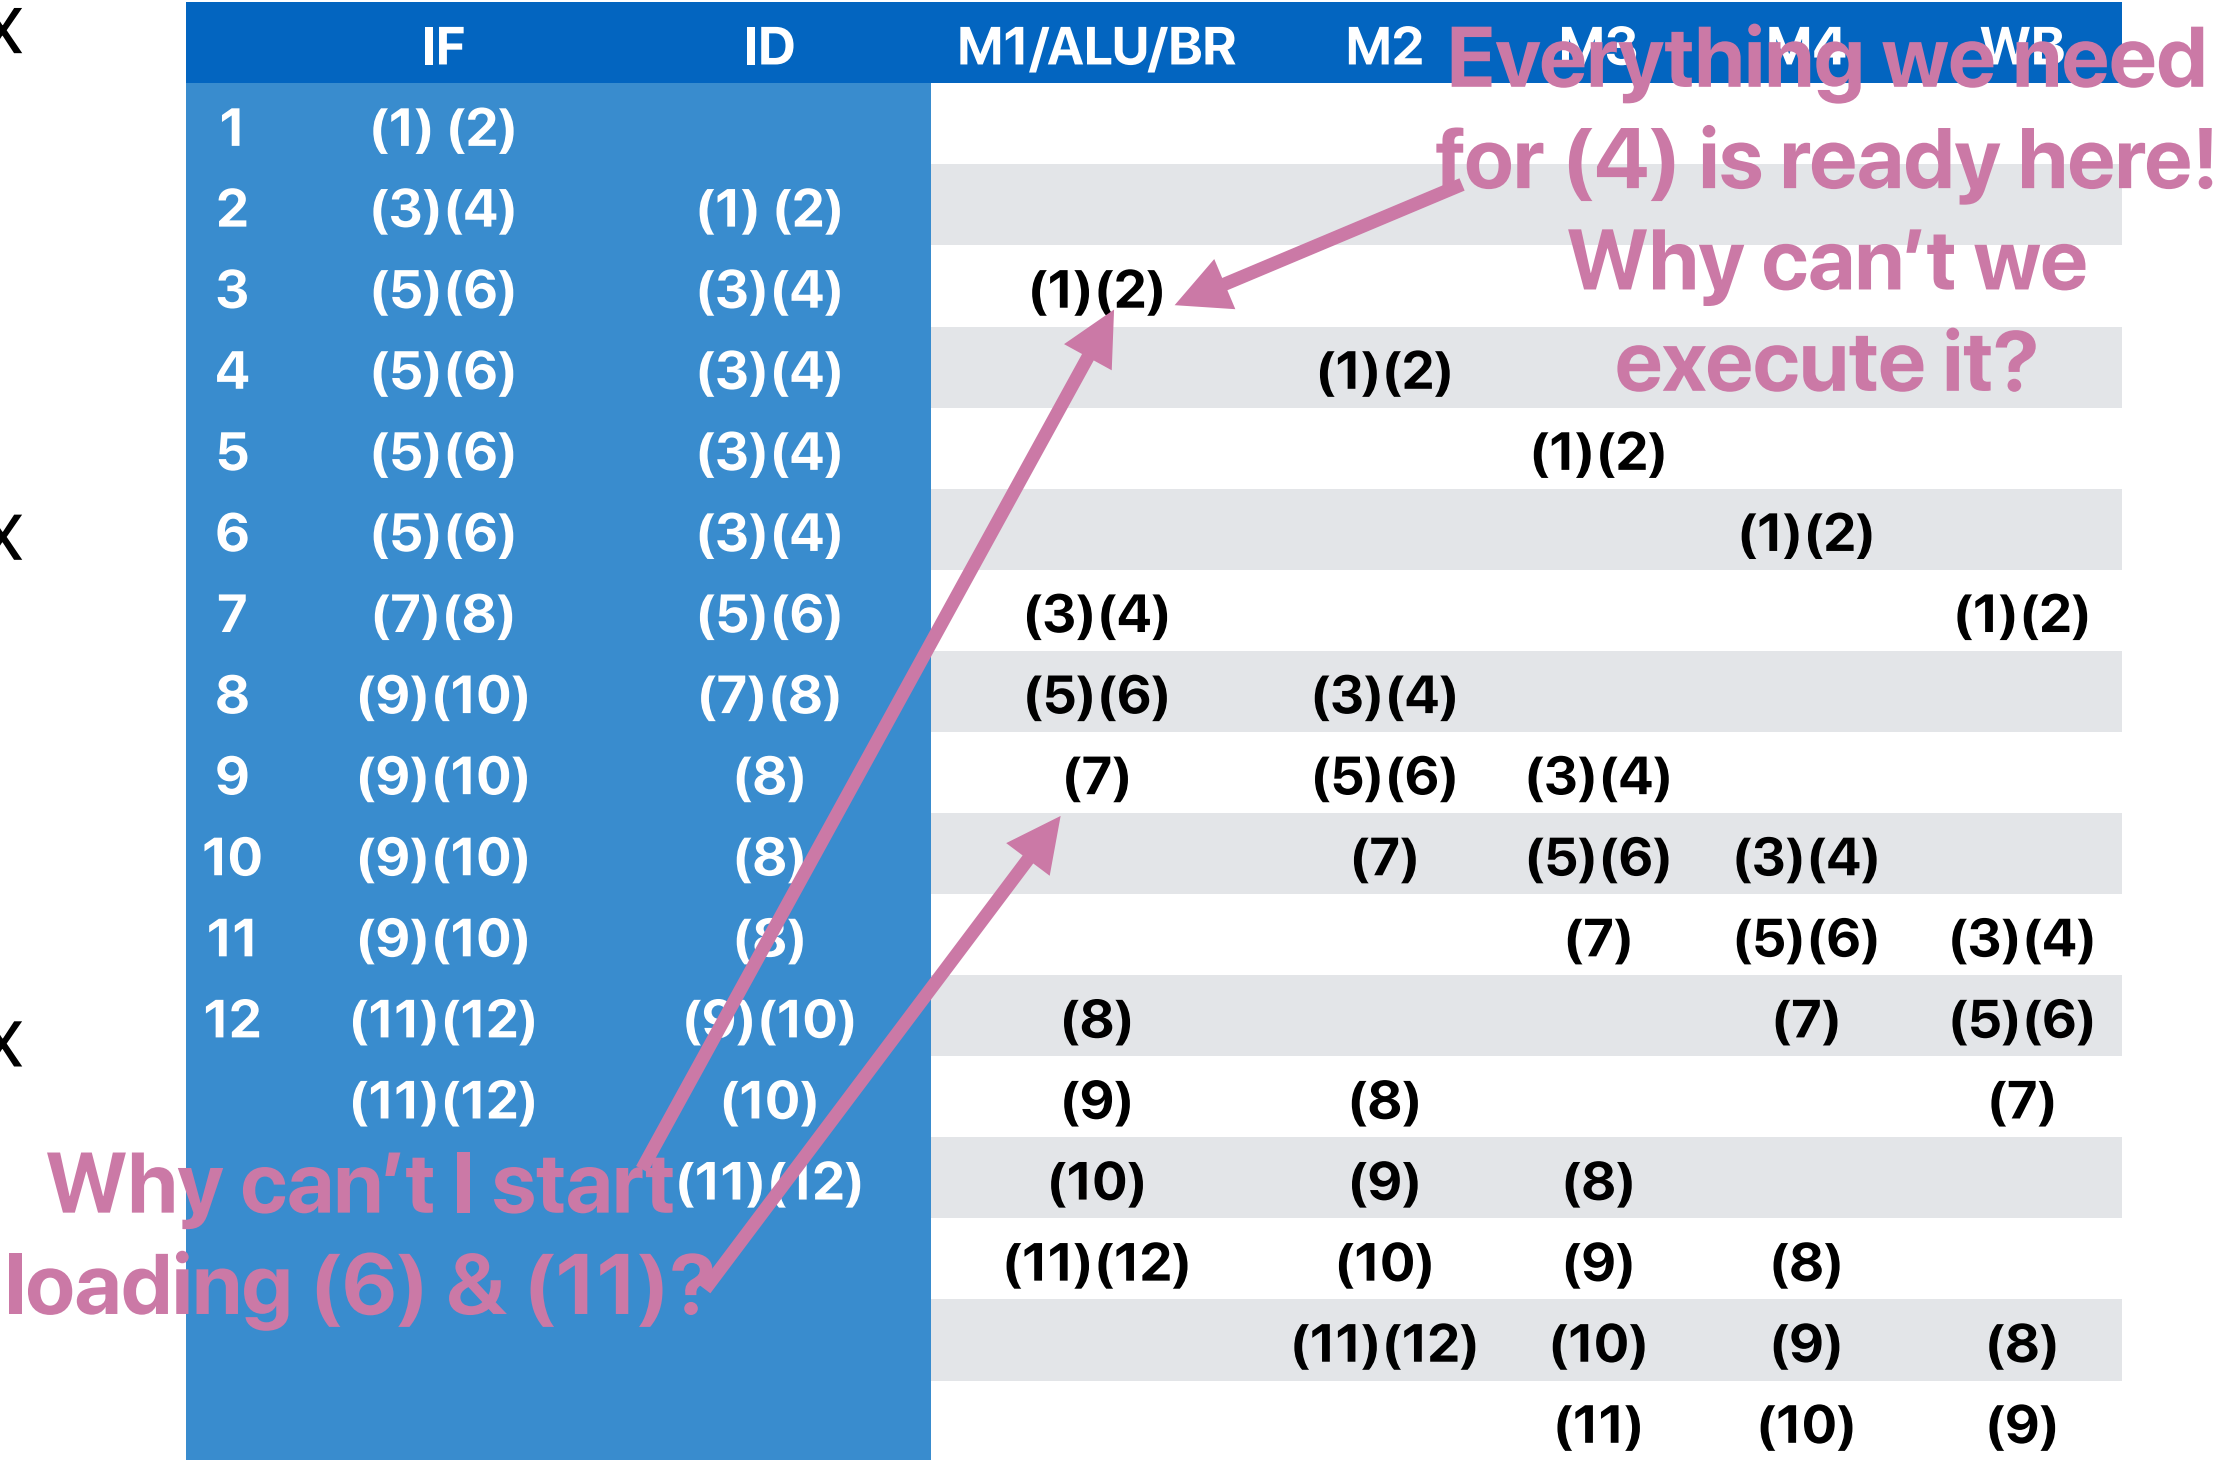
	IF	ID	M1/ALU/BR	M2	M3	M4	WB
1	(1) (2)						
2	(3) (4)	(1) (2)					
3	(5)	(3) (4)	(1) (2)				
4	(5)	(3) (4)		(1) (2)			
5	(5)	(3) (4)			(1) (2)		
6	(5)	(3) (4)				(1) (2)	
7		(5)	(3) (4)				(1) (2)
8			(5)	(3) (4)			
9				(5)	(3) (4)		
10					(5)	(3) (4)	
11						(5)	(3) (4)
12							(5)

Everything we need for (4) is ready here

Why can't we execute it?

If we loop many times (assume perfect predictor)

```
① movl    (%rdi), %ecx
② addq    $4, %rdi
③ addl    %ecx, %eax
④ cmpq    %rdx, %rdi
⑤ jne     .L3
⑥ movl    (%rdi), %ecx
⑦ addq    $4, %rdi
⑧ addl    %ecx, %eax
⑨ cmpq    %rdx, %rdi
⑩ jne     .L3
⑪ movl    (%rdi), %ecx
⑫ addq    $4, %rdi
⑬ addl    %ecx, %eax
⑭ cmpq    %rdx, %rdi
⑮ jne     .L3
```



Limitations of Compiler Optimizations

- If the hardware (e.g., pipeline changes), the same compiler optimization may not be that helpful
- The compiler can only optimize on static instructions, but cannot optimize dynamic instruction
 - Compiler cannot predict branches
 - Compiler does not know if cache has the data/instructions

What do you need to execution an instruction?

- Whenever the instruction is decoded — put decoded instruction somewhere
- Whenever the inputs are ready — **all data dependencies are resolved**
- Whenever the target functional unit is available

Dynamic instruction scheduling/ Out-of-order (OoO) execution


What do you need to execution an instruction?

- Whenever the instruction is decoded — put decoded instruction somewhere
- Whenever the inputs are ready — **all data dependencies are resolved**
- Whenever the target functional unit is available

Scheduling instructions: based on data dependencies

- Draw the data dependency graph, put an arrow if an instruction depends on the other.

```
① movl    (%rdi), %ecx
② addq    $4, %rdi
③ addl    %ecx, %eax
④ cmpq    %rdx, %rdi
⑤ jne     .L3
⑥ movl    (%rdi), %ecx
⑦ addq    $4, %rdi
⑧ addl    %ecx, %eax
⑨ cmpq    %rdx, %rdi
⑩ jne     .L3
```



- In theory**, instructions without dependencies can be executed in parallel or out-of-order
- Instructions with dependencies can never be reordered




If we can predict the future ...

- Consider the following dynamic instructions:


```
① movl    (%rdi), %ecx
② addq    $4, %rdi
③ addl    %ecx, %eax
④ cmpq    %rdx, %rdi
⑤ jne     .L3
⑥ movl    (%rdi), %ecx
⑦ addq    $4, %rdi
⑧ addl    %ecx, %eax
⑨ cmpq    %rdx, %rdi
⑩ jne     .L3
```

Which of the following pair can we reorder without affecting the correctness if the **branch prediction is perfect**?

- A. (1) and (2)
- B. (3) and (4)
- C. (3) and (6)
- D. (4) and (7)
- E. (6) and (7)




0%




A

0%




B

0%




C

0%



D

0%



E




If we can predict the future ...

- Consider the following dynamic instructions:


```
① movl    (%rdi), %ecx
② addq    $4, %rdi
③ addl    %ecx, %eax
④ cmpq    %rdx, %rdi
⑤ jne     .L3
⑥ movl    (%rdi), %ecx
⑦ addq    $4, %rdi
⑧ addl    %ecx, %eax
⑨ cmpq    %rdx, %rdi
⑩ jne     .L3
```

Which of the following pair can we reorder without affecting the correctness if the **branch prediction is perfect**?

- A. (1) and (2)
- B. (3) and (4)
- C. (3) and (6)
- D. (4) and (7)
- E. (6) and (7)




0%




A

0%




B

0%




C

0%



D

0%



E

If we can predict the future ...

- Consider the following dynamic instructions:

```
① movl    (%rdi), %ecx
② addq    $4, %rdi
③ addl    %ecx, %eax
④ cmpq    %rdx, %rdi
⑤ jne     .L3
⑥ movl    (%rdi), %ecx
⑦ addq    $4, %rdi
⑧ addl    %ecx, %eax
⑨ cmpq    %rdx, %rdi
⑩ jne     .L3
```

Can we use "branch prediction" to predict the future and reorder instructions across the branch?

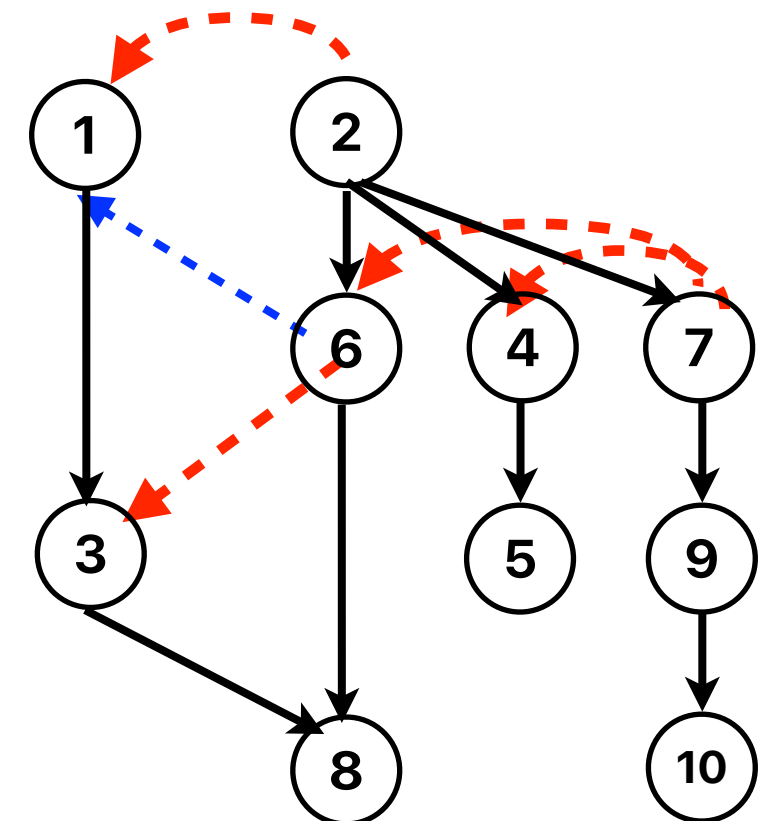
Which of the following pair can we reorder without affecting the correctness if the **branch prediction is perfect**?

- A. (1) and (2)
- B. (3) and (4)
- C. (3) and (6)
- D. (4) and (7)
- E. (6) and (7)

False dependencies

- We are still limited by **false dependencies**
- They are not “true” dependencies because they don’t have an arrow in data dependency graph
 - **WAR (Write After Read):** a later instruction overwrites the source of an earlier one
 - 2 and 1, 6 and 3, 7 and 4, 7 and 6
 - **WAW (Write After Write):** a later instruction overwrites the output of an earlier one
 - 6 and 1

```
①  movl    (%rdi), %ecx
②  addq    $4, %rdi
③  addl    %ecx, %eax
④  cmpq    %rdx, %rdi
⑤  jne     .L3
⑥  movl    (%rdi), %ecx
⑦  addq    $4, %rdi
⑧  addl    %ecx, %eax
⑨  cmpq    %rdx, %rdi
⑩  jne     .L3
```



False dependencies

- We are still limited by **false dependencies**
- They are not “true” dependencies because they don’t have an arrow in data dependency graph


- **WAR (Write After Read):** a later instruction overwrites the source of an earlier one

- 2 and 1, 6 and 3, 7 and 4, 7 and 6

- **WAW (Write After Write):** a later instruction overwrites the output of an earlier one

- 6 and 1

```
①  movl    (%rdi), %ecx
②  addq    $4, %rdi
③  addl    %ecx, %eax
④  cmpq    %rdx, %rdi
⑤  jne     .L3
⑥  movl    (%rdi), %ecx
⑦  addq    $4, %rdi
⑧  addl    %ecx, %eax
⑨  cmpq    %rdx, %rdi
⑩  jne     .L3
```



Limitations of Compiler Optimizations

- If the hardware (e.g., pipeline changes), the same compiler optimization may not be that helpful
- The compiler can only optimize on static instructions, but cannot optimize dynamic instructions
- Compilers are limited by the registers an ISA provides

Recap: False dependencies

- We are still limited by **false dependencies**
- They are not “true” dependencies because they don’t have an arrow in data dependency graph

- **WAR (Write After Read):** a later instruction overwrites the source of an earlier one


- 2 and 1, 6 and 3, 7 and 4, 7 and 6

- **WAW (Write After Write):** a later instruction overwrites the output of an earlier one

We need to give each output a new register!!!


```
① movl    (%rdi), %ecx
② addq    $4, %rdi
③ addq    %ecx, %eax
④ cmpq    %rdx, %rdi
⑤ jne     .L3
⑥ movl    (%rdi), %ecx
⑦ addq    $4, %rdi
⑧ addl    %ecx, %eax
⑨ cmpq    %rdx, %rdi
⑩ jne     .L3
```

80



What if we can use more registers...

```
① movl    (%rdi), %ecx
② addq    $4, %rdi
③ addl    %ecx, %eax
④ cmpq    %rdx, %rdi
⑤ jne     .L3
⑥ movl    (%rdi), %ecx
⑦ addq    $4, %rdi
⑧ addl    %ecx, %eax
⑨ cmpq    %rdx, %rdi
⑩ jne     .L3
```



```
① movl    (%rdi), %ecx
② addq    $4, %rdi, %t0
③ addl    %ecx, %eax, %t1
④ cmpq    %rdx, %t0
⑤ jne     .L3
⑥ movl    (%t0), %t2
⑦ addq    $4, %t0, %t3
⑧ addl    %t1, %t2, %t4
⑨ cmpq    %rdx, %t3
⑩ jne     .L3
```

All false dependencies are gone!!!

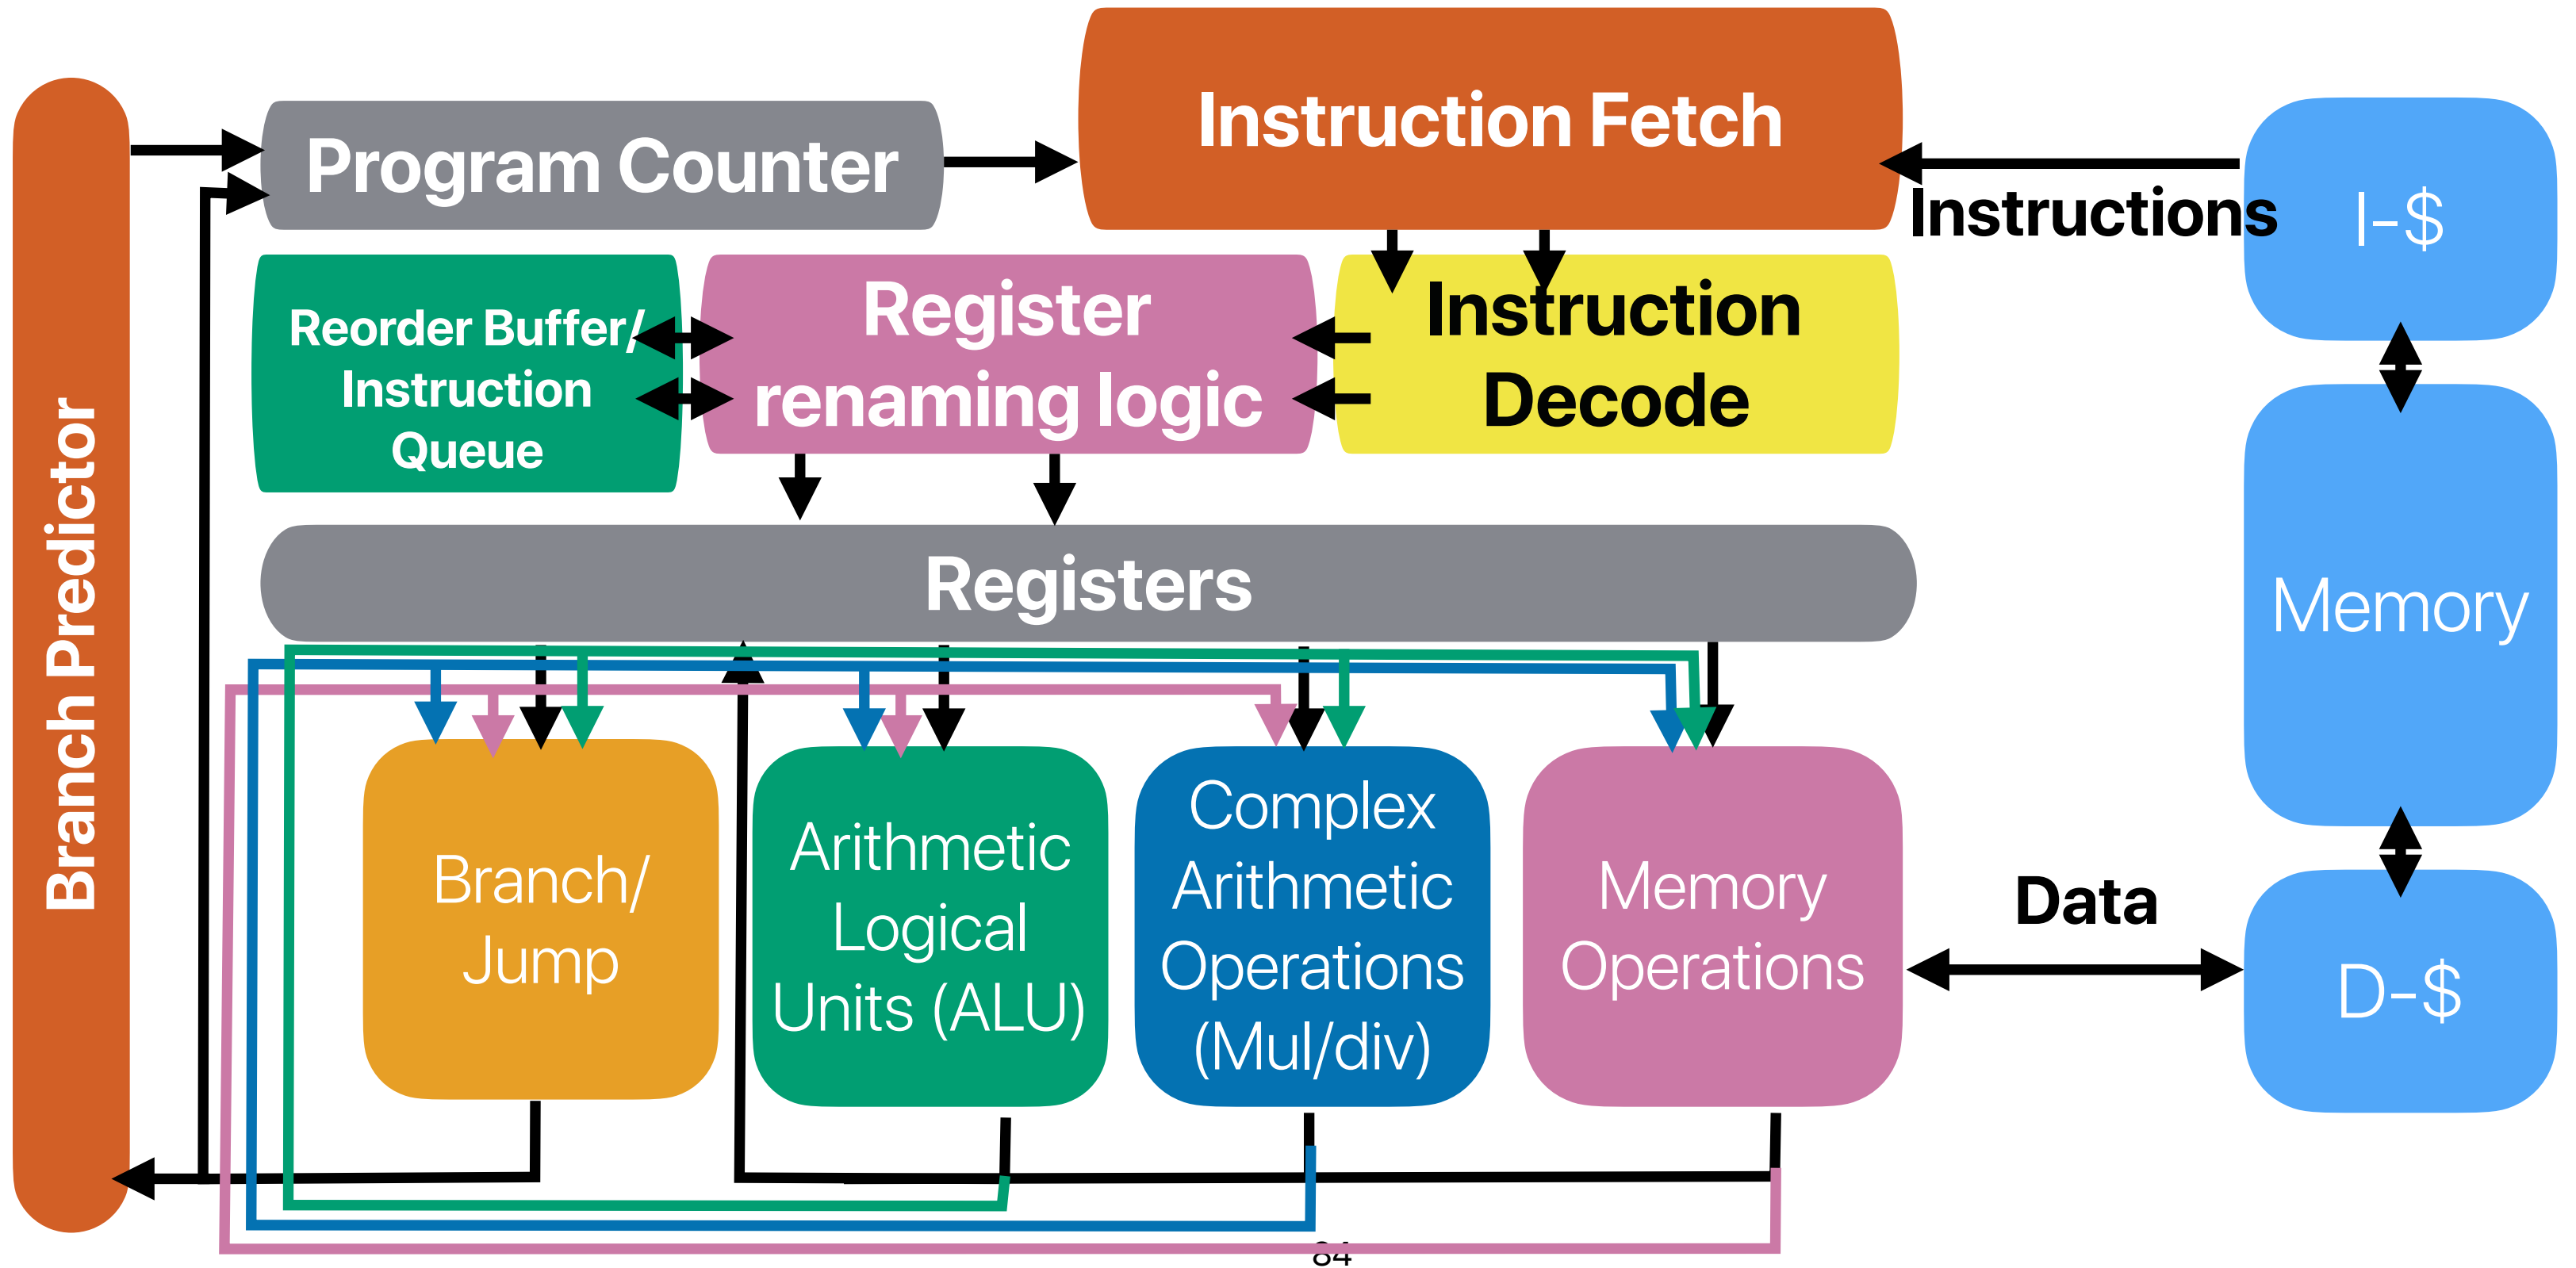
Register renaming + speculative execution

- K. C. Yeager, "The Mips R10000 superscalar microprocessor," in IEEE Micro, vol. 16, no. 2, pp. 28-41, April 1996.

Speculative Execution

- Exceptions (e.g. divided by 0, page fault) may occur anytime
 - A later instruction cannot write back its own result otherwise the architectural states won't be correct
- Hardware can schedule instruction across branch instructions with the help of branch prediction
 - Fetch instructions according to the branch prediction
 - However, branch predictor can never be perfect
- Execute instructions across branches
 - Speculative execution: execute an instruction before the processor know if we need to execute or not
 - Execute an instruction all operands are ready (the values of depending physical registers are generated)
 - Store results in **reorder buffer** before the processor knows if the instruction is going to be executed or not.

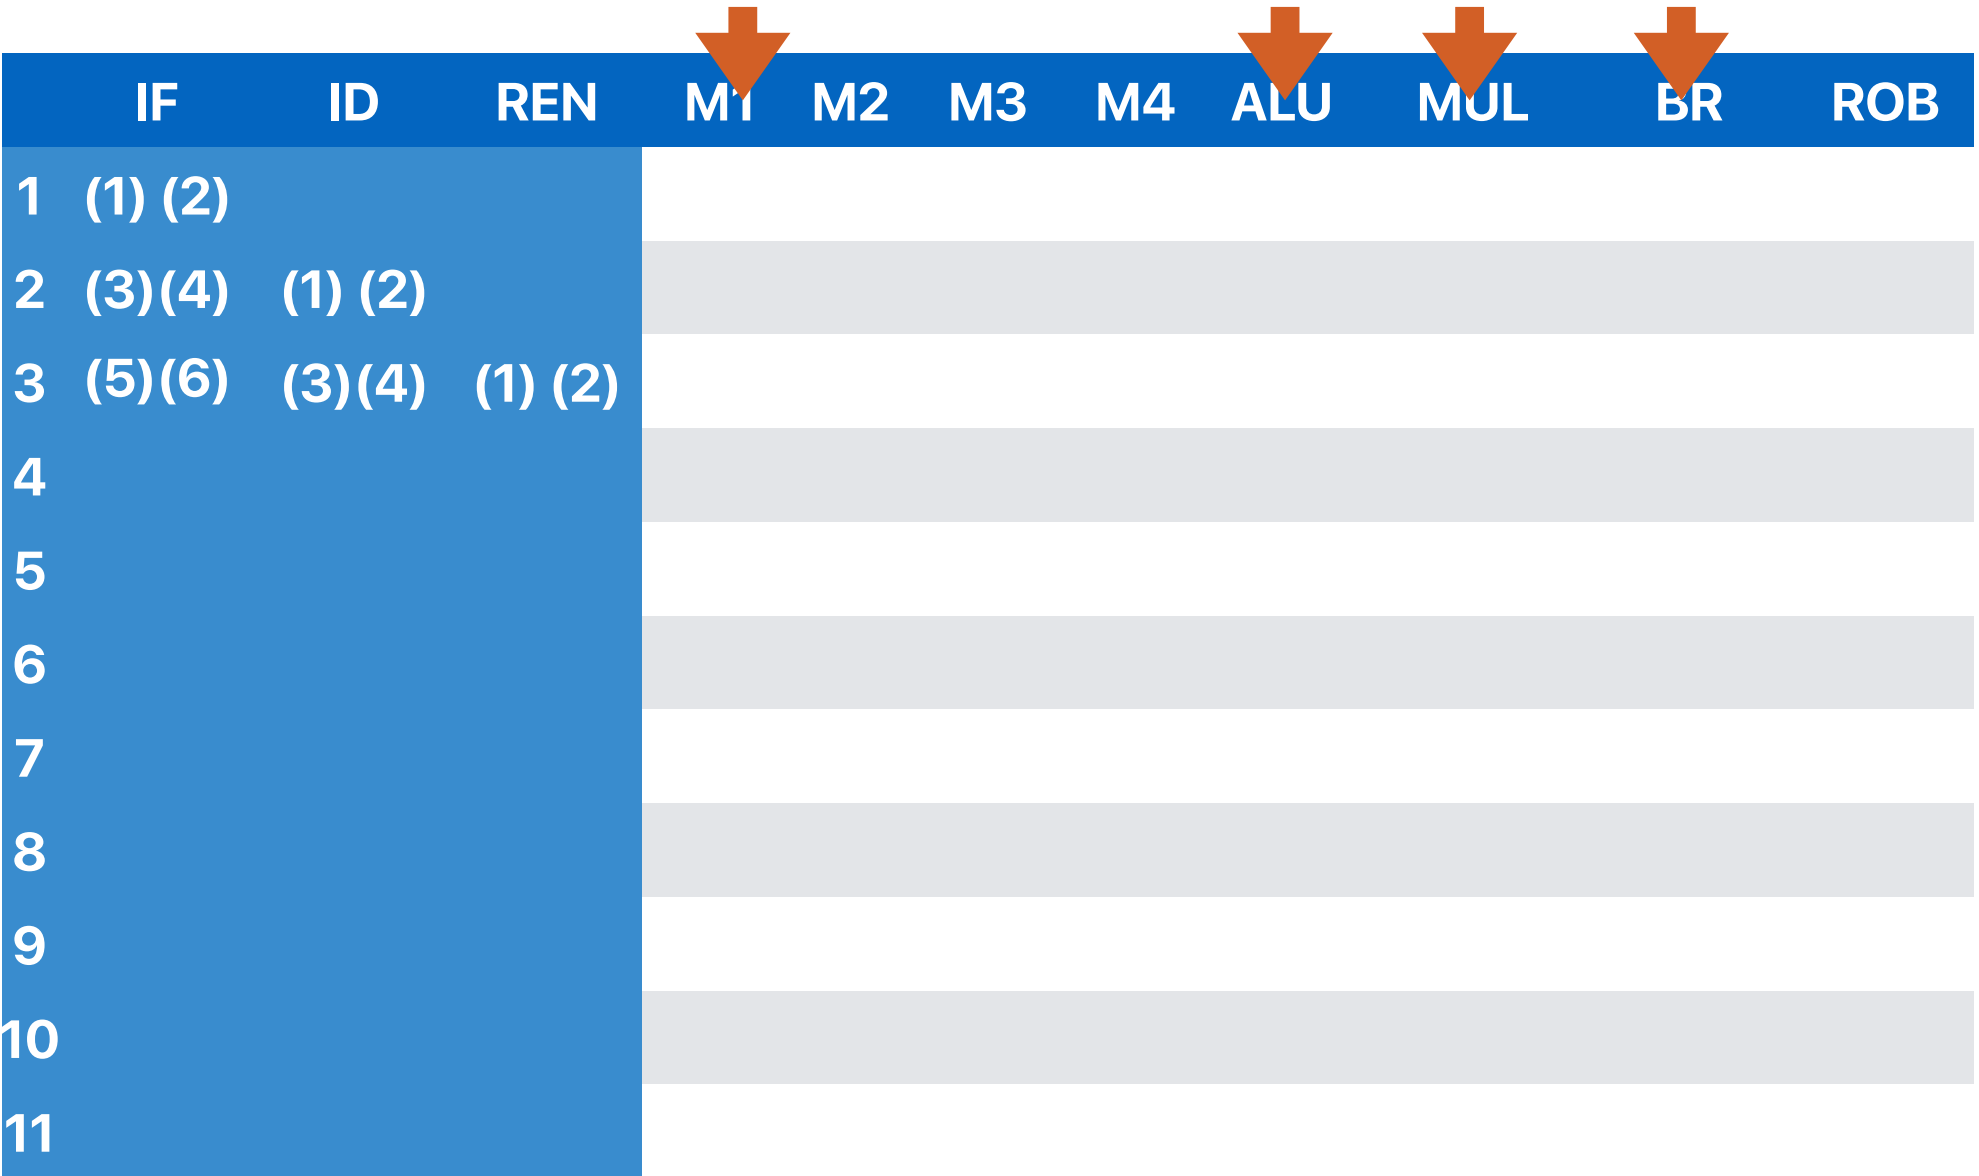
Register renaming



Register renaming

2-issue: Only 2 of them can have instructions at the same cycle

```
① movl    (%rdi), %ecx
② addq    $4, %rdi
③ addl    %ecx, %eax
④ cmpq    %rdx, %rdi
⑤ jne     .L3
⑥ movl    (%rdi), %ecx
⑦ addq    $4, %rdi
⑧ addl    %ecx, %eax
⑨ cmpq    %rdx, %rdi
⑩ jne     .L3
⑪ movl    (%rdi), %ecx
⑫ addq    $4, %rdi
⑬ addl    %ecx, %eax
⑭ cmpq    %rdx, %rdi
⑮ jne     .L3
```



Physical Register	
eax	
ecx	
rdi	
rdx	

Valid	Value	In use	Valid	Value	In use
P1			P6		
P2			P7		
P3			P8		
P4			P9		
P5			P10		

Register renaming

2-issue: Only 2 of them can have instructions at the same cycle

- ① movl (%rdi), %ecx → P1
- ② addq \$4, %rdi → P2
- ③ addl %ecx, %eax
- ④ cmpq %rdx, %rdi
- ⑤ jne .L3
- ⑥ movl (%rdi), %ecx
- ⑦ addq \$4, %rdi
- ⑧ addl %ecx, %eax
- ⑨ cmpq %rdx, %rdi
- ⑩ jne .L3
- ⑪ movl (%rdi), %ecx
- ⑫ addq \$4, %rdi
- ⑬ addl %ecx, %eax
- ⑭ cmpq %rdx, %rdi
- ⑮ jne .L3

	IF	ID	REN	M1	M2	M3	M4	ALU	MUL	BR	ROB
1	(1)	(2)									
2	(3)	(4)	(1)	(2)							
3	(5)	(6)	(3)	(4)	(1)	(2)					
4		(5)	(6)	(3)	(4)						
5											
6											
7											
8											
9											
10											
11											

Physical Register	
eax	
ecx	P1
rdi	P2
rdx	

	Valid	Value	In use		Valid	Value	In use
P1	0		1	P6			
P2	0		1	P7			
P3				P8			
P4				P9			
P5				P10			

Register renaming

2-issue: Only 2 of them can have instructions at the same cycle

- ① movl (%rdi), %ecx → P1
- ② addq \$4, %rdi → P2
- ③ addl %ecx, %eax → P3
- ④ cmpq %rdx, %rdi
- ⑤ jne .L3
- ⑥ movl (%rdi), %ecx → P4
- ⑦ addq \$4, %rdi
- ⑧ addl %ecx, %eax
- ⑨ cmpq %rdx, %rdi
- ⑩ jne .L3
- ⑪ movl (%rdi), %ecx
- ⑫ addq \$4, %rdi
- ⑬ addl %ecx, %eax
- ⑭ cmpq %rdx, %rdi
- ⑮ jne .L3

	IF	ID	REN	M1	M2	M3	M4	ALU	MUL	BR	ROB
1	(1)	(2)									
2	(3)	(4)	(1)	(2)							
3	(5)	(6)	(3)	(4)	(1)			(2)			
4	(7)	(8)	(5)	(6)	(3)	(4)					
5	(9)	(10)	(7)	(8)	(3)	(5)	(6)	(4)			(2)
6											
7											
8											
9											
10											
11											

(4) is now
executing before
(3)!

Physical Register	
eax	
ecx	P1
rdi	P2
rdx	P4

	Valid	Value	In use		Valid	Value	In use
P1	0		1	P6			
P2	1		1	P7			
P3	0		1	P8			
P4	0		1	P9			
P5				P10			

Register renaming

2-issue: Only 2 of them can have instructions at the same cycle

- ① movl (%rdi), %ecx → P1

② addq \$4, %rdi → P2

③ addl %ecx, %eax → P3

④ cmpq %rdx, %rdi

⑤ jne .L3

⑥ movl (%rdi), %ecx → P4

⑦ addq \$4, %rdi → P5

⑧ addl %ecx, %eax → P6

⑨ cmpq %rdx, %rdi

⑩ jne .L3

⑪ movl (%rdi), %ecx

⑫ addq \$4, %rdi

⑬ addl %ecx, %eax

⑭ cmpq %rdx, %rdi

⑮ jne .L3

	IF	ID	REN	M1	M2	M3	M4	ALU	MUL	BR	ROB
1	(1)	(2)									
2	(3)(4)	(1)(2)									
3	(5)(6)	(3)(4)	(1)(2)								
4	(7)(8)	(5)(6)	(3)(4)	(1)				(2)			
5	(9)(10)	(7)(8)	(3)(5)(6)		(1)			(4)			(2)
6	(11)(12)	(9)(10)	(3)(7)(8)	(6)		(1)				(5)	(2)(4)
7											
8											
9											
10											
11											

Physical Register	
eax	P6
ecx	P1
rdi	P5
rdx	P4

	Valid	Value	In use		Valid	Value	In use
P1	0		1	P6	0		1
P2	1		1	P7			
P3	0		1	P8			
P4	0		1	P9			
P5	0		1	P10			

Register renaming

2-issue: Only 2 of them can have instructions at the same cycle

- ①

movl

(%rdi), %ecx → P1
- ②

addq

\$4, %rdi → P2
- ③

addl

%ecx, %eax → P3
- ④

cmpq

%rdx, %rdi
- ⑤

jne

.L3
- ⑥

movl

(%rdi), %ecx → P4
- ⑦

addq

\$4, %rdi → P5
- ⑧

addl

%ecx, %eax → P6
- ⑨

cmpq

%rdx, %rdi
- ⑩

jne

.L3
- ⑪

movl

(%rdi), %ecx
- ⑫

addq

\$4, %rdi
- ⑬

addl

%ecx, %eax
- ⑭

cmpq

%rdx, %rdi
- ⑮

jne

.L3

	IF	ID	REN	M1	M2	M3	M4	ALU	MUL	BR	ROB
1	(1)	(2)									
2	(3)(4)	(1)(2)									
3	(5)(6)	(3)(4)	(1)(2)								
4	(7)(8)	(5)(6)	(3)(4)	(1)				(2)			
5	(9)(10)	(7)(8)	(3)(5)(6)		(1)			(4)			(2)
6	(11)(12)	(9)(10)	(3)(7)(8)	(6)		(1)				(5)	(2)(4)
7	(13)(14)	(11)(12)	(3)(8)(9)(10)		(6)		(1)	(7)			(2)(4)(5)
8											
9											
10											
11											

Physical Register	
eax	P6
ecx	P1
rdi	P5
rdx	P4

	Valid	Value	In use		Valid	Value	In use
P1	0		1	P6	0		1
P2	1		1	P7			
P3	0		1	P8			
P4	0		1	P9			
P5	0		1	P10			

Register renaming

2-issue: Only 2 of them can have instructions at the same cycle

- ①

movl

(%rdi), %ecx → P1
- ②

addq

\$4, %rdi → P2
- ③

addl

%ecx, %eax → P3
- ④

cmpq

%rdx, %rdi
- ⑤

jne

.L3
- ⑥

movl

(%rdi), %ecx → P4
- ⑦

addq

\$4, %rdi → P5
- ⑧

addl

%ecx, %eax → P6
- ⑨

cmpq

%rdx, %rdi
- ⑩

jne

.L3
- ⑪

movl

(%rdi), %ecx → P7
- ⑫

addq

\$4, %rdi → P8
- ⑬

addl

%ecx, %eax
- ⑭

cmpq

%rdx, %rdi
- ⑮

jne

.L3

	IF	ID	REN	M1	M2	M3	M4	ALU	MUL	BR	ROB				
1	(1)	(2)													
2	(3)	(4)	(1)	(2)											
3	(5)	(6)	(3)	(4)	(1)	(2)									
4	(7)	(8)	(5)	(6)	(3)	(4)	(1)	(2)							
5	(9)	(10)	(7)	(8)	(3)	(5)	(6)	(4)			(2)				
6	(11)	(12)	(9)	(10)	(3)	(7)	(8)			(5)	(2)	(4)			
7	(13)	(14)	(11)	(12)	(3)	(8)	(9)	(10)			(2)	(4)	(5)		
8	(15)	(16)	(13)	(14)	(8)	(9)	(10)	(11)	(12)		(1)	(2)	(4)	(5)	(7)
9															
10															
11															

Physical Register	
eax	P6
ecx	P7
rdi	P8
rdx	P4

	Valid	Value	In use		Valid	Value	In use
P1	1		1	P6	0		1
P2	1		1	P7	0		1
P3	0		1	P8	0		1
P4	0		1	P9			
P5	1		1	P10			

Register renaming

2-issue: Only 2 of them can have instructions at the same cycle

- ① movl (%rdi), %ecx → P1
- ② addq \$4, %rdi → P2
- ③ addl %ecx, %eax → P3
- ④ cmpq %rdx, %rdi
- ⑤ jne .L3
- ⑥ movl (%rdi), %ecx → P4
- ⑦ addq \$4, %rdi → P5
- ⑧ addl %ecx, %eax → P6
- ⑨ cmpq %rdx, %rdi
- ⑩ jne .L3
- ⑪ movl (%rdi), %ecx → P7
- ⑫ addq \$4, %rdi → P8
- ⑬ addl %ecx, %eax → P9
- ⑭ cmpq %rdx, %rdi
- ⑮ jne .L3

	IF	ID	REN	M1	M2	M3	M4	ALU	MUL	BR	ROB
1	(1)	(2)									
2	(3)	(4)	(1)								
3	(5)	(6)	(3)								
4	(7)	(8)	(5)	(1)				(2)			
5	(9)	(10)	(7)		(1)			(4)			(2)
6	(11)	(12)	(9)	(6)		(1)				(5)	(2)(4)
7	(13)	(14)	(11)		(6)		(1)	(7)			(2)(4)(5)
8	(15)	(16)	(13)			(6)		(3)			(1)(2)(4)(5) (7)
9	(17)	(18)	(15)	(11)			(6)	(9)			(3)(4)(5)(7)
10											
11											

Physical Register	
eax	P9
ecx	P7
rdi	P8
rdx	P4

	Valid	Value	In use		Valid	Value	In use
P1	1		0	P6	0		1
P2	1		0	P7	0		1
P3	1		1	P8	0		1
P4	0		1	P9	0		1
P5	1		1	P10			

Register renaming

2-issue: Only 2 of them can have instructions at the same cycle

- ①

movl

(%rdi), %ecx → P1
- ②

addq

\$4, %rdi → P2
- ③

addl

%ecx, %eax → P3
- ④

cmpq

%rdx, %rdi
- ⑤

jne

.L3
- ⑥

movl

(%rdi), %ecx → P4
- ⑦

addq

\$4, %rdi → P5
- ⑧

addl

%ecx, %eax → P6
- ⑨

cmpq

%rdx, %rdi
- ⑩

jne

.L3
- ⑪

movl

(%rdi), %ecx → P7
- ⑫

addq

\$4, %rdi → P8
- ⑬

addl

%ecx, %eax → P9
- ⑭

cmpq

%rdx, %rdi
- ⑮

jne

.L3

	IF	ID	REN	M1	M2	M3	M4	ALU	MUL	BR	ROB
1	(1)	(2)									
2	(3)(4)	(1)(2)									
3	(5)(6)	(3)(4)	(1)(2)								
4	(7)(8)	(5)(6)	(3)(4)	(1)				(2)			
5	(9)(10)	(7)(8)	(3)(5)(6)		(1)			(4)			(2)
6	(11)(12)	(9)(10)	(3)(7)(8)	(6)		(1)				(5)	(2)(4)
7	(13)(14)	(11)(12)	(3)(8)(9)(10)		(6)		(1)	(7)			(2)(4)(5)
8	(15)(16)	(13)(14)	(8)(9)(10)(11)(12)			(6)		(3)			(1)(2)(4)(5)(7)
9	(17)(18)	(15)(16)	(8)(10)(12)(13)(14)	(11)			(6)	(9)			(3)(4)(5)(7)
10	(19)(20)	(17)(18)	(12)(13)(14)(15)(16)		(11)			(8)		(10)	(6)(7)(9)
11											

Physical Register	
eax	P9
ecx	P7
rdi	P8
rdx	P4

	Valid	Value	In use		Valid	Value	In use
P1	1		0	P6	0		1
P2	1		0	P7	0		1
P3	1		0	P8	0		1
P4	0		1	P9	0		1
P5	1		1	P10	0		1

Register renaming

2-issue: Only 2 of them can have instructions at the same cycle

- ① movl (%rdi), %ecx → P1
- ② addq \$4, %rdi → P2
- ③ addl %ecx, %eax → P3
- ④ cmpq %rdx, %rdi
- ⑤ jne .L3
- ⑥ movl (%rdi), %ecx → P4
- ⑦ addq \$4, %rdi → P5
- ⑧ addl %ecx, %eax → P6
- ⑨ cmpq %rdx, %rdi
- ⑩ jne .L3
- ⑪ movl (%rdi), %ecx → P7
- ⑫ addq \$4, %rdi → P8
- ⑬ addl %ecx, %eax → P9
- ⑭ cmpq %rdx, %rdi
- ⑮ jne .L3

	IF	ID	REN	M1	M2	M3	M4	ALU	MUL	BR	ROB
1	(1)	(2)									
2	(3)(4)	(1)(2)									
3	(5)(6)	(3)(4)	(1)(2)								
4	(7)(8)	(5)(6)	(3)(4)	(1)				(2)			
5	(9)(10)	(7)(8)	(3)(5)(6)		(1)			(4)			(2)
6	(11)(12)	(9)(10)	(3)(7)(8)	(6)		(1)				(5)	(2)(4)
7	(13)(14)	(11)(12)	(3)(8)(9)(10)		(6)		(1)	(7)			(2)(4)(5)
8	(15)(16)	(13)(14)	(8)(9)(10)(11)(12)			(6)		(3)			(1)(2)(4)(5)(7)
9	(17)(18)	(15)(16)	(8)(10)(12)(13)(14)	(11)			(6)	(9)			(3)(4)(5)(7)
10	(19)(20)	(17)(18)	(12)(13)(14)(15)(16)		(11)			(8)		(10)	(6)(7)(9)
11		(19)(20)	(13)(14)(15)(16)(17)(18)			(11)		(12)			(8)(9)(10)

Physical Register	
eax	P6
ecx	P1
rdi	P5
rdx	P4

	Valid	Value	In use		Valid	Value	In use
P1	1		0	P6	1		1
P2	1		0	P7	0		1
P3	1		0	P8	0		1
P4	1		0	P9	0		1
P5	1		1	P10	0		1

Register renaming

2-issue: Only 2 of them can have instructions at the same cycle

- ①

movl

(%rdi), %ecx → P1
- ②

addq

\$4, %rdi → P2
- ③

addl

%ecx, %eax → P3
- ④

cmpq

%rdx, %rdi
- ⑤

jne

.L3
- ⑥

movl

(%rdi), %ecx → P4
- ⑦

addq

\$4, %rdi → P5
- ⑧

addl

%ecx, %eax → P6
- ⑨

cmpq

%rdx, %rdi
- ⑩

jne

.L3
- ⑪

movl

(%rdi), %ecx → P7
- ⑫

addq

\$4, %rdi → P8
- ⑬

addl

%ecx, %eax → P9
- ⑭

cmpq

%rdx, %rdi
- ⑮

jne

.L3

	IF	ID	REN	M1	M2	M3	M4	ALU	MUL	BR	ROB
1	(1)	(2)									
2	(3)(4)	(1)(2)									
3	(5)(6)	(3)(4)	(1)(2)								
4	(7)(8)	(5)(6)	(3)(4)	(1)				(2)			
5	(9)(10)	(7)(8)	(3)(5)(6)		(1)			(4)			(2)
6	(11)(12)	(9)(10)	(3)(7)(8)	(6)		(1)				(5)	(2)(4)
7	(13)(14)	(11)(12)	(3)(8)(9)(10)		(6)		(1)	(7)			(2)(4)(5)
8	(15)(16)	(13)(14)	(8)(9)(10)(11)(12)			(6)		(3)			(1)(2)(4)(5)(7)
9	(17)(18)	(15)(16)	(8)(10)(12)(13)(14)	(11)			(6)	(9)			(3)(4)(5)(7)
10	(19)(20)	(17)(18)	(12)(13)(14)(15)(16)		(11)			(8)		(10)	(6)(7)(9)
11		(19)(20)	(13)(14)(15)(16)(17)(18)			(11)		(12)			(8)(9)(10)
12			(13)(15)(17)(18)(19)(20)	(16)			(11)	(14)			(12)
13											
14											
15											

Register renaming

2-issue: Only 2 of them can have instructions at the same cycle

- ①

movl

(%rdi), %ecx → P1
- ②

addq

\$4, %rdi → P2
- ③

addl

%ecx, %eax → P3
- ④

cmpq

%rdx, %rdi
- ⑤

jne

.L3
- ⑥

movl

(%rdi), %ecx → P4
- ⑦

addq

\$4, %rdi → P5
- ⑧

addl

%ecx, %eax → P6
- ⑨

cmpq

%rdx, %rdi
- ⑩

jne

.L3
- ⑪

movl

(%rdi), %ecx → P7
- ⑫

addq

\$4, %rdi → P8
- ⑬

addl

%ecx, %eax → P9
- ⑭

cmpq

%rdx, %rdi
- ⑮

jne

.L3

	IF	ID	REN	M1	M2	M3	M4	ALU	MUL	BR	ROB
1	(1)	(2)									
2	(3)(4)	(1)(2)									
3	(5)(6)	(3)(4)	(1)(2)								
4	(7)(8)	(5)(6)	(3)(4)	(1)				(2)			
5	(9)(10)	(7)(8)	(3)(5)(6)		(1)			(4)			(2)
6	(11)(12)	(9)(10)	(3)(7)(8)	(6)		(1)				(5)	(2)(4)
7	(13)(14)	(11)(12)	(3)(8)(9)(10)		(6)		(1)	(7)			(2)(4)(5)
8	(15)(16)	(13)(14)	(8)(9)(10)(11)(12)			(6)		(3)			(1)(2)(4)(5)(7)
9	(17)(18)	(15)(16)	(8)(10)(12)(13)(14)	(11)			(6)	(9)			(3)(4)(5)(7)
10	(19)(20)	(17)(18)	(12)(13)(14)(15)(16)		(11)			(8)		(10)	(6)(7)(9)
11		(19)(20)	(13)(14)(15)(16)(17)(18)			(11)		(12)			(8)(9)(10)
12			(13)(15)(17)(18)(19)(20)	(16)			(11)	(14)			(12)
13			(17)(18)(19)(20)		(16)			(13)		(15)	(11)(12)(14)
14											
15											

Register renaming

2-issue: Only 2 of them can have instructions at the same cycle

- ①

movl

(%rdi), %ecx → P1
- ②

addq

\$4, %rdi → P2
- ③

addl

%ecx, %eax → P3
- ④

cmpq

%rdx, %rdi
- ⑤

jne

.L3
- ⑥

movl

(%rdi), %ecx → P4
- ⑦

addq

\$4, %rdi → P5
- ⑧

addl

%ecx, %eax → P6
- ⑨

cmpq

%rdx, %rdi
- ⑩

jne

.L3
- ⑪

movl

(%rdi), %ecx → P7
- ⑫

addq

\$4, %rdi → P8
- ⑬

addl

%ecx, %eax → P9
- ⑭

cmpq

%rdx, %rdi
- ⑮

jne

.L3

	IF	ID	REN	M1	M2	M3	M4	ALU	MUL	BR	ROB
1	(1)	(2)									
2	(3)(4)	(1)(2)									
3	(5)(6)	(3)(4)	(1)(2)								
4	(7)(8)	(5)(6)	(3)(4)	(1)				(2)			
5	(9)(10)	(7)(8)	(3)(5)(6)		(1)			(4)			(2)
6	(11)(12)	(9)(10)	(3)(7)(8)	(6)		(1)				(5)	(2)(4)
7	(13)(14)	(11)(12)	(3)(8)(9)(10)		(6)		(1)	(7)			(2)(4)(5)
8	(15)(16)	(13)(14)	(8)(9)(10)(11)(12)			(6)		(3)			(1)(2)(4)(5)(7)
9	(17)(18)	(15)(16)	(8)(10)(12)(13)(14)	(11)			(6)	(9)			(3)(4)(5)(7)
10	(19)(20)	(17)(18)	(12)(13)(14)(15)(16)		(11)			(8)		(10)	(6)(7)(9)
11		(19)(20)	(13)(14)(15)(16)(17)(18)			(11)		(12)			(8)(9)(10)
12			(13)(15)(17)(18)(19)(20)	(16)			(11)	(14)			(12)
13			(17)(18)(19)(20)		(16)			(13)		(15)	(11)(12) (14)
14						(16)		(17)			(13)(14)(15)
15											

Register renaming

2-issue: Only 2 of them can have instructions at the same cycle

- ①

movl

(%rdi), %ecx → P1
- ②

addq

\$4, %rdi → P2
- ③

addl

%ecx, %eax → P3
- ④

cmpq

%rdx, %rdi
- ⑤

jne

.L3
- ⑥

movl

(%rdi), %ecx → P4
- ⑦

addq

\$4, %rdi → P5
- ⑧

addl

%ecx, %eax → P6
- ⑨

cmpq

%rdx, %rdi
- ⑩

jne

.L3
- ⑪

movl

(%rdi), %ecx → P7
- ⑫

addq

\$4, %rdi → P8
- ⑬

addl

%ecx, %eax → P9
- ⑭

cmpq

%rdx, %rdi
- ⑮

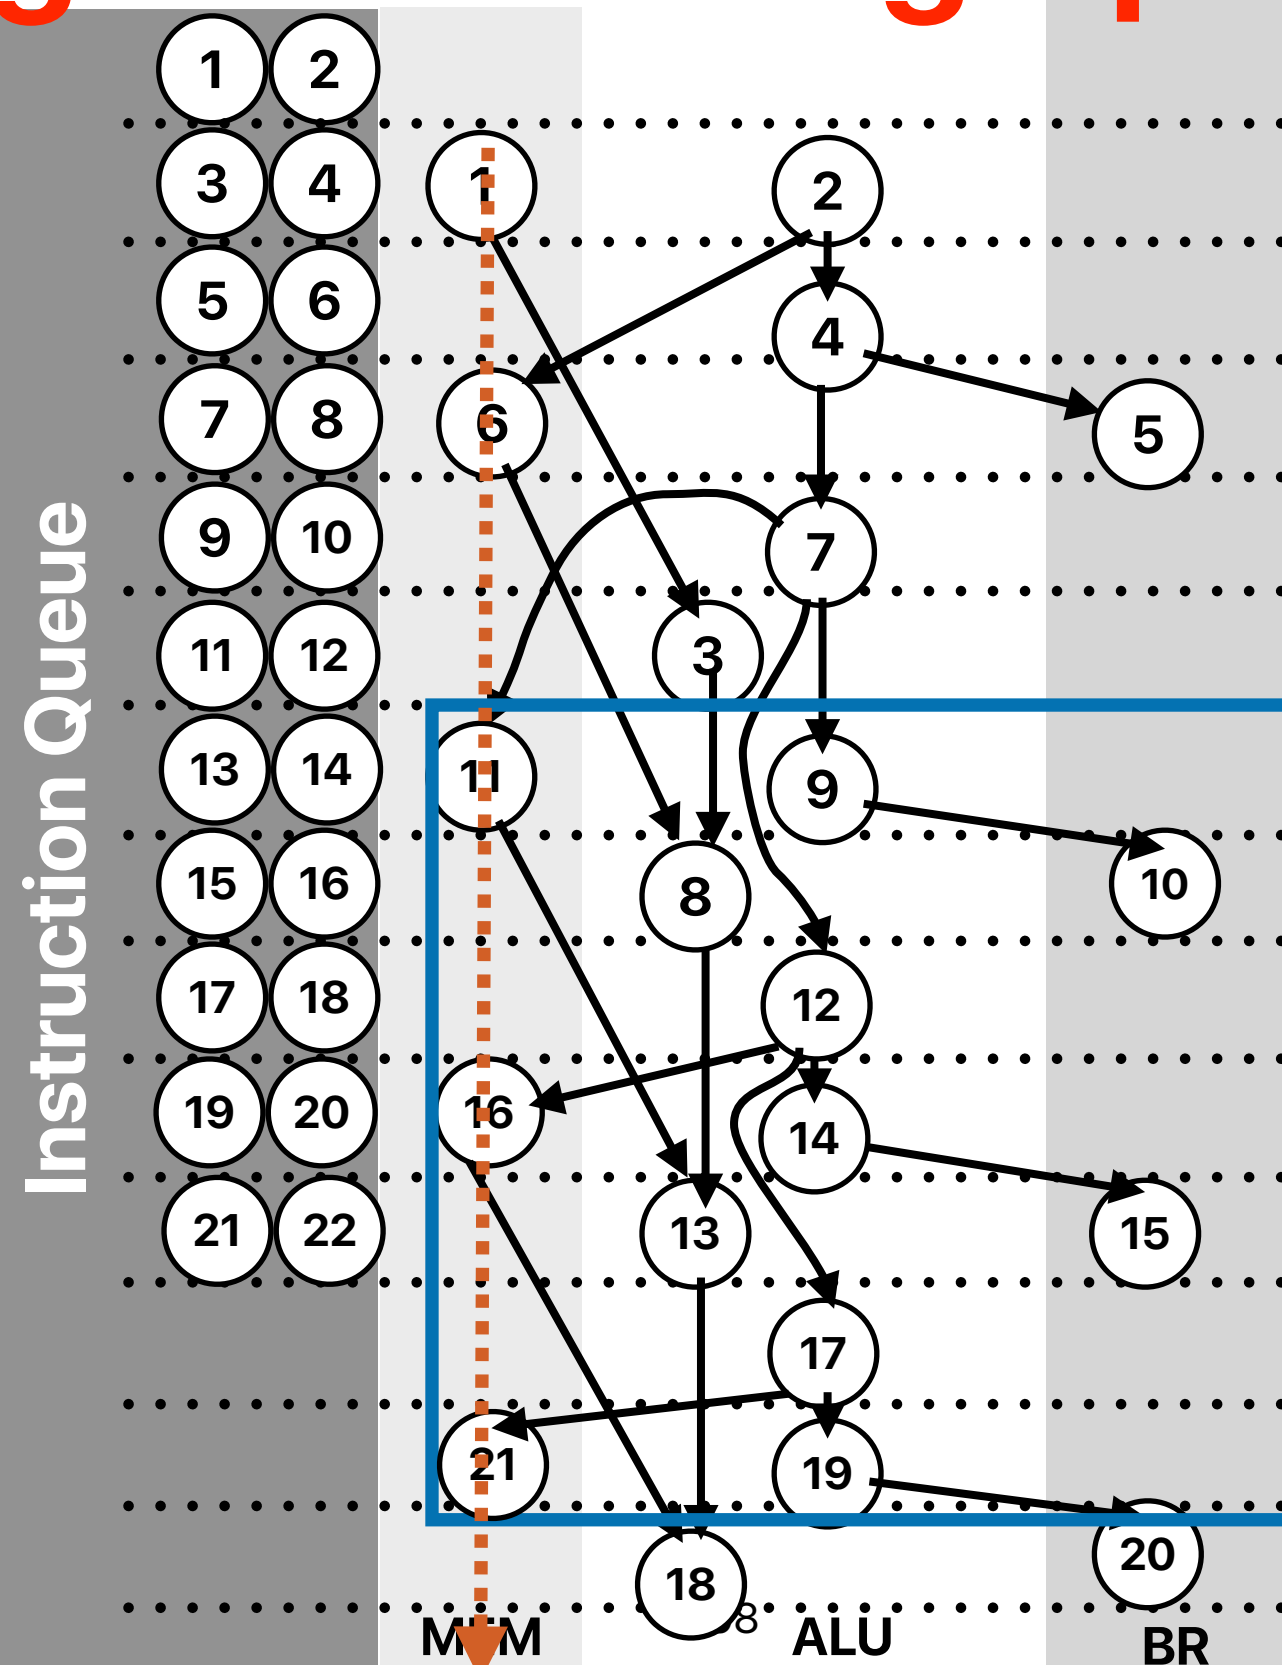
jne

.L3

	IF	ID	REN	M1	M2	M3	M4	ALU	MUL	BR	ROB
1	(1)	(2)									
2	(3)(4)	(1)(2)									
3	(5)(6)	(3)(4)	(1)(2)								
4	(7)(8)	(5)(6)	(3)(4)	(1)				(2)			
5	(9)(10)	(7)(8)	(3)(5)(6)		(1)			(4)			(2)
6	(11)(12)	(9)(10)	(3)(7)(8)	(6)		(1)				(5)	(2)(4)
7	(13)(14)	(11)(12)	(3)(8)(9)(10)		(6)		(1)	(7)			(2)(4)(5)
8	(15)(16)	(13)(14)	(8)(9)(10)(11)(12)			(6)		(3)			(1)(2)(4)(5)(7)
9	(17)(18)	(15)(16)	(8)(10)(12)(13)(14)	(11)			(6)	(9)			(3)(4)(5)(7)
10	(19)(20)	(17)(18)	(12)(13)(14)(15)(16)		(11)			(8)		(10)	(6)(7)(9)
11		(19)(20)	(13)(14)(15)(16)(17)(18)			(11)		(12)			(8)(9)(10)
12			(13)(15)(17)(18)(19)(20)	(16)			(11)	(14)			(12)
13			(17)(18)(19)(20)		(16)			(13)		(15)	(11)(12)(14)
14						(16)		(17)			(13)(14)(15)
15							(16)	(19)			(17)

Through data flow graph analysis

```
① movl (%rdi), %ecx
② addq $4, %rdi
③ addl %ecx, %eax
④ cmpq %rdx, %rdi
⑤ jne .L3
⑥ movl (%rdi), %ecx
⑦ addq $4, %rdi
⑧ addl %ecx, %eax
⑨ cmpq %rdx, %rdi
⑩ jne .L3
⑪ movl (%rdi), %ecx
⑫ addq $4, %rdi
⑬ addl %ecx, %eax
⑭ cmpq %rdx, %rdi
⑮ jne .L3
⑯ movl (%rdi), %ecx
⑰ addq $4, %rdi
⑱ addl %ecx, %eax
⑲ cmpq %rdx, %rdi
⑳ jne .L3
㉑ movl (%rdi), %ecx
```



Execution time is determined
by the "critical path"
composed by 1, 6, 11, ..., 1+5n

3 cycles every iteration
$$CPI = \frac{3}{5} = 0.6!$$



What about "linked list"


- Assume the current PC is already at instruction (1) and this linked list has only three nodes. This processor can fetch and issue 2 instructions per cycle, with exactly the same register renaming hardware and pipeline as we showed previously.

Which of the following C state of the code snippet determines the performance?

- A. do {
- B. number_of_nodes++;
- C. current = current->next;
- D. } while (current != NULL);




0%




A

0%




B

0%




C

0%



D

0%



E

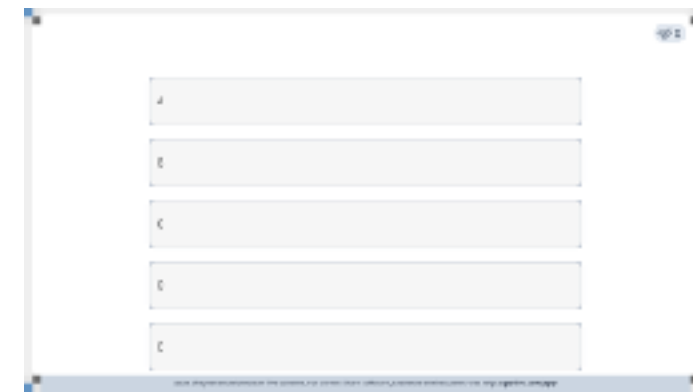


What about "linked list"


- Assume the current PC is already at instruction (1) and this linked list has only three nodes. This processor can fetch and issue 2 instructions per cycle, with exactly the same register renaming hardware and pipeline as we showed previously.

Which of the following C state of the code snippet determines the performance?

- A. do {
- B. number_of_nodes++;
- C. current = current->next;
- D. } while (current != NULL);




0%




A

0%




B

0%




C

0%



D

0%




E

What about "linked list"

Dynamic instructions

①	.L3:	movq	8(%rdi), %rdi
②		addl	\$1, %eax
③		testq	%rdi, %rdi
④		jne	.L3
⑤	.L3:	movq	8(%rdi), %rdi
⑥		addl	\$1, %eax
⑦		testq	%rdi, %rdi
⑧		jne	.L3
⑨	.L3:	movq	8(%rdi), %rdi
⑩		addl	\$1, %eax
⑪		testq	%rdi, %rdi
⑫		jne	.L3
⑬	.L3:	movq	8(%rdi), %rdi
⑭		addl	\$1, %eax
⑮		testq	%rdi, %rdi
⑯		jne	.L3

103

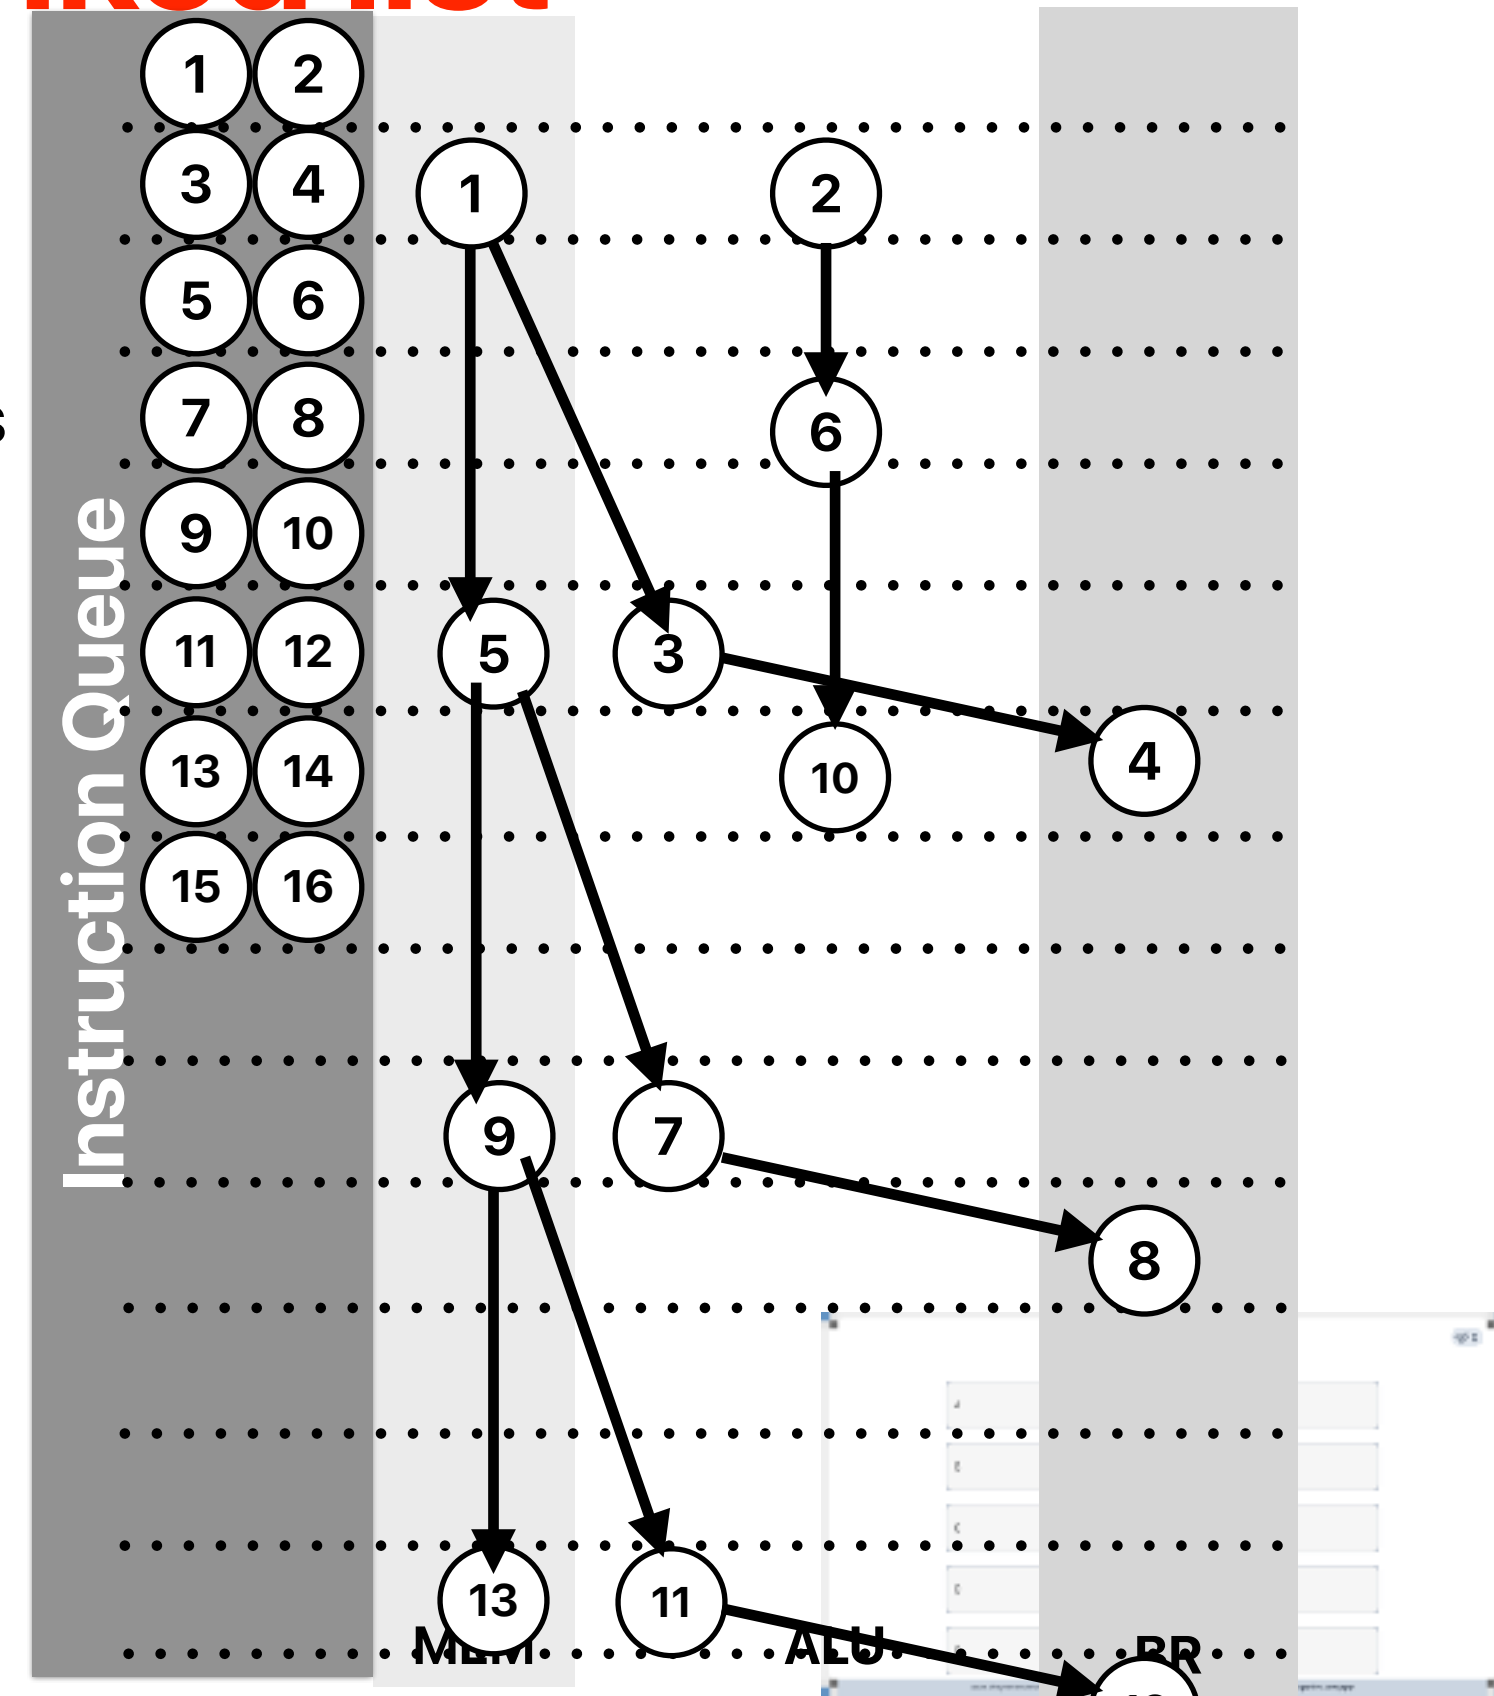


What about "linked list"


- For the following C code and it's translation in x86, **what's average CPI?** Assume the current PC is already at instruction (1) and this linked list has thousands of nodes. This processor can fetch and issue **2** instructions per cycle, with exactly the same register renaming hardware and pipeline as we showed previously.

```
do {  
    number_of_nodes++;  
    current = current->next;  
} while ( current != NULL )
```

- A. 0.5
- B. 0.8
- C. 1.0
- D. 1.2
- E. 1.5




0%




A

0%




B

0%




C

0%



D

0%




E

What about "linked list"


- For the following C code and it's translation in x86, **what's average CPI?** Assume the current PC is already at instruction (1) and this linked list has thousands of nodes. This processor can fetch and issue **2** instructions per cycle, with exactly the same register renaming hardware and pipeline as we showed previously.

```
do {  
    number_of_nodes++;  
    current = current->next;  
} while ( current != NULL )
```

- A. 0.5
- B. 0.8
- C. 1.0
- D. 1.2
- E. 1.5




0%




A

0%




B

0%




C

0%



D

0%



E

What about "linked list"

Performance determined by the critical path


4 cycles each iteration

4 instructions per iteration

$$CPI = \frac{4}{4} = 1$$

```
do {  
    number_of_nodes++;  
    current = current->next;  
} while ( current != NULL );
```

①	.L3:	movq	8(%rdi), %rdi
②		addl	\$1, %eax
③		testq	%rdi, %rdi
④		jne	.L3



The pipelines of Modern Processors

Intel Skylake

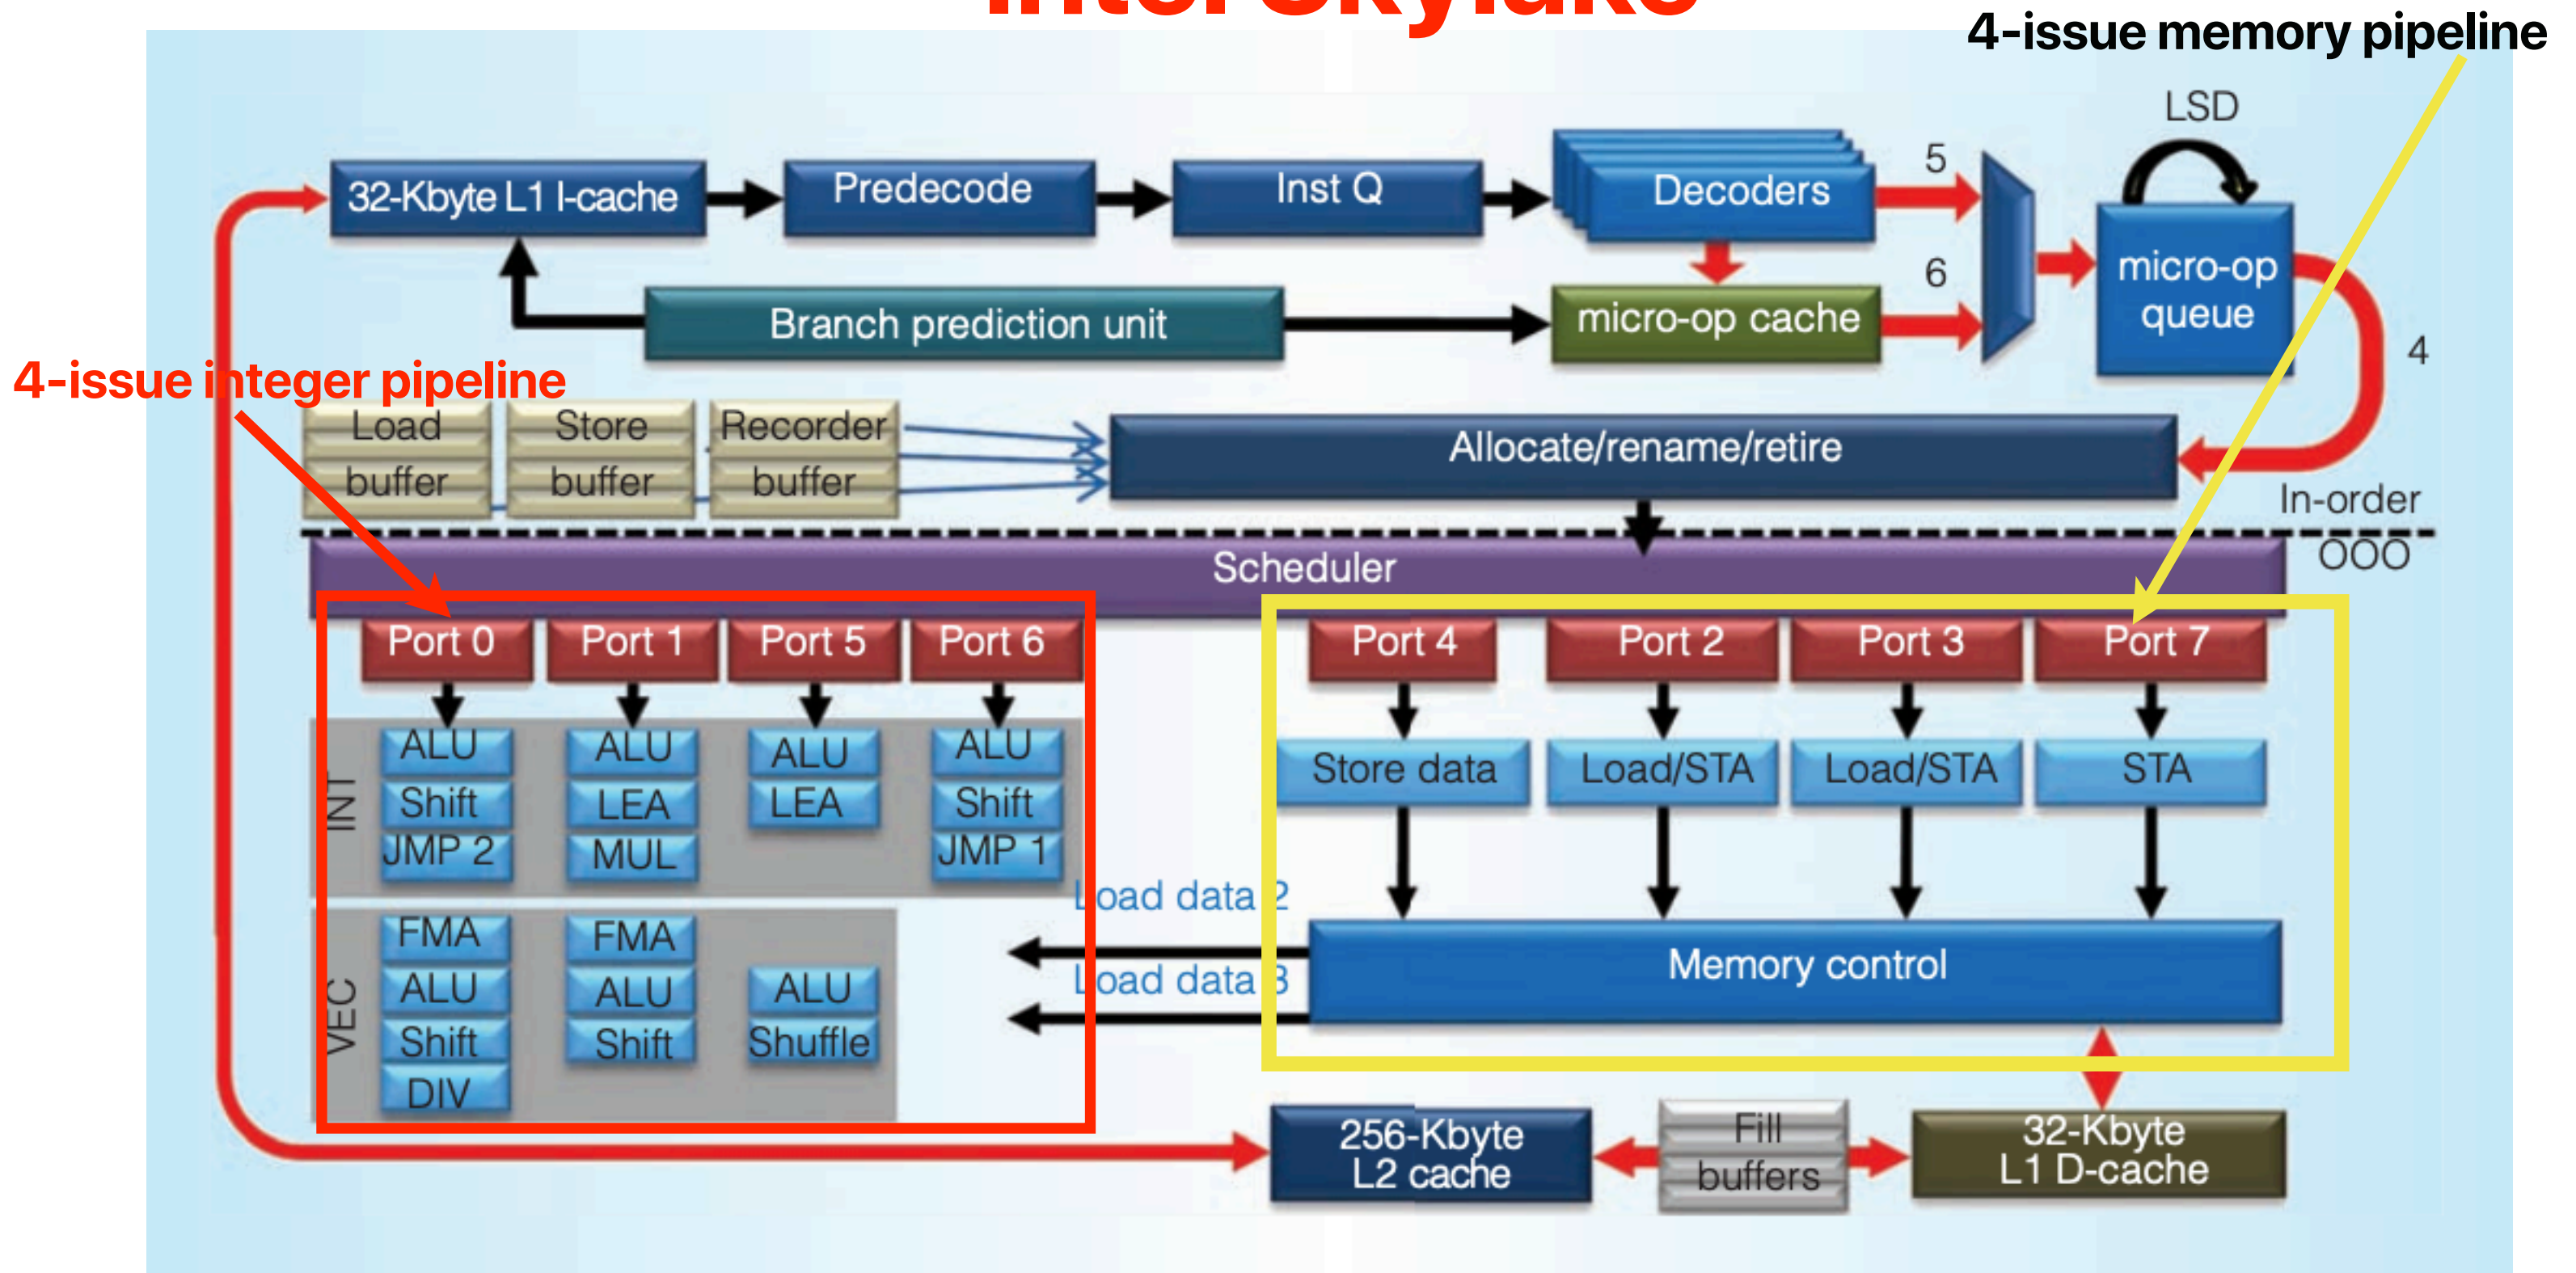


Figure 4. Skylake core block diagram.

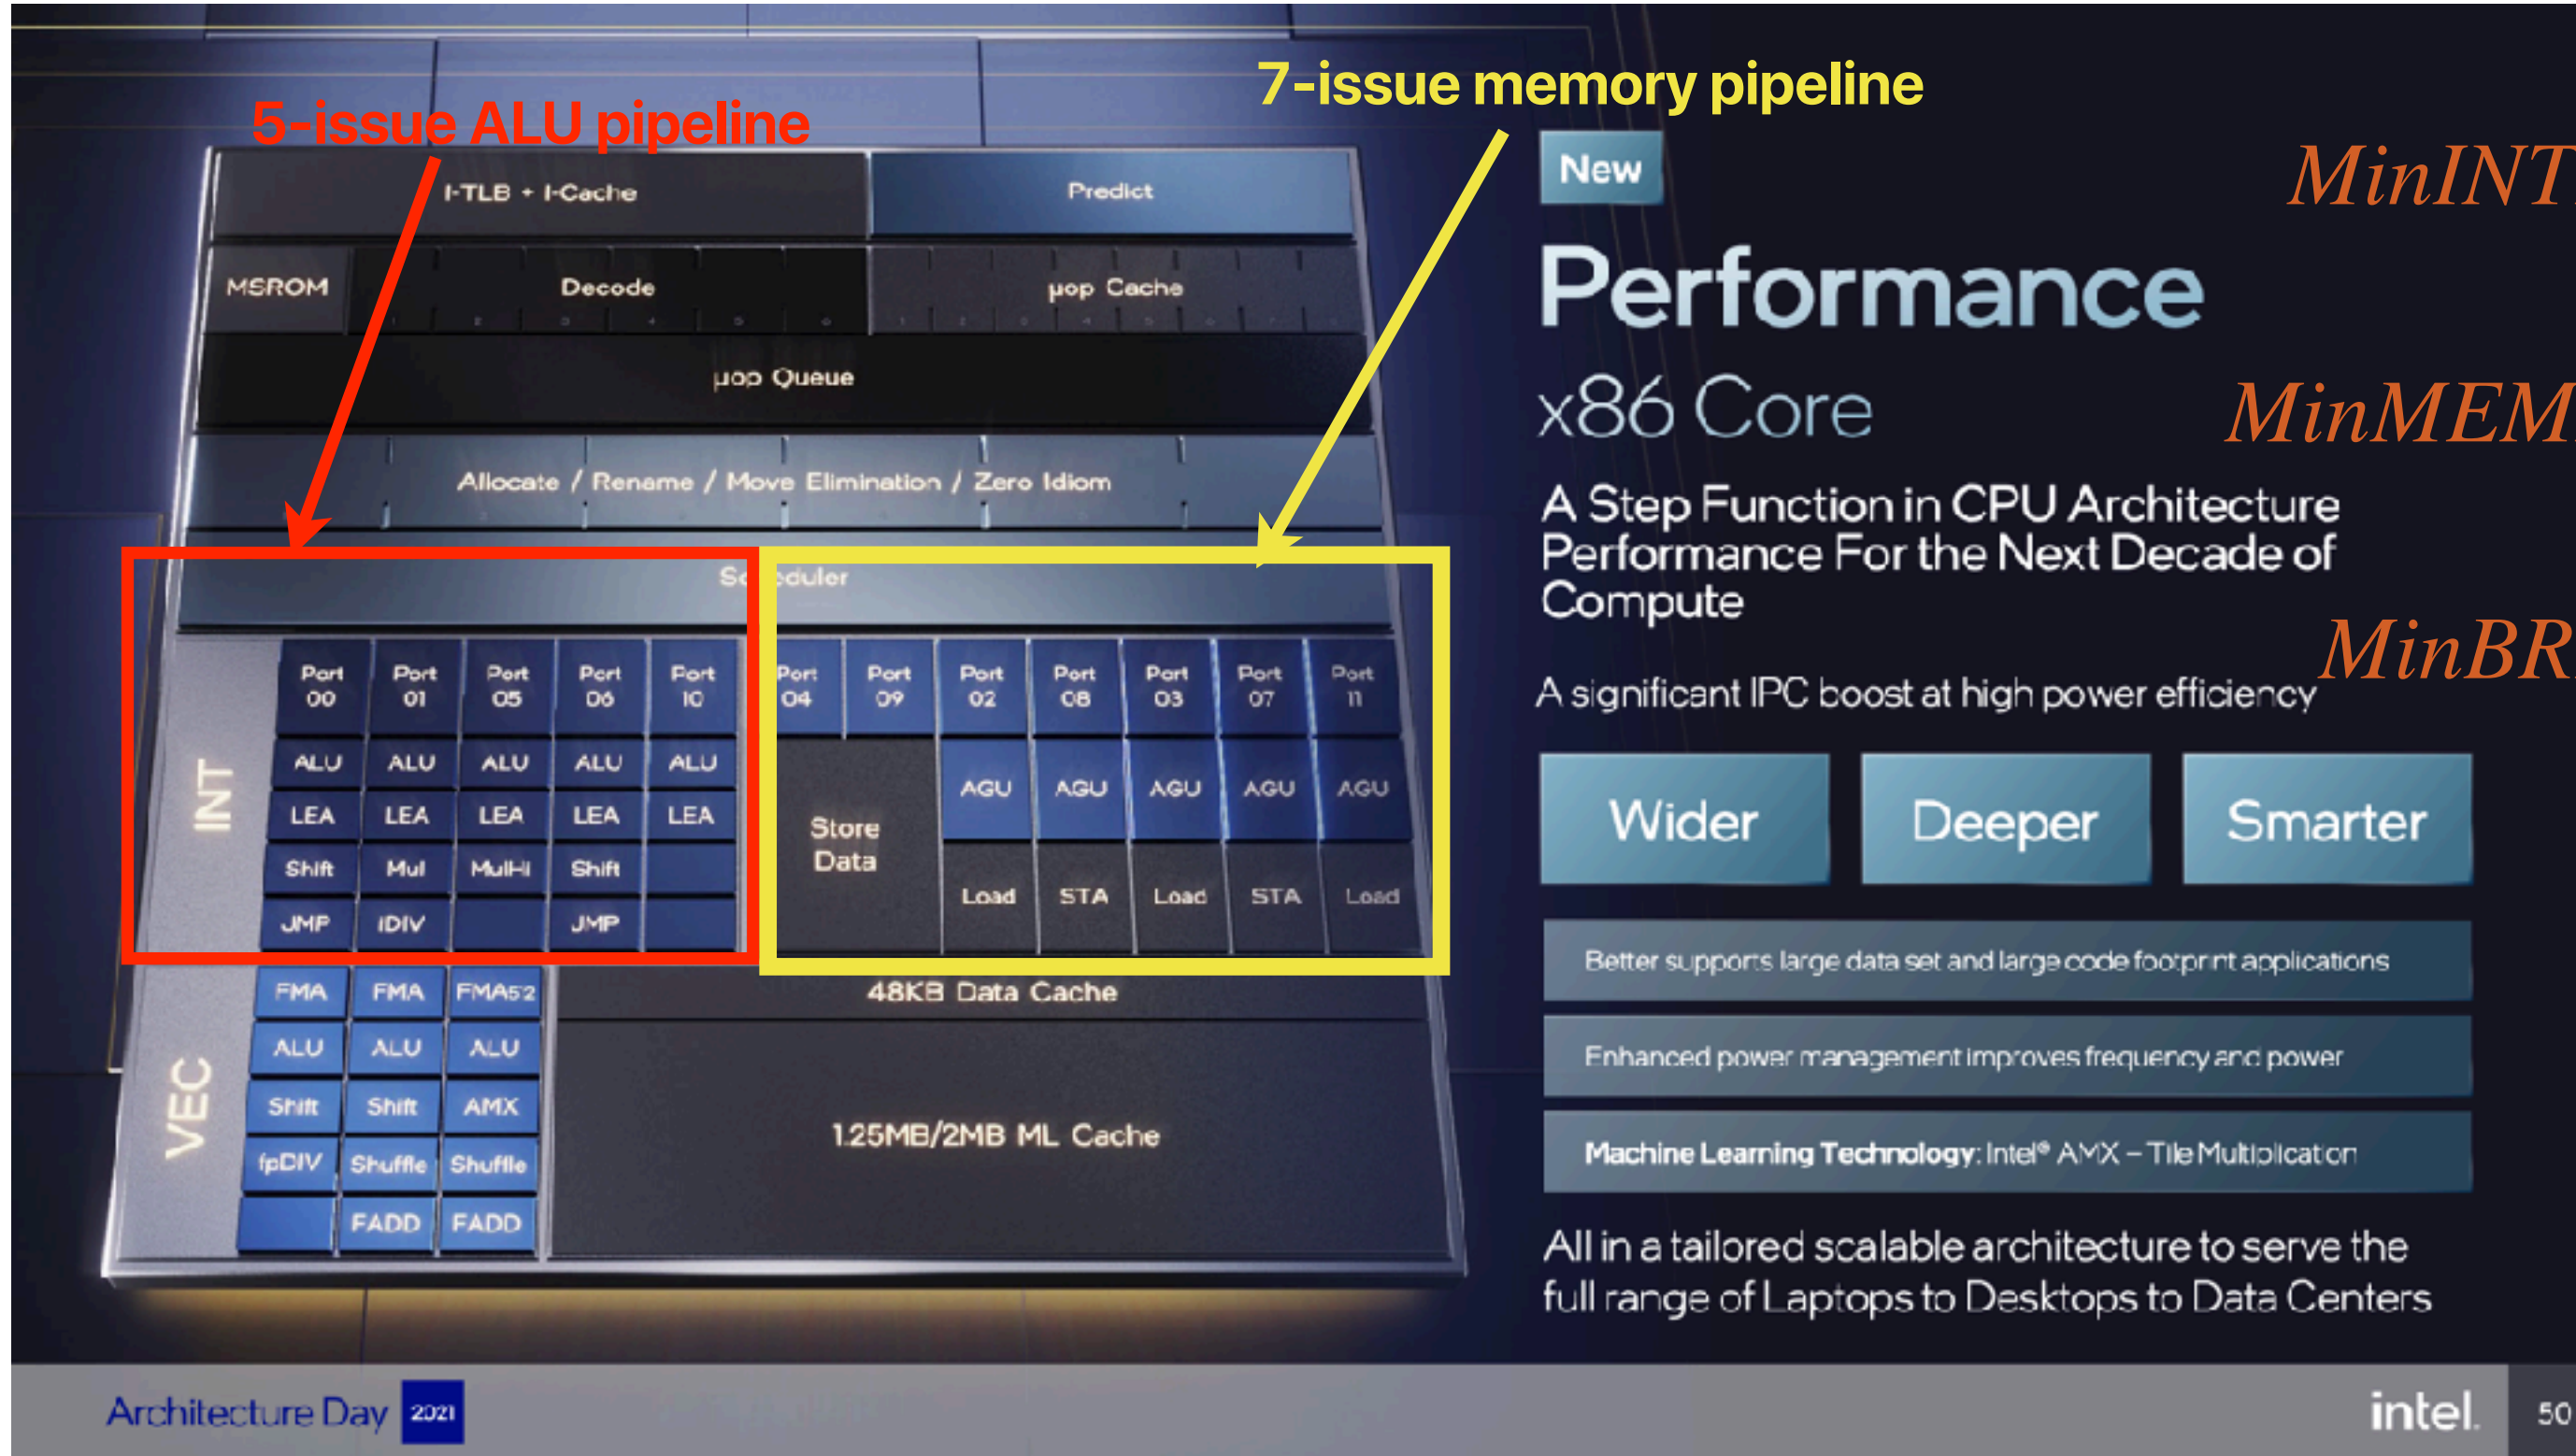
Intel Alder Lake

$$MinCPI = \frac{1}{12}$$

$$MinINTInst . CPI = \frac{1}{5}$$

$$MinMEMInst . CPI = \frac{1}{7}$$

$$MinBRInst . CPI = \frac{1}{2}$$



The image displays the Intel Alder Lake architecture diagram and a performance slide. The architecture diagram on the left shows the internal components of the processor, including the I-TLB + I-Cache, MSROM, Decode, μop Cache, μop Queue, Allocate / Rename / Move Elimination / Zero Idiom, Scheduler, and various execution units like ALU, LEA, Shift, MUL, MULHI, JMP, IDIV, FMA, FMA52, AMX, fpDIV, Shuffle, FADD, and AGU. A red box highlights the 5-issue ALU pipeline, and a yellow box highlights the 7-issue memory pipeline. The performance slide on the right, titled 'New Performance x86 Core', describes the processor as 'A Step Function in CPU Architecture Performance For the Next Decade of Compute' and 'A significant IPC boost at high power efficiency'. It lists three key features: 'Wider', 'Deeper', and 'Smarter', with corresponding benefits: 'Better supports large data set and large code footprint applications', 'Enhanced power management improves frequency and power', and 'Machine Learning Technology: Intel® AMX – Tile Multiplication'. The slide concludes with 'All in a tailored scalable architecture to serve the full range of Laptops to Desktops to Data Centers'. The Intel logo and 'Architecture Day 2021' are visible at the bottom left, and the Intel logo and '50' are at the bottom right.

5-issue ALU pipeline

7-issue memory pipeline

New Performance x86 Core

A Step Function in CPU Architecture Performance For the Next Decade of Compute

A significant IPC boost at high power efficiency

Wider Deeper Smarter

Better supports large data set and large code footprint applications

Enhanced power management improves frequency and power


Machine Learning Technology: Intel® AMX – Tile Multiplication

All in a tailored scalable architecture to serve the full range of Laptops to Desktops to Data Centers

Architecture Day 2021 intel 50

Project the performance of this code on Alder Lake

```
① .L10:
   cmpq  $0, 8(%rdi)
②   je  .L9
③   movslq (%rdi), %rdx
④   addq  %rdx, %rax
⑤   .L9:
   addq  $16, %rdi
⑥   cmpq  %rdi, %rcx
⑦   jne .L10
⑧   .L10:
   cmpq  $0, 8(%rdi)
⑨   je  .L9
⑩   movslq (%rdi), %rdx
⑪   addq  %rdx, %rax
⑫   .L9:
   addq  $16, %rdi
⑬   cmpq  %rdi, %rcx
⑭   jne .L10
⑮   .L10:
   cmpq  $0, 8(%rdi)
⑯   je  .L9
⑰   movslq (%rdi), %rdx
⑱   addq  %rdx, %rax
⑲   .L9:
   addq  $16, %rdi
⑳   cmpq  %rdi, %rcx
㉑   jne .L10
```



Summary: Characteristics of modern processor architectures

- Multiple-issue pipelines with multiple functional units available
 - Multiple ALUs
 - Multiple Load/store units
 - Dynamic OoO scheduling to reorder instructions whenever possible
- Cache — very high hit rate if your code has good locality
 - Very matured data/instruction prefetcher
- Branch predictors — very high accuracy if your code is predictable
 - Perceptron
 - Variable history predictors



Announcements

- Tips for answering examine question (or more importantly, technical interviews)
 - Be precise — pulling everything you know simply shows you don't really understanding what the question is asking
 - Important things go first — interviewers and graders lose patience very quick, you need to give a summary or conclusion before going into detail
 - We cannot read your mind. Explaining your thoughts after the examine (e.g., regrading requests) or interview (e.g., through follow-up) **does not** help win anything back
 - Show your work — simply tell the interviewer we can use DP to solve this problem won't give you positive review results
- **Last reading quiz** due this Wednesday before the lecture

Computer Science & Engineering

142

つづく

