

# **Virtual Memory**

Hung-Wei Tseng

# Recap: Let's dig into this code

```
int main(int argc, char *argv[])
{
    int i,j;
    double **a;
    double sum=0, average;
    int dim=32768;
    if(argc < 2)
    {
        fprintf(stderr, "Usage: %s dimension\n", argv[0]);
        exit(1);
    }
    dim = atoi(argv[1]);
    a = (double **)malloc(sizeof(double *)*dim);
    for(i = 0 ; i < dim; i++)
        a[i] = (double *)malloc(sizeof(double)*dim);
    for(i = 0 ; i < dim; i++)
        for(j = 0 ; j < dim; j++)
            a[i][j] = rand();
    for(i = 0 ; i < dim; i++)
        for(j = 0 ; j < dim; j++)
            sum+=a[i][j];
    average = sum/(dim*dim);
    fprintf(stderr,"average: %lf\n",average);
    for(i = 0 ; i < dim; i++)
        free(a[i]);
    free(a);
    return 0;
}
```

```
終端機 — -tosh — 102x25
top ... -tosh + E

#include <stdlib.h>
#include <cassert.h>
#include <sched.h>
#include <sys/syscall.h>
#include <time.h>

//#define dim 32768
//#define dim 49152
int main(int argc, char *argv[])
{
    int i,j;
    double **a;
    double sum=0, average;
    int dim=32768;
    if(argc < 2)
    {
        fprintf(stderr, "Usage: %s dimension\n", argv[0]);
        exit(1);
    }
    dim = atoi(argv[1]);
    a = (double **)malloc(sizeof(double *)*dim);
File memory_allocation.c not changed so no update needed
BunnyMACProRetina [/Users/bunny/Dropbox/CSE203/GitHub/demo/virtual_memory] -bunny- ./memory_allocation

BunnyMACProRetina [/Users/bunny/Dropbox/CSE203/GitHub/demo/virtual_memory] -bunny- □
```

# Recap: Let's dig into this code

```
#define _GNU_SOURCE
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>
#include <assert.h>
#include <sched.h>
#include <sys/syscall.h>
#include <time.h>

double a;

int main(int argc, char *argv[])
{
    int i, number_of_total_processes=4;
    number_of_total_processes = atoi(argv[1]);
    // Create processes
    for(i = 0; i< number_of_total_processes-1 && fork(); i++);
    // Generate rand seed
    srand((int)time(NULL)+(int)getpid());
    a = rand();
    fprintf(stderr, "\nProcess %d. Value of a is %lf and address of a is %p\n", getpid(), a, &a);
    sleep(10);
    fprintf(stderr, "\nProcess %d. Value of a is %lf and address of a is %p\n", getpid(), a, &a);
    return 0;
}
```

# Demo revisited

```
sleep(10);
fprintf(stderr, "\nProcess %d: Value of a is %lf and address of a is %p\n", (int)getpid(), a, &a);
return 0;
}
```

File virtualization.c not changed so no update needed

```
BunnyMACProRetina [/Users/bunny/Dropbox/CSE203/GitHub/demo/virtual_memory] -bunny- make
gcc -O3 virtualization.c -o virtualization
```

```
BunnyMACProRetina [/Users/bunny/Dropbox/CSE203/GitHub/demo/virtual_memory] -bunny- ./virtualization 4
```

```
Process 19719: Value of a is 1671139616.000000 and address of a is 0x104967050
```

```
Process 19720: Value of a is 1671156423.000000 and address of a is 0x104967050
```

```
Process 19718: Value of a is 1671122809.000000 and address of a is 0x104967050
```

```
Process 19721: Value of a is 1671173230.000000 and address of a is 0x104967050
    Different values
```

```
Process 19719: Value of a is 1671139616.000000 and address of a is 0x104967050
```

```
Process 19721: Value of a is 1671173230.000000 and address of a is 0x104967050
```

```
Process 19720: Value of a is 1671156423.000000 and address of a is 0x104967050
```

```
Process 19718: Value of a is 1671122809.000000 and address of a is 0x104967050
```

```
BunnyMACProRetina [/Users/bunny/Dropbox/CSE203/GitHub/demo/virtual_memory] -bunny-
```

Different values are  
preserved

The same memory  
address!

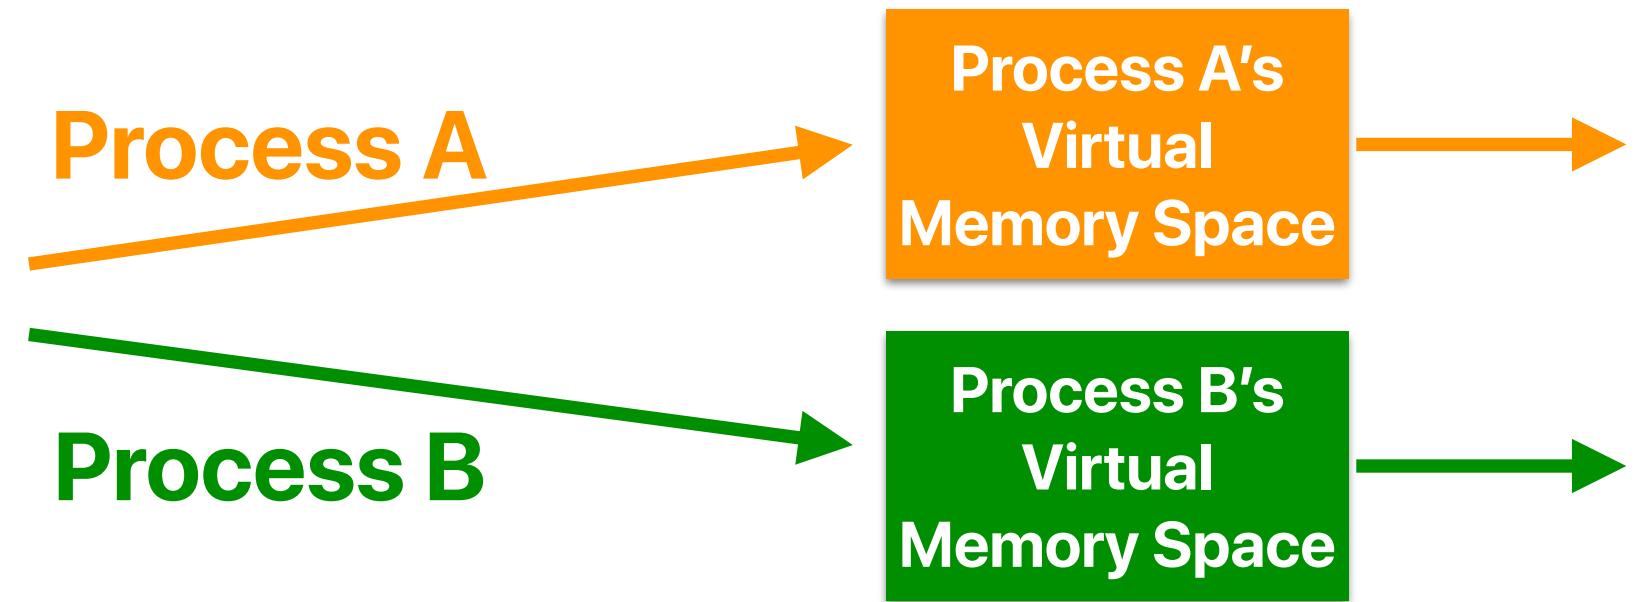
# Recap: Demo revisited

```
#define _GNU_SOURCE
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>
#include <assert.h>
#include <sched.h>
#include <sys/syscall.h>
#include <time.h>

double a;

int main(int argc, char *argv[])
{
    int i, number_of_total_processes=4;
    number_of_total_processes = atoi(argv[1]);
    for(i = 0; i< number_of_total_processes-1 && fork(); i++);
    srand((int)time(NULL)+(int)getpid());
    fprintf(stderr, "\nProcess %d. Value of a is %lf and address of a is %p\n",getpid(), a, &a);
    sleep(10);
    fprintf(stderr, "\nProcess %d. Value of a is %lf and address of a is %p\n",getpid(), a, &a);
    return 0;
}
```

**&a = 0x601090**

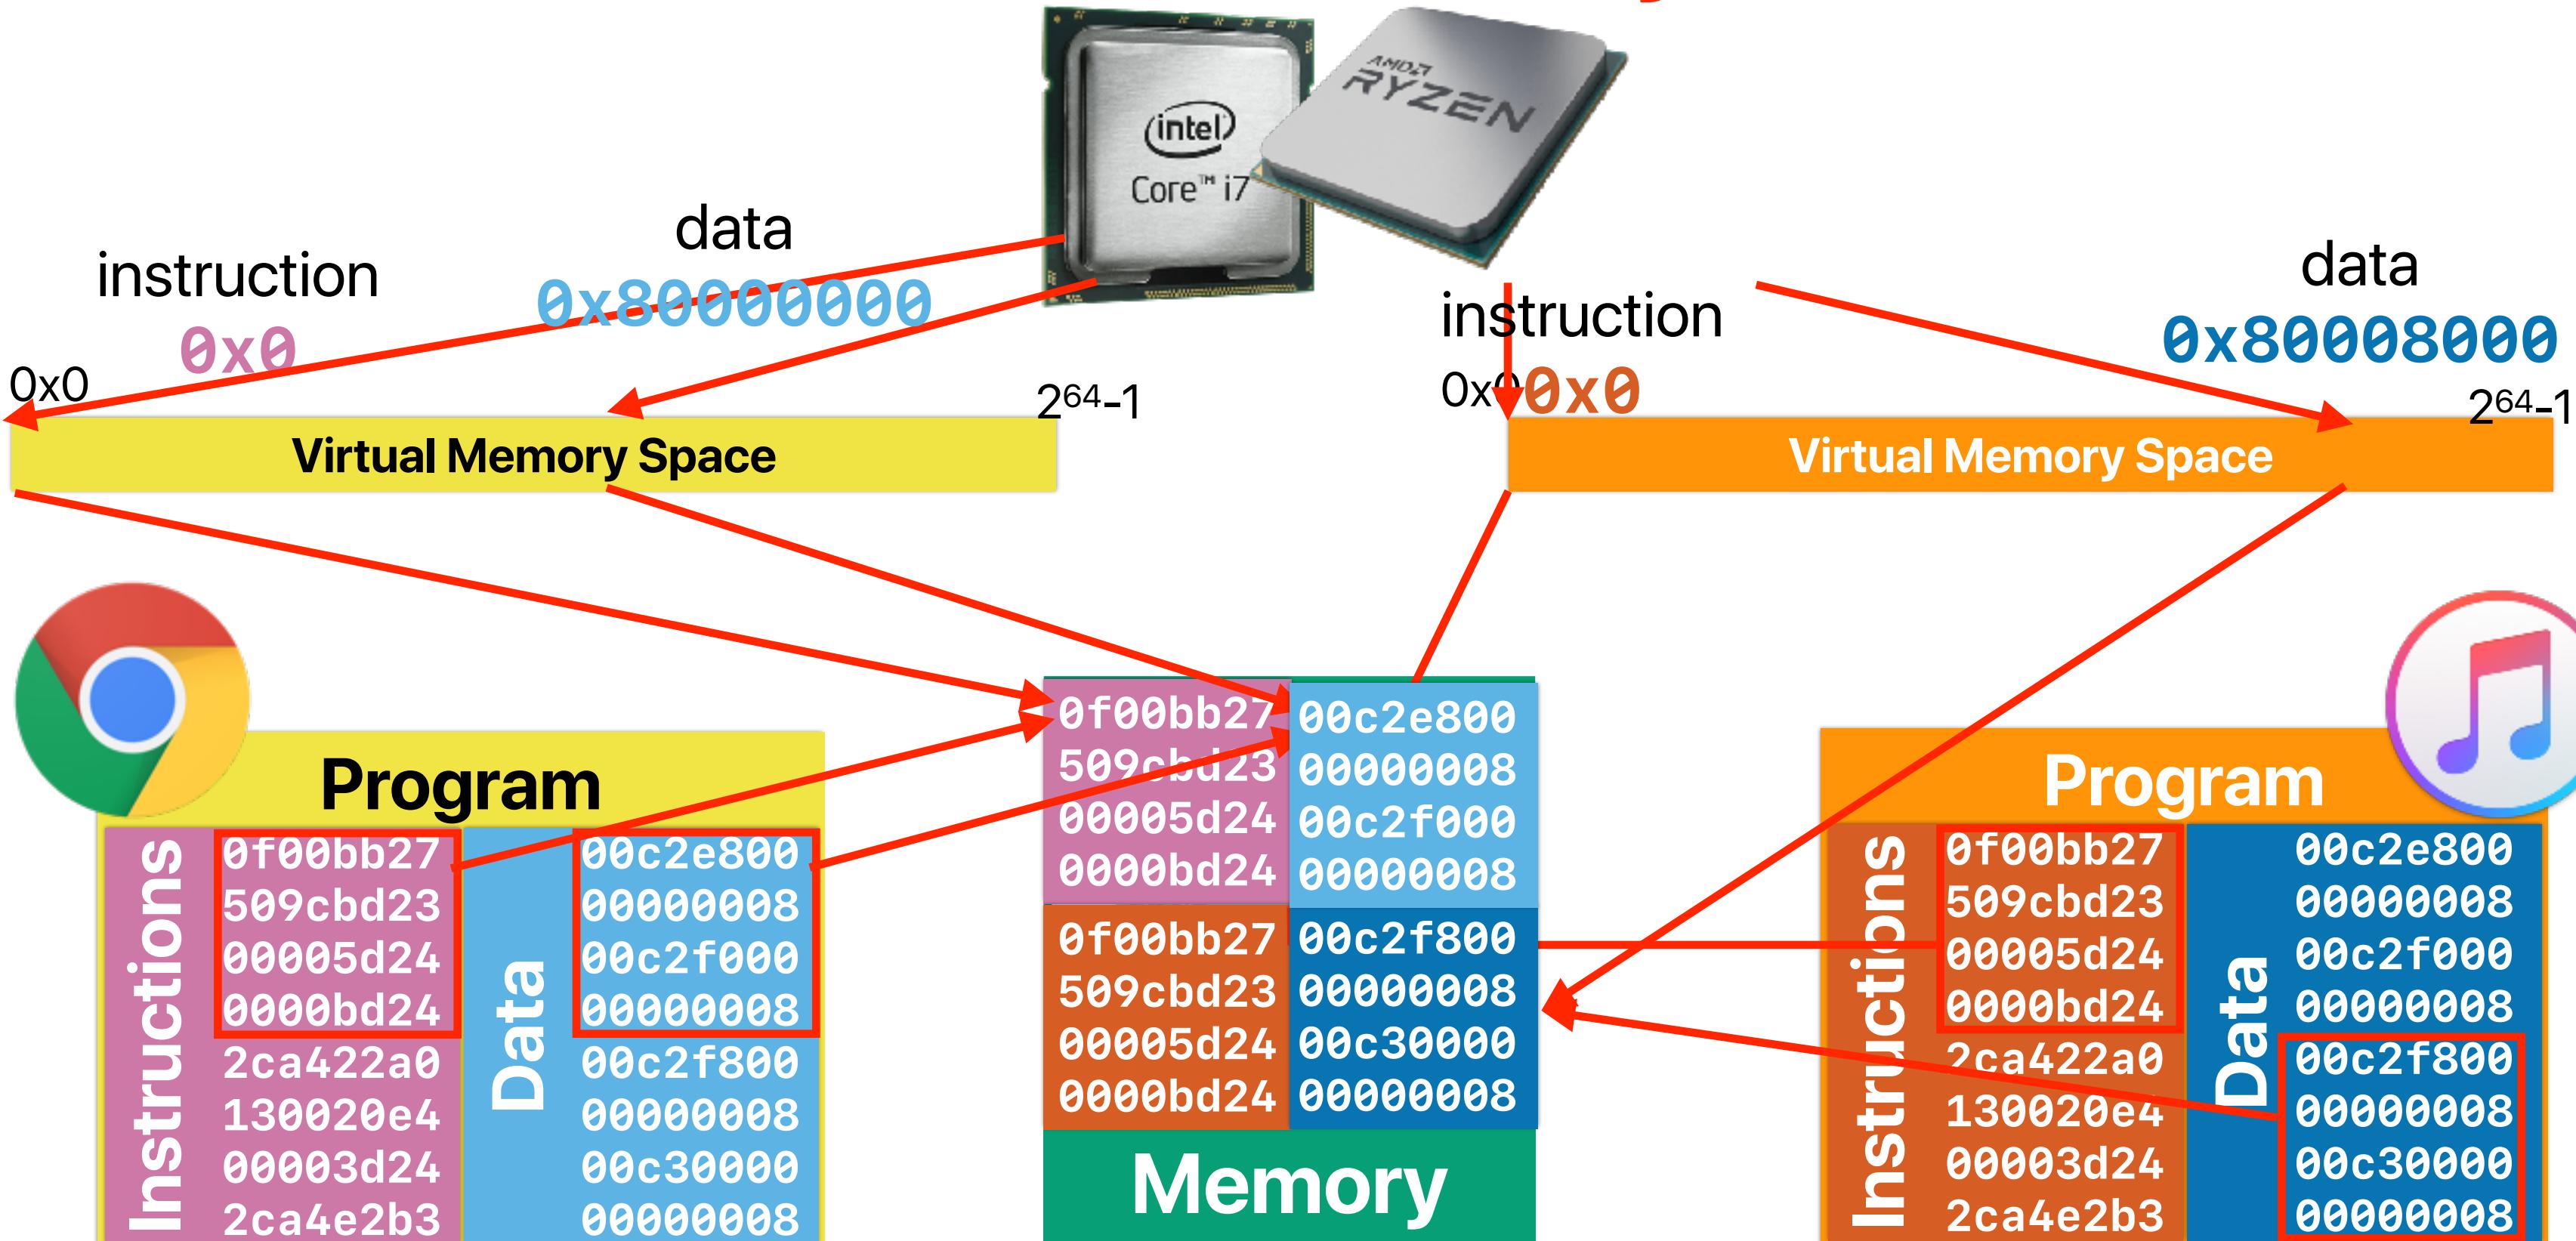


# Outline

- Virtual memory
- Architectural support for virtual memory

# **Virtual Memory**

# Virtual memory



# Virtual memory

- An **abstraction** of memory space available for programs/software/programmer
- Programs execute using virtual memory address
- The operating system and hardware work together to handle the mapping between virtual memory addresses and real/physical memory addresses
- Virtual memory organizes memory locations into “**pages**”

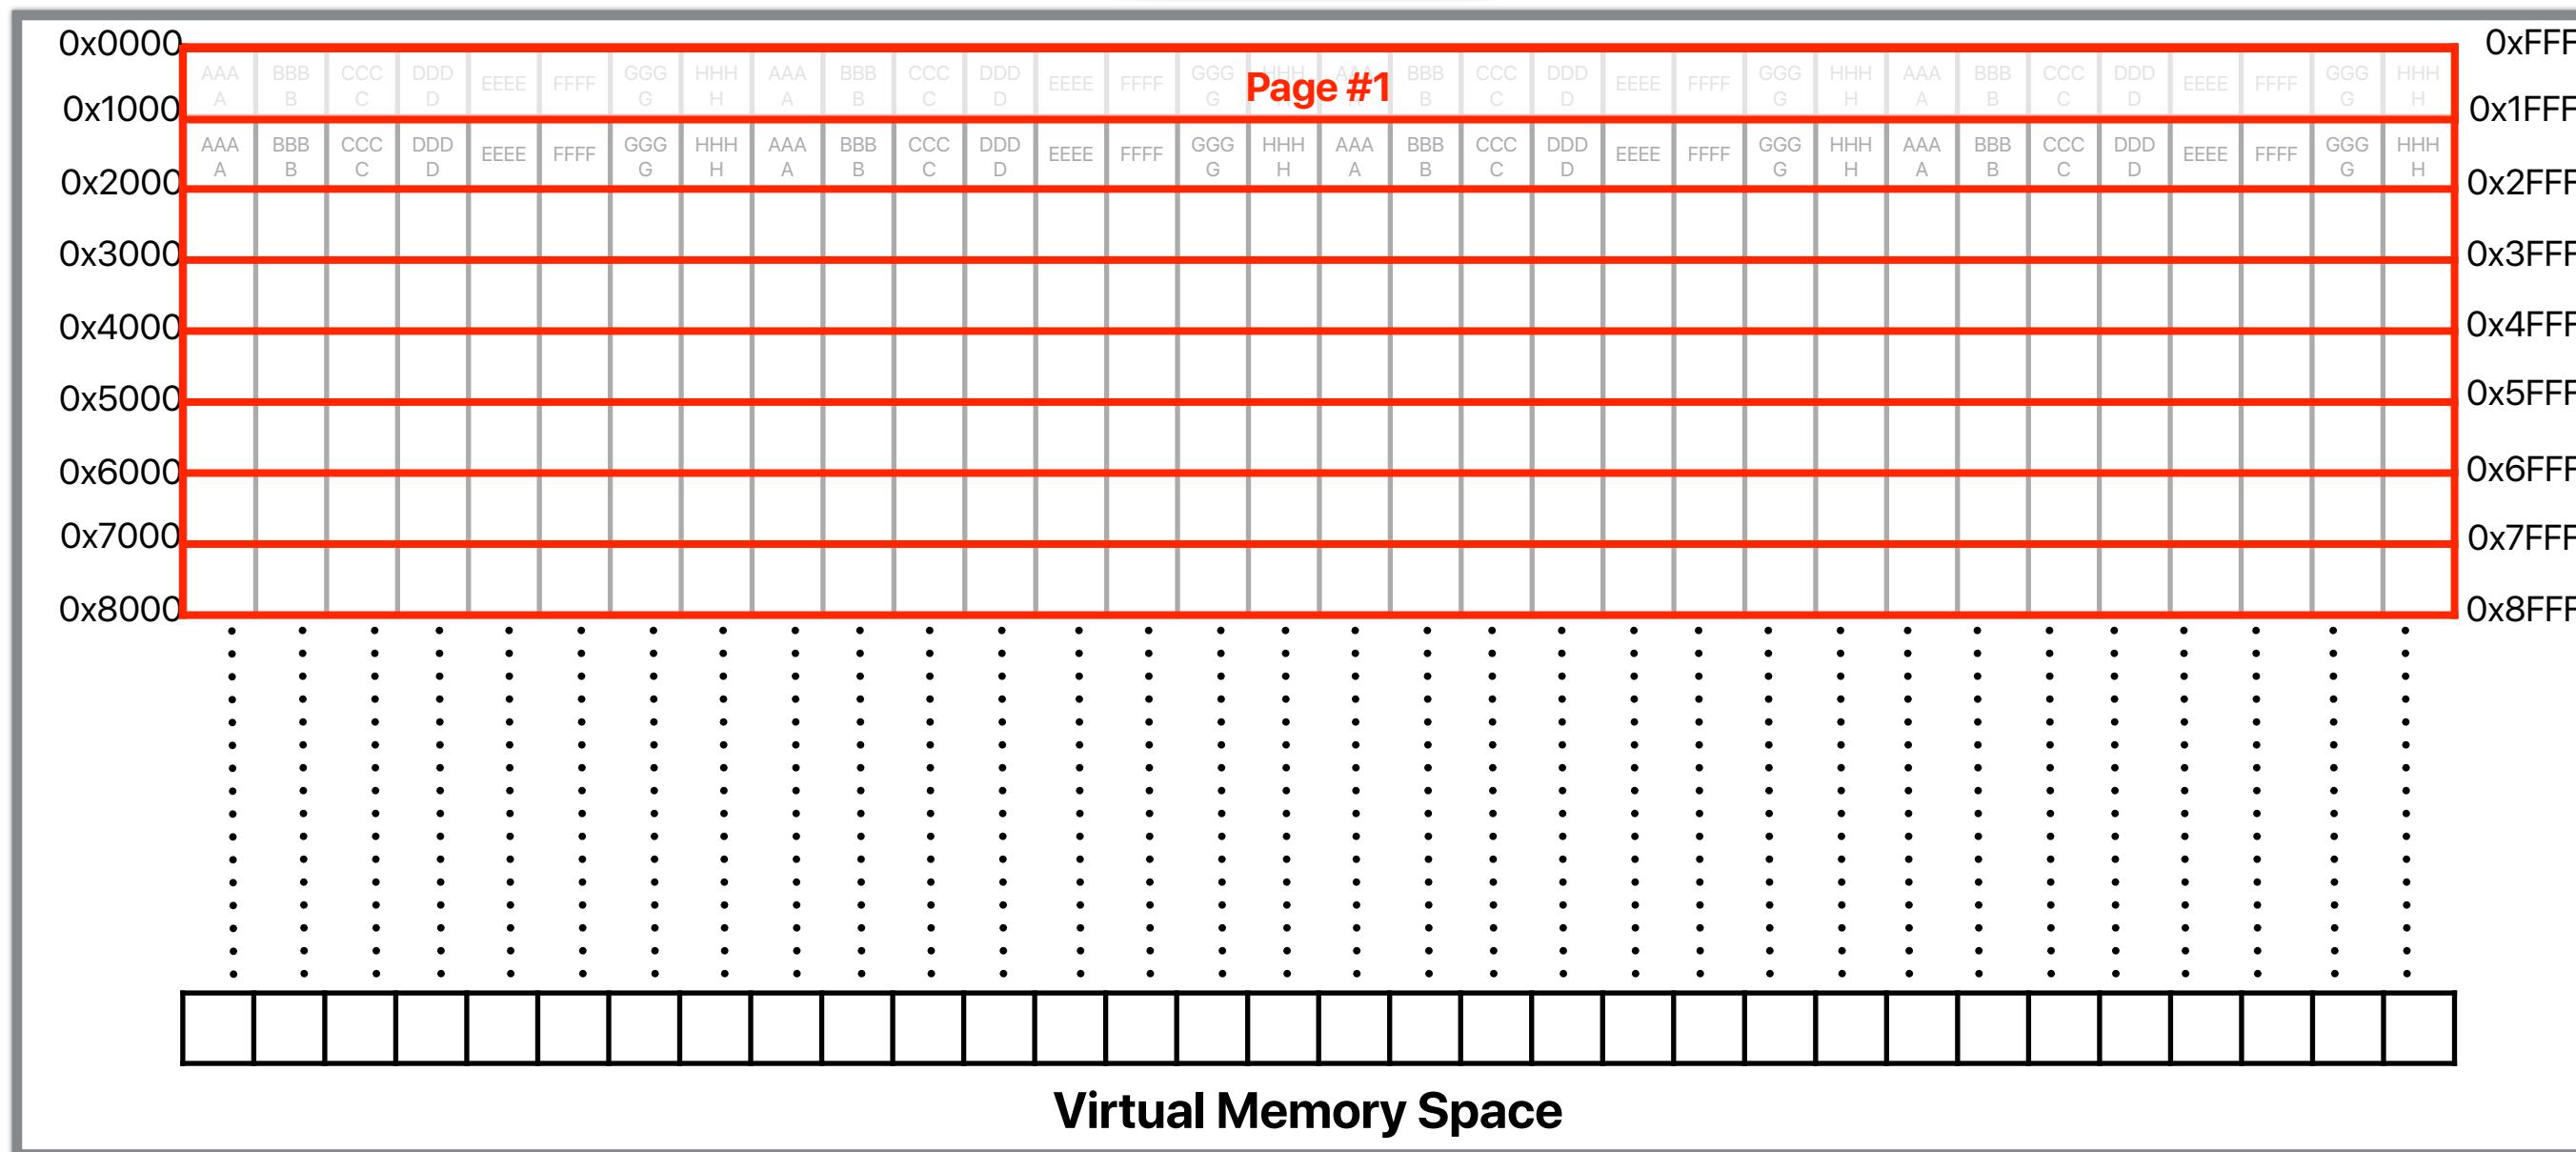
Processor  
Core  
Registers

# The virtual memory abstraction

load 0x0009

Page table

Main memory  
(DRAM)



Processor

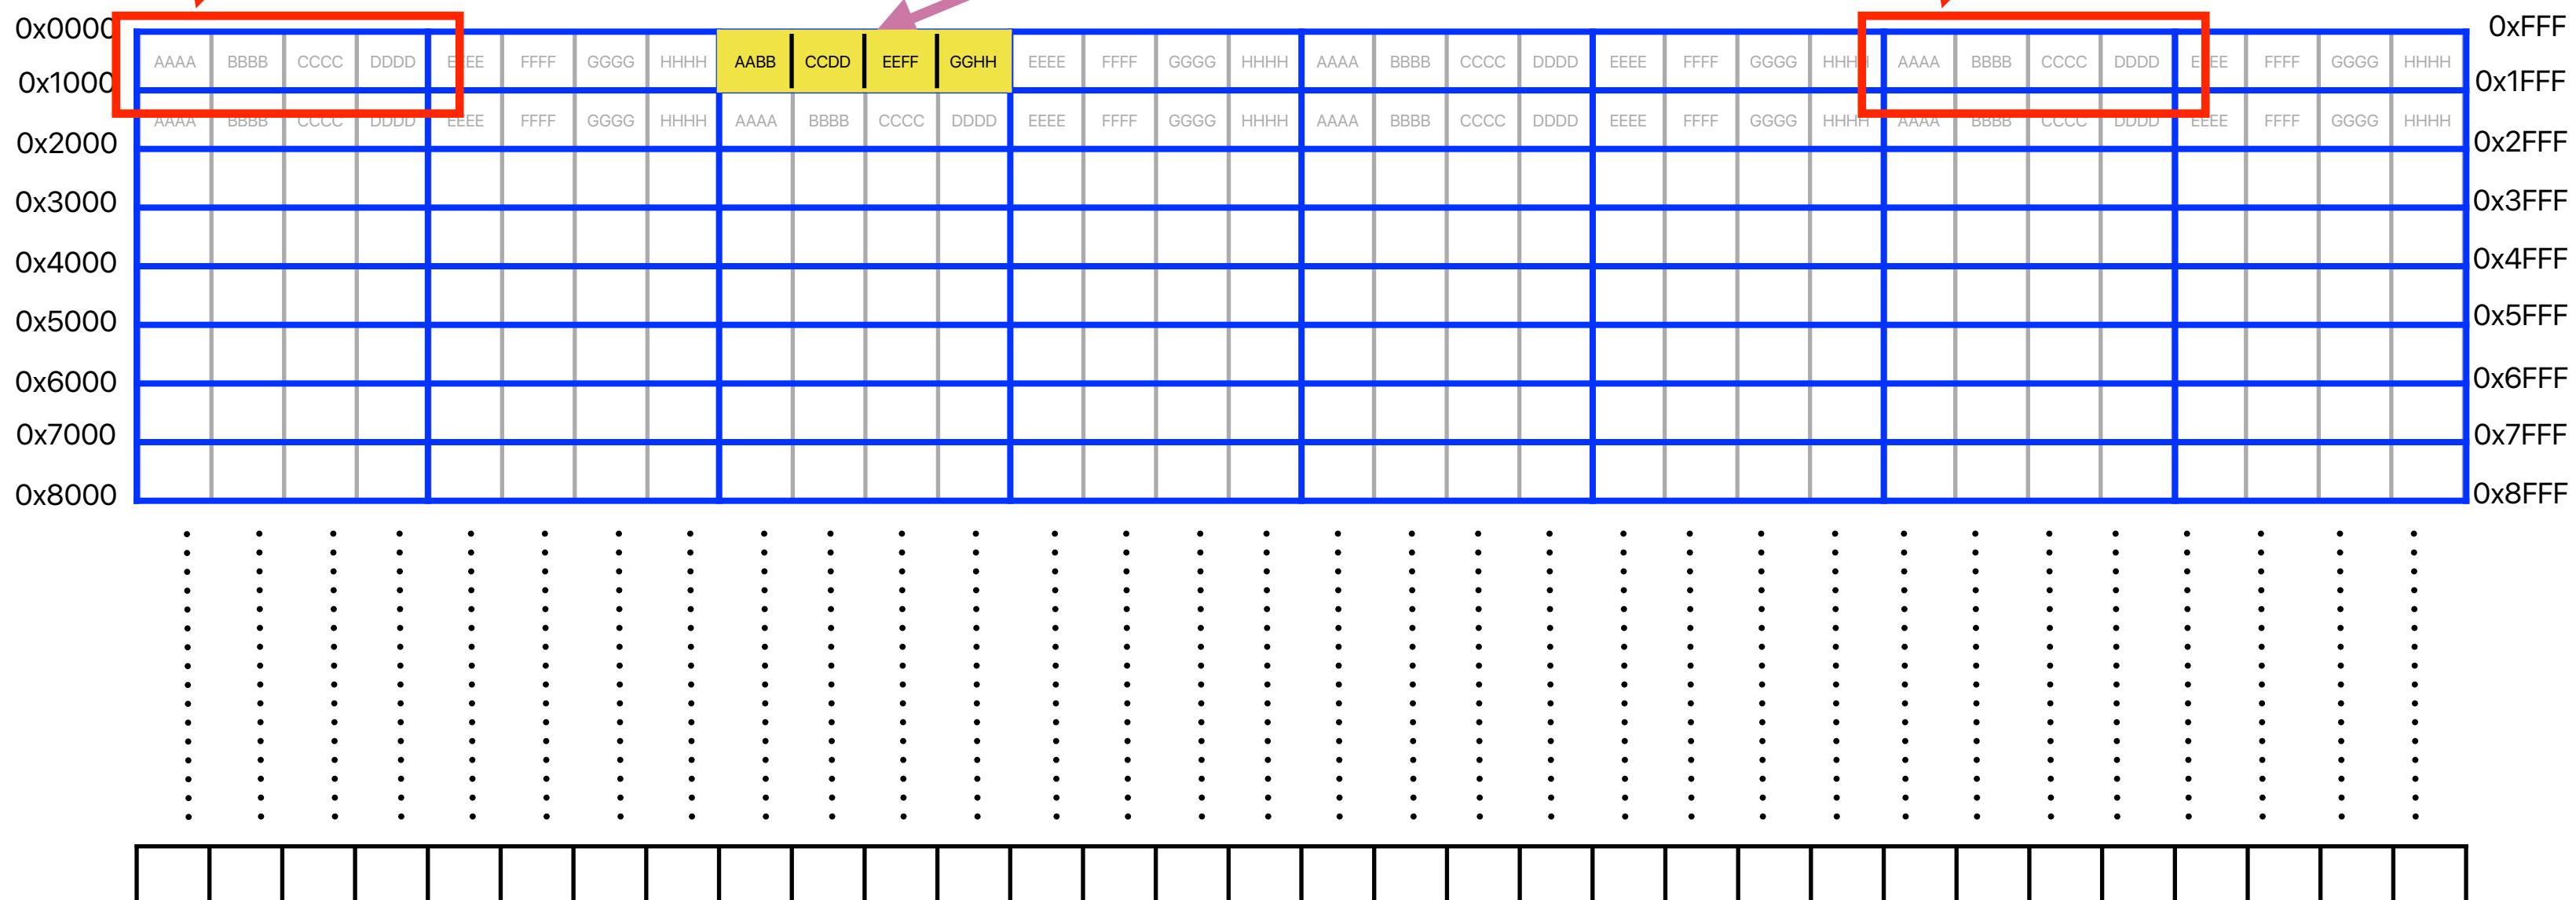
Core

Registers

Recap: To capture “spatial” locality, \$ fetch a “block”

1w 0x0024

Assume each block is 16 bytes



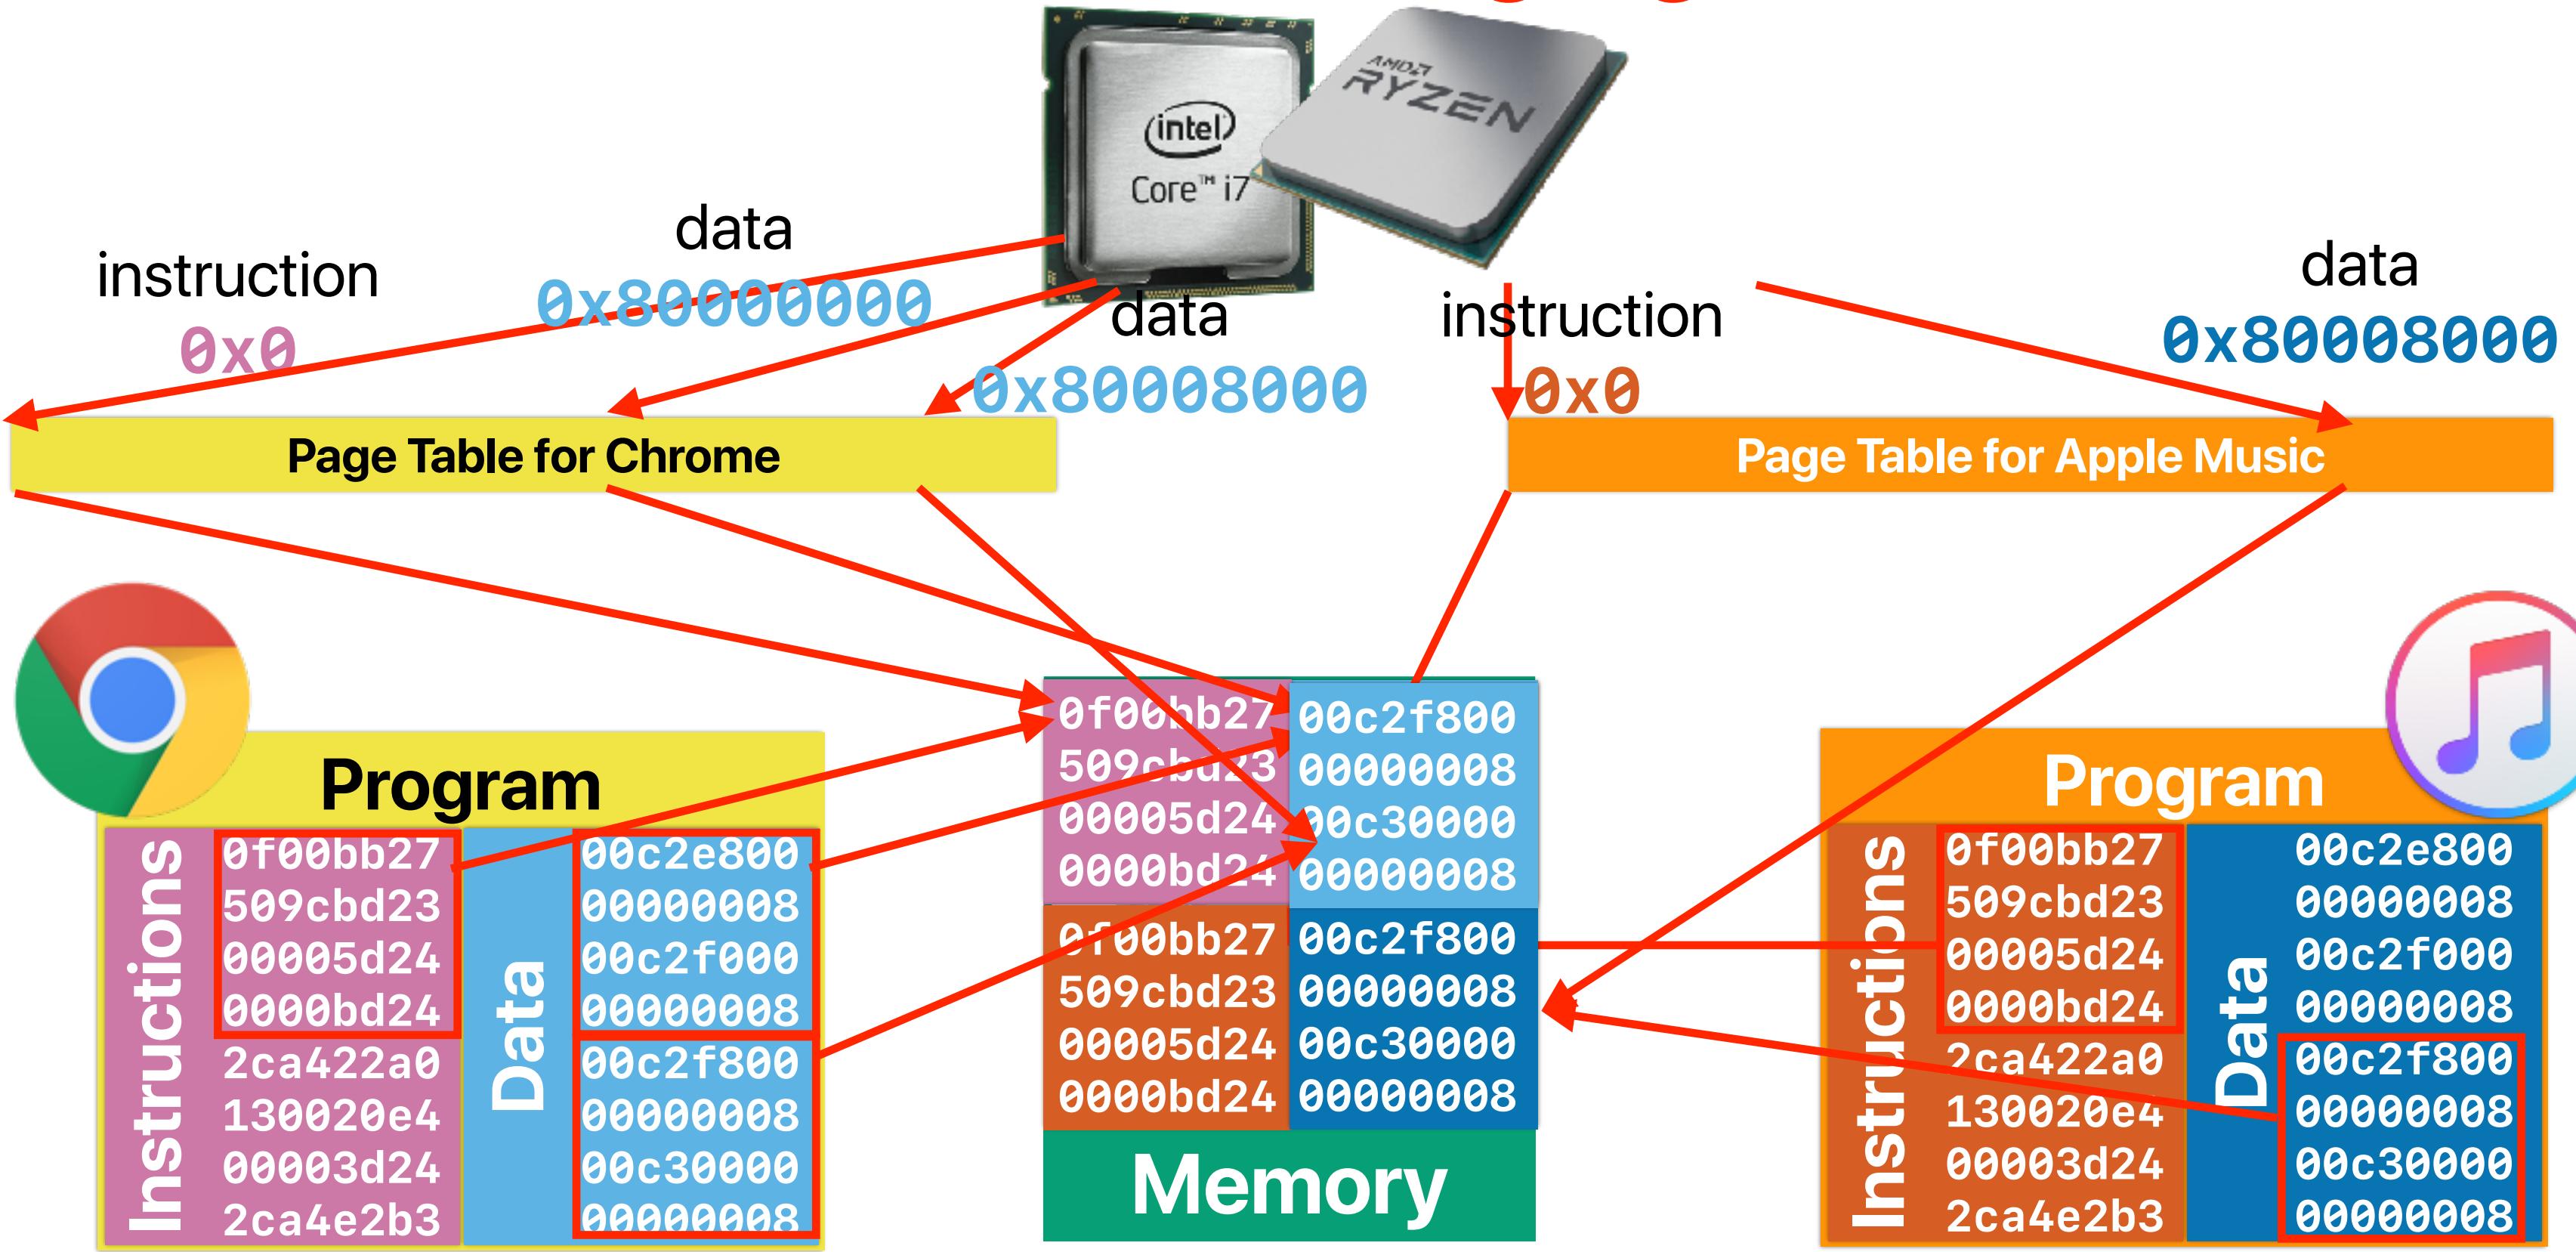
# Why Virtual memory?

- Allowing multiple applications to share physical main memory
  - Memory protection/isolation among programs/processes is automatically achieved
- Allowing applications to work even though the installed physical memory or available physical memory is smaller than the working set of the application
  - Programmer does not need to worry about the physical memory capacity of different machines — make compiled program compatible
  - Multiple programs can work concurrently even though their total memory demand is larger than the installed physical memory

# Demand paging

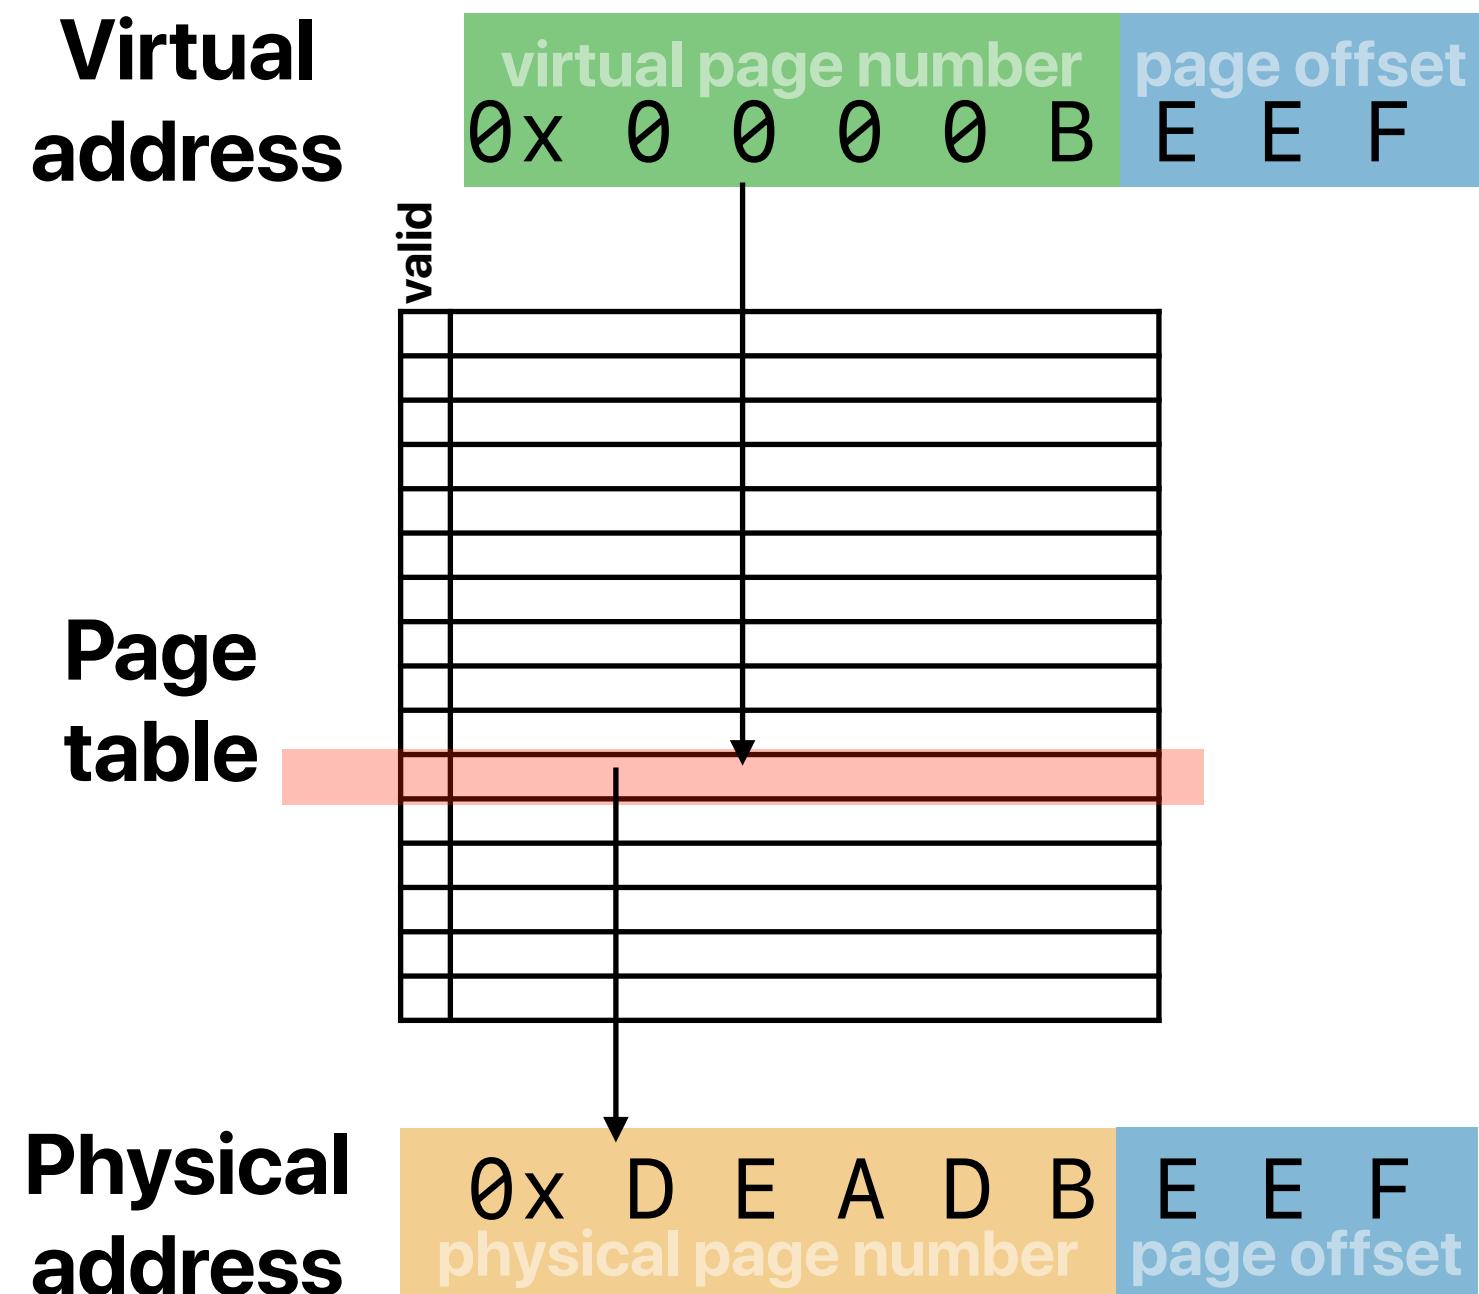
- Treating physical main memory as a “cache” of virtual memory
- The block size is the “page size”
- The page table is the “tag array”
- It’s a “fully-associate” cache — a virtual page can go anywhere in the physical main memory

# Demand paging



# Address translation

- Processor receives virtual addresses from the running code, main memory uses physical memory addresses
- Virtual address space is organized into “pages”
- The system references the **page table** to translate addresses
  - Each process has its own page table
  - The page table content is maintained by OS





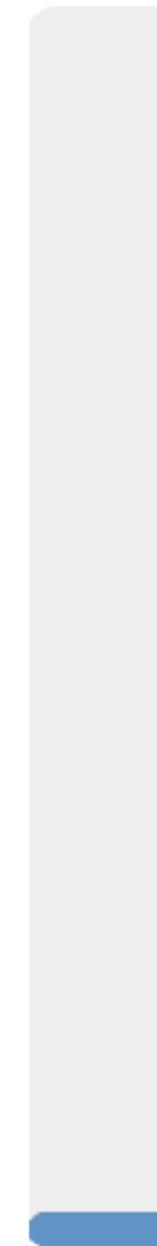
# Size of page table

- Assume that we have **64-bit** virtual address space, each page is 4KB, each page table entry is 8 Bytes, what magnitude in size is the page table for a process?
  - A. MB —  $2^{20}$  Bytes
  - B. GB —  $2^{30}$  Bytes
  - C. TB —  $2^{40}$  Bytes
  - D. PB —  $2^{50}$  Bytes
  - E. EB —  $2^{60}$  Bytes



 0

0



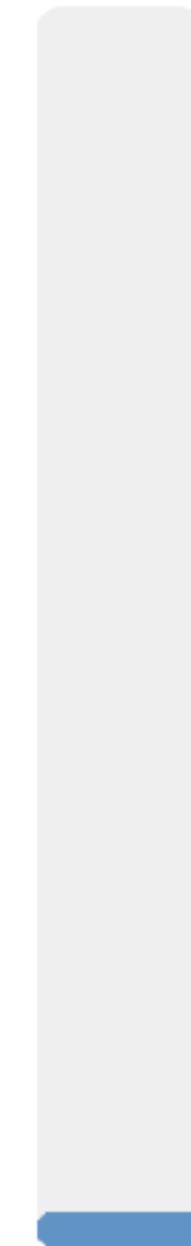
A

0



B

0



C

0



D

0



E



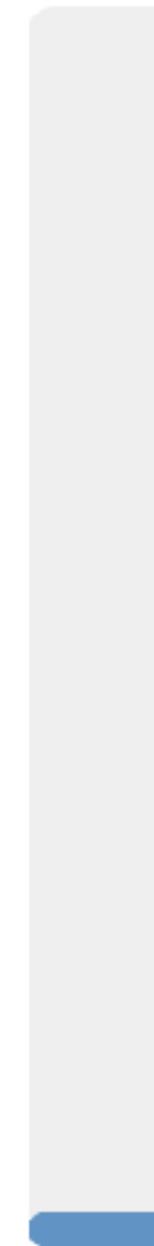
# Size of page table

- Assume that we have **64-bit** virtual address space, each page is 4KB, each page table entry is 8 Bytes, what magnitude in size is the page table for a process?
  - A. MB —  $2^{20}$  Bytes
  - B. GB —  $2^{30}$  Bytes
  - C. TB —  $2^{40}$  Bytes
  - D. PB —  $2^{50}$  Bytes
  - E. EB —  $2^{60}$  Bytes



 0

0



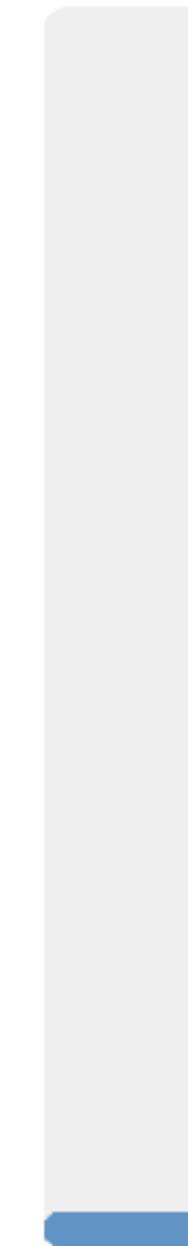
A

0



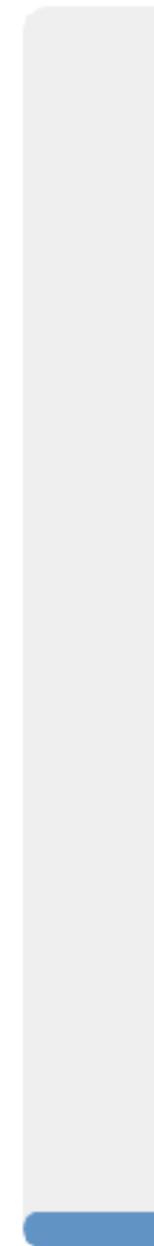
B

0



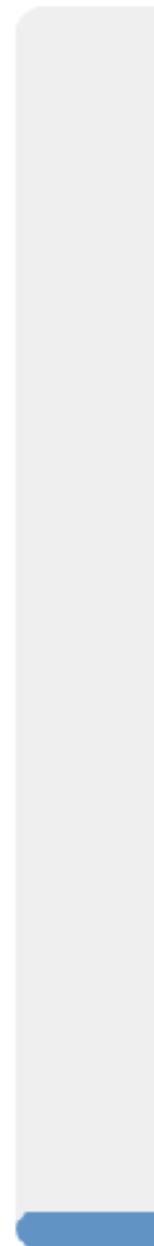
C

0



D

0



E

# Size of page table

- Assume that we have **64-bit** virtual address space, each page is 4KB, each page table entry is 8 Bytes, what magnitude in size is the page table for a process?
  - MB —  $2^{20}$  Bytes
  - GB —  $2^{30}$  Bytes
  - TB —  $2^{40}$  Bytes
  - PB —  $2^{50}$  Bytes
  - EB —  $2^{60}$  Bytes

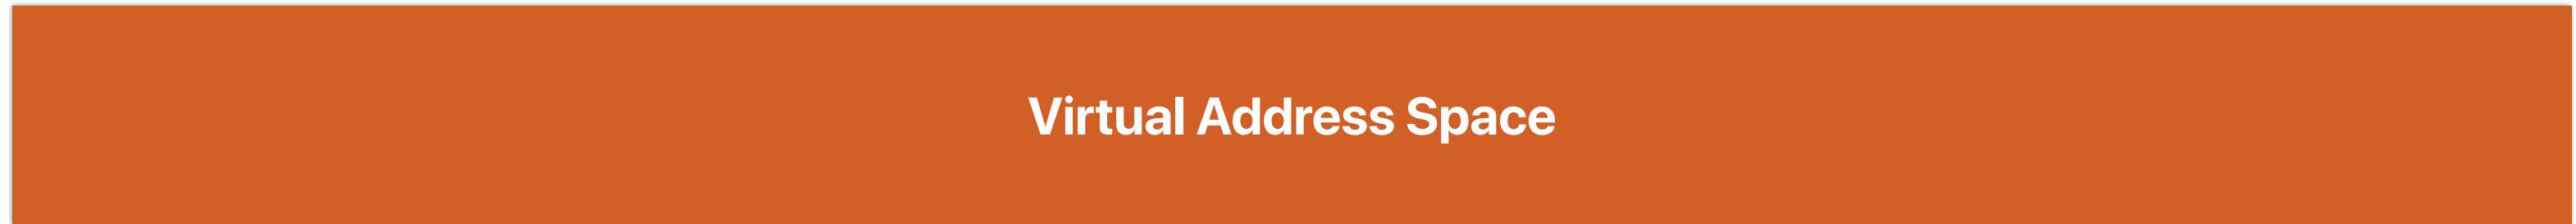
$$\frac{2^{64} \text{ Bytes}}{4 \text{ KB}} \times 8 \text{ Bytes} = 2^{55} \text{ Bytes} = 32 \text{ PB}$$

If you still don't know why — you need to take CS202

# Conventional page table

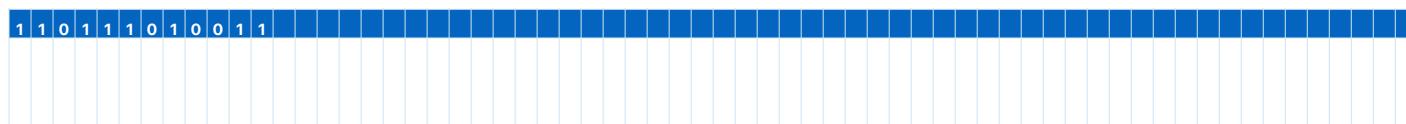
0x0

0xFFFFFFFFFFFFFFFFF

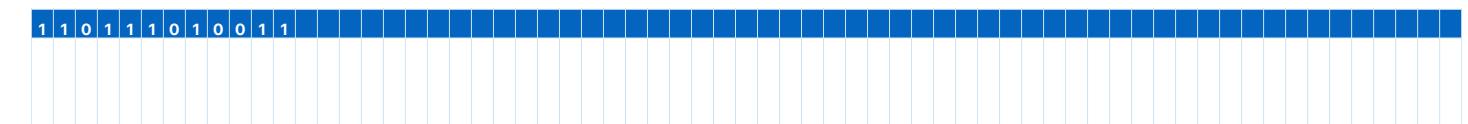


- must be consecutive in the physical memory
- need a big segment! — difficult to find a spot
- simply too big to fit in memory if address space is large!

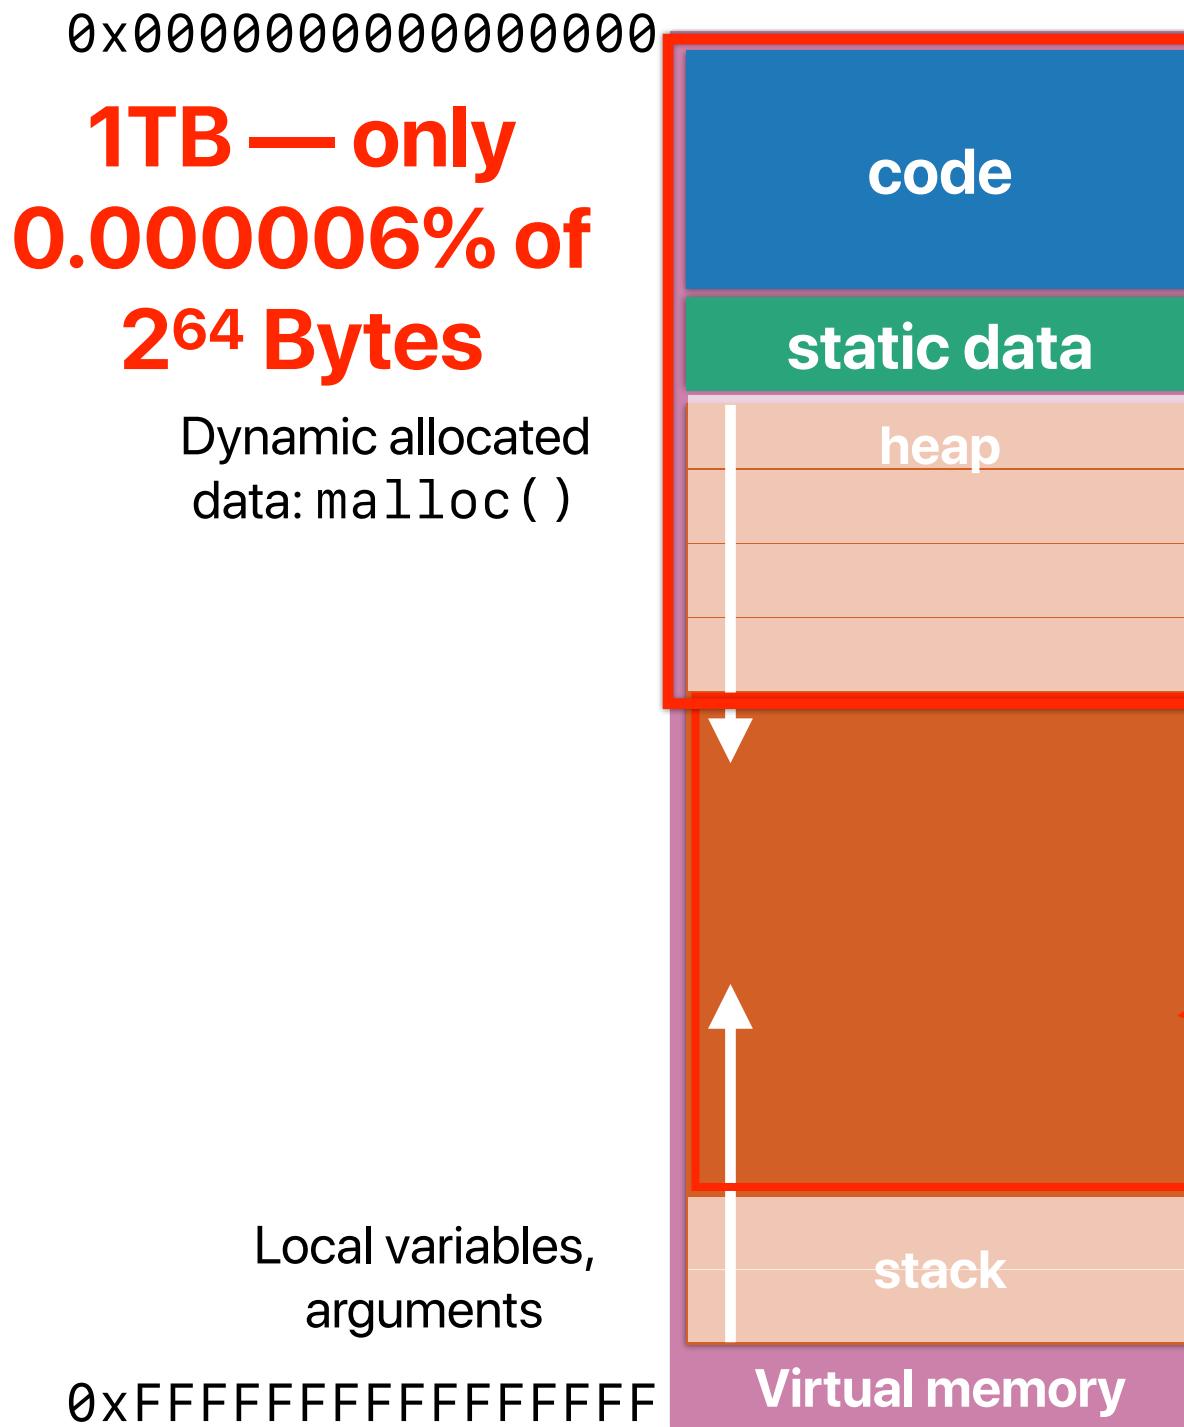
—  $\frac{2^{64} \text{ } B}{2^{12} \text{ } B}$  page table entries/leaf nodes —



.....



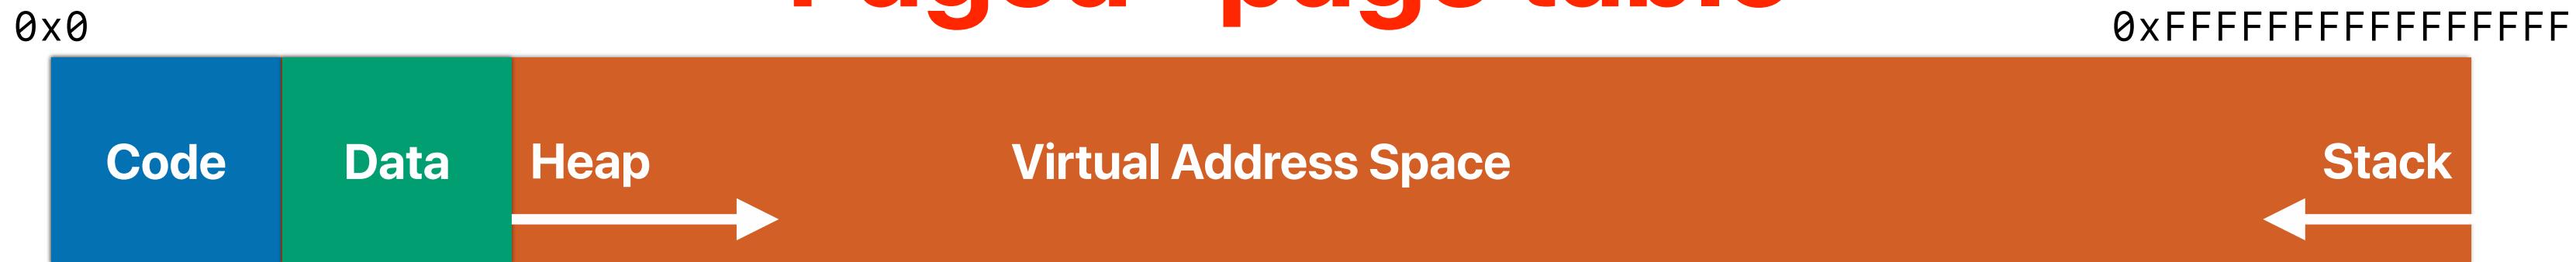
# Do we really need a large table?



Your program probably  
never uses this huge area!

If you still don't know why — you need to take CS202

# "Paged" page table

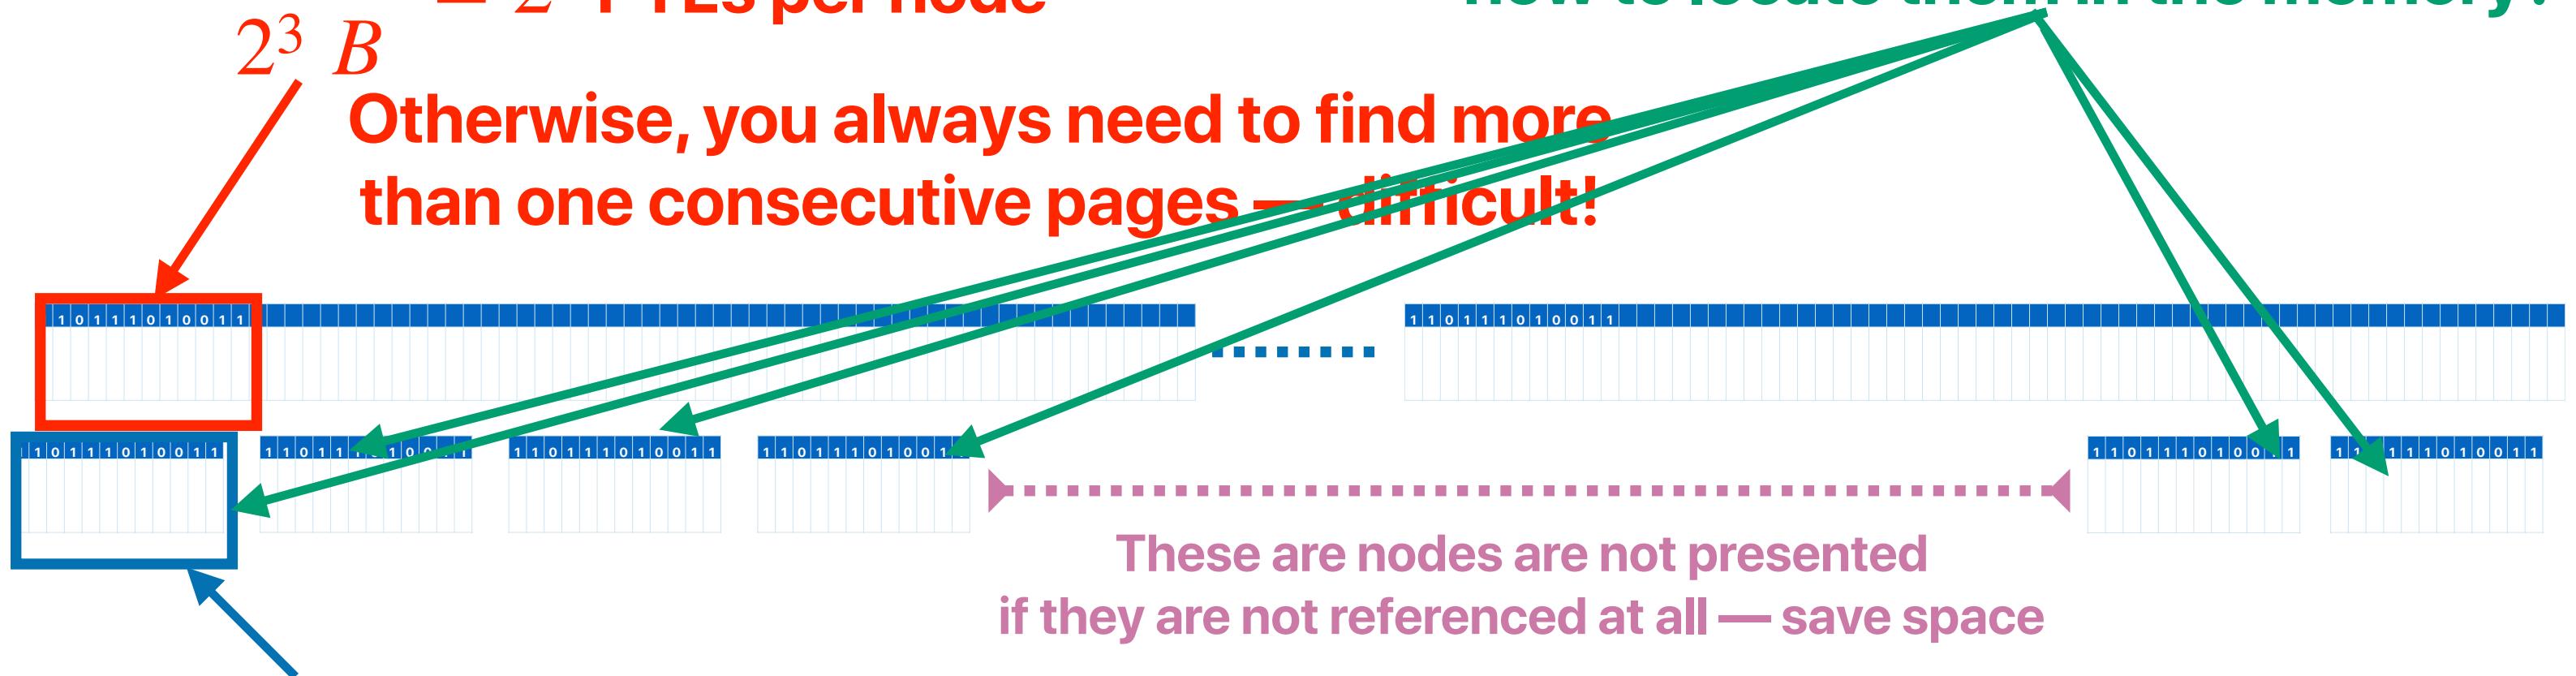


**Break up entries into pages!**

**Each of these occupies exactly a page**

$$-\frac{2^{12} \text{ B}}{2^3 \text{ B}} = 2^9 \text{ PTEs per node}$$

**Otherwise, you always need to find more than one consecutive pages — difficult!**



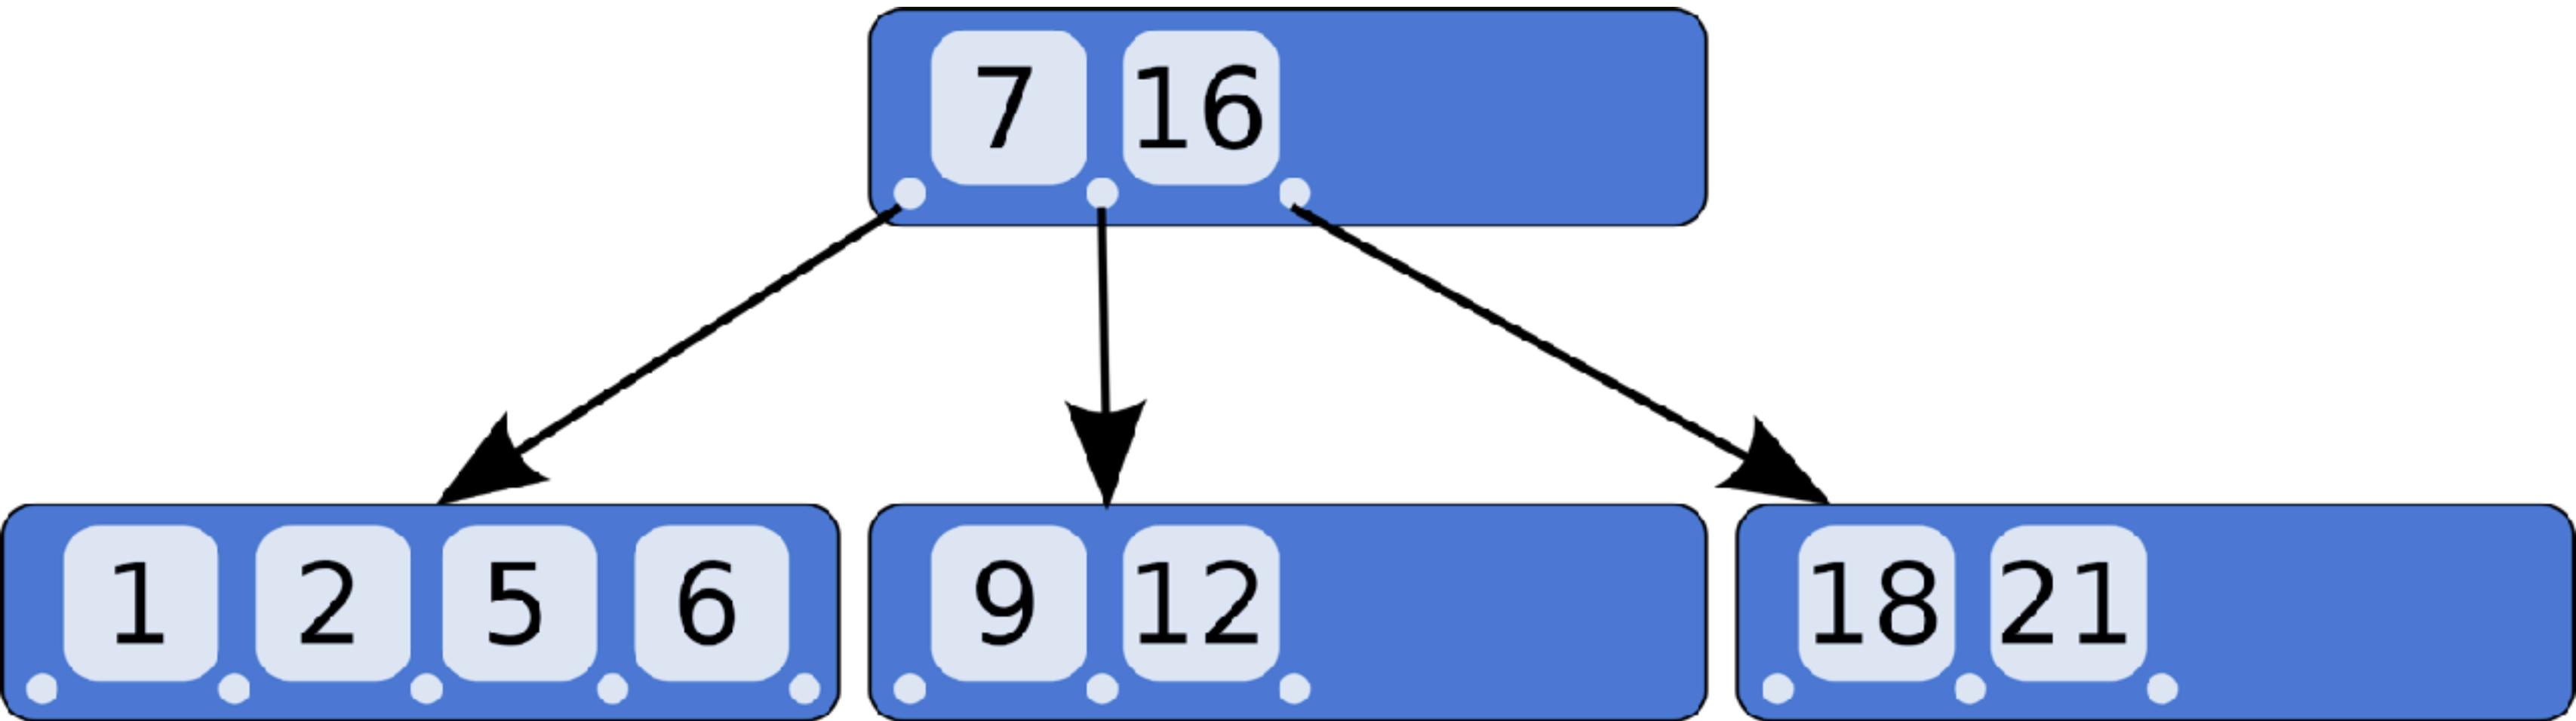
**Allocate page table entry nodes “on demand”**

**Question:**

**These nodes are spread out, how to locate them in the memory?**

These are nodes are not presented  
if they are not referenced at all — save space

# B-tree

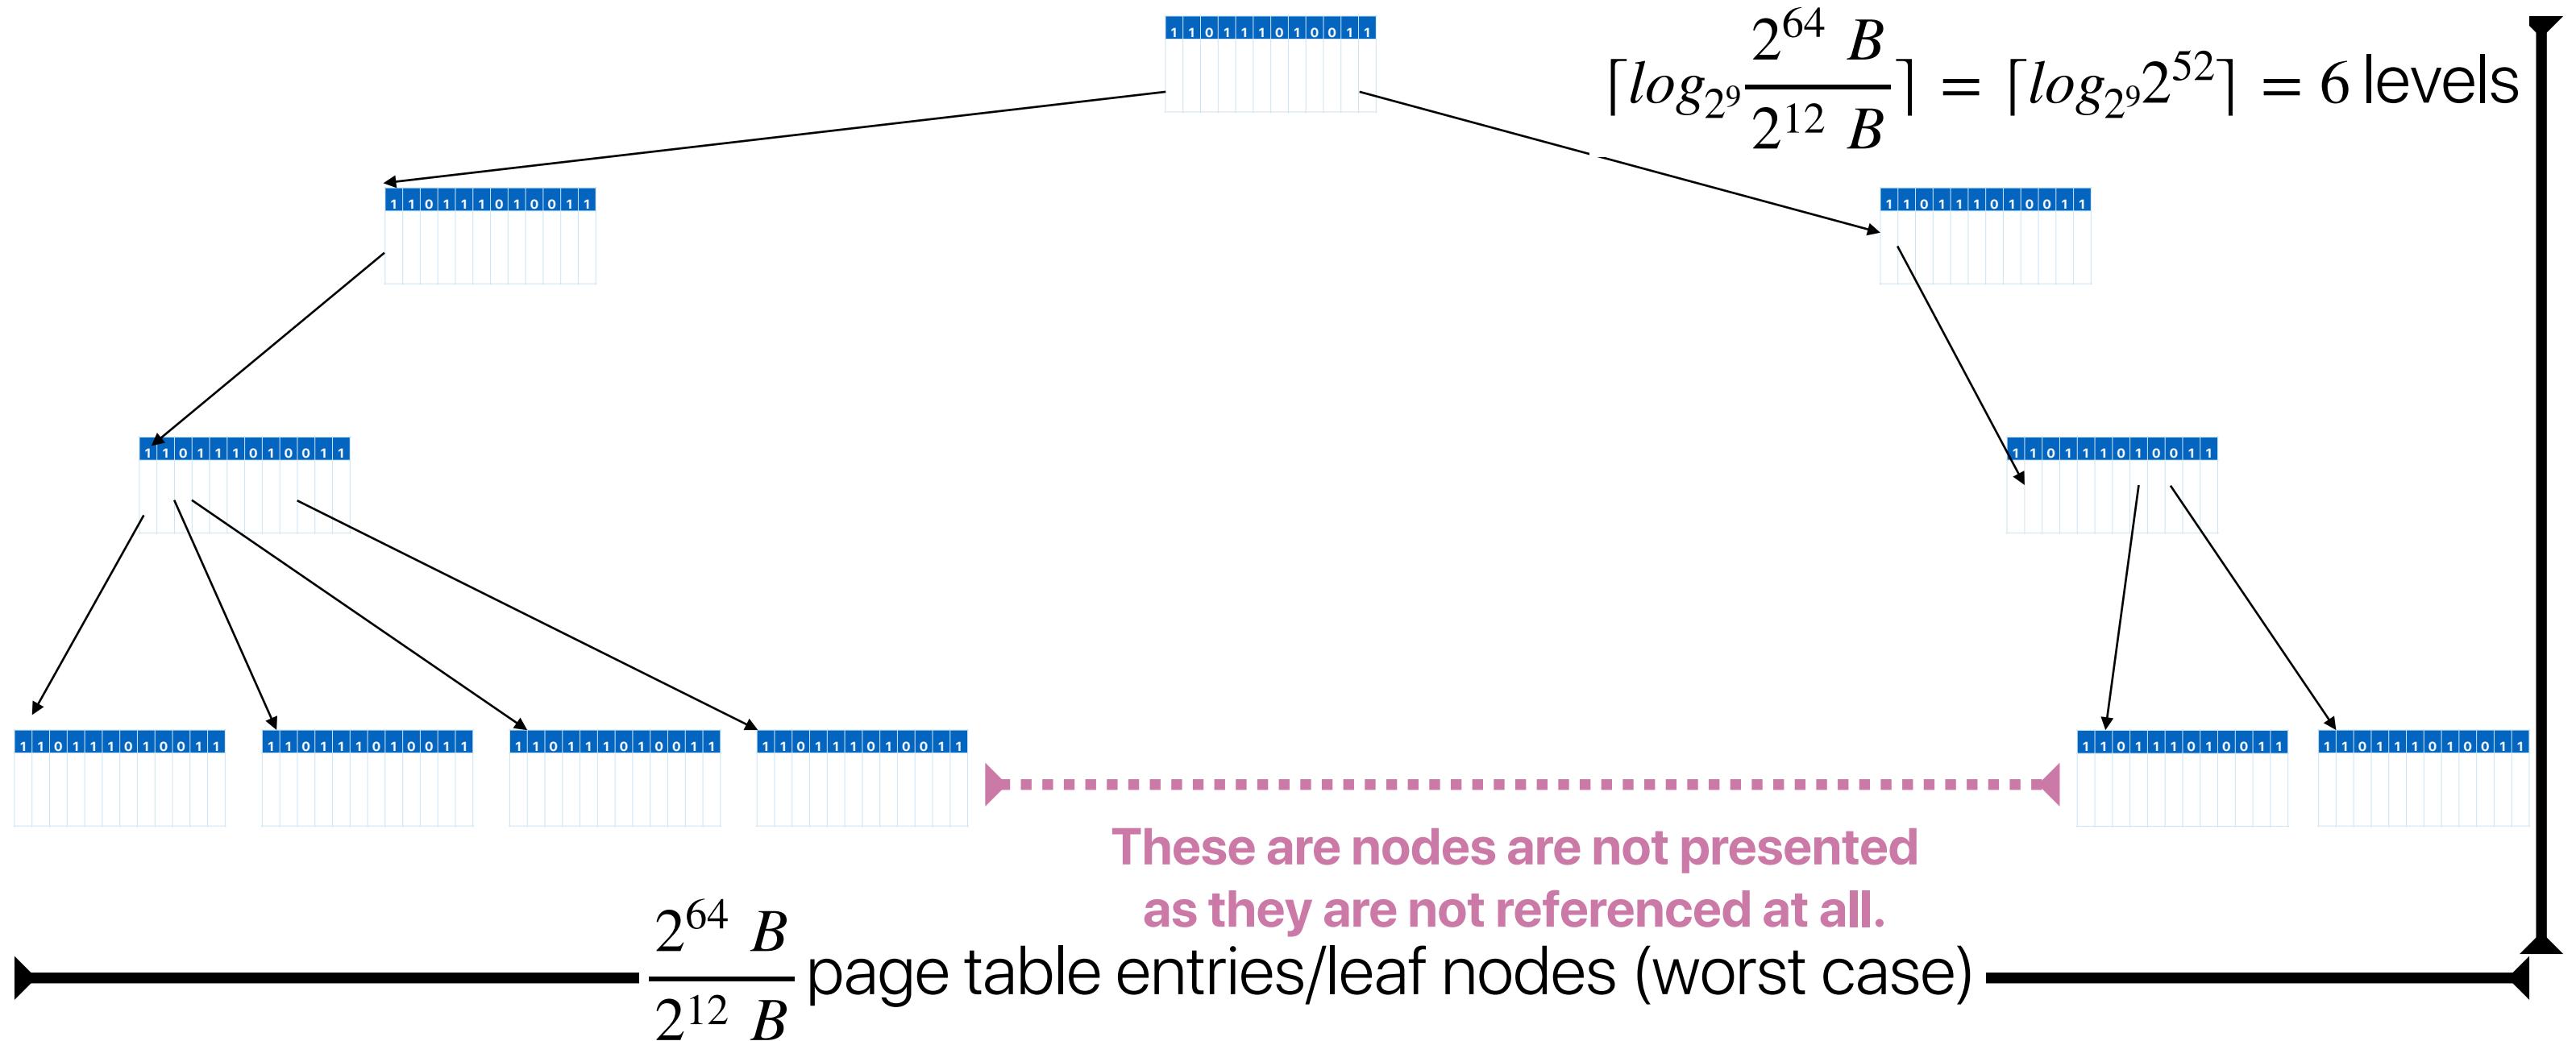


<https://en.wikipedia.org/wiki/B-tree#/media/File:B-tree.svg>

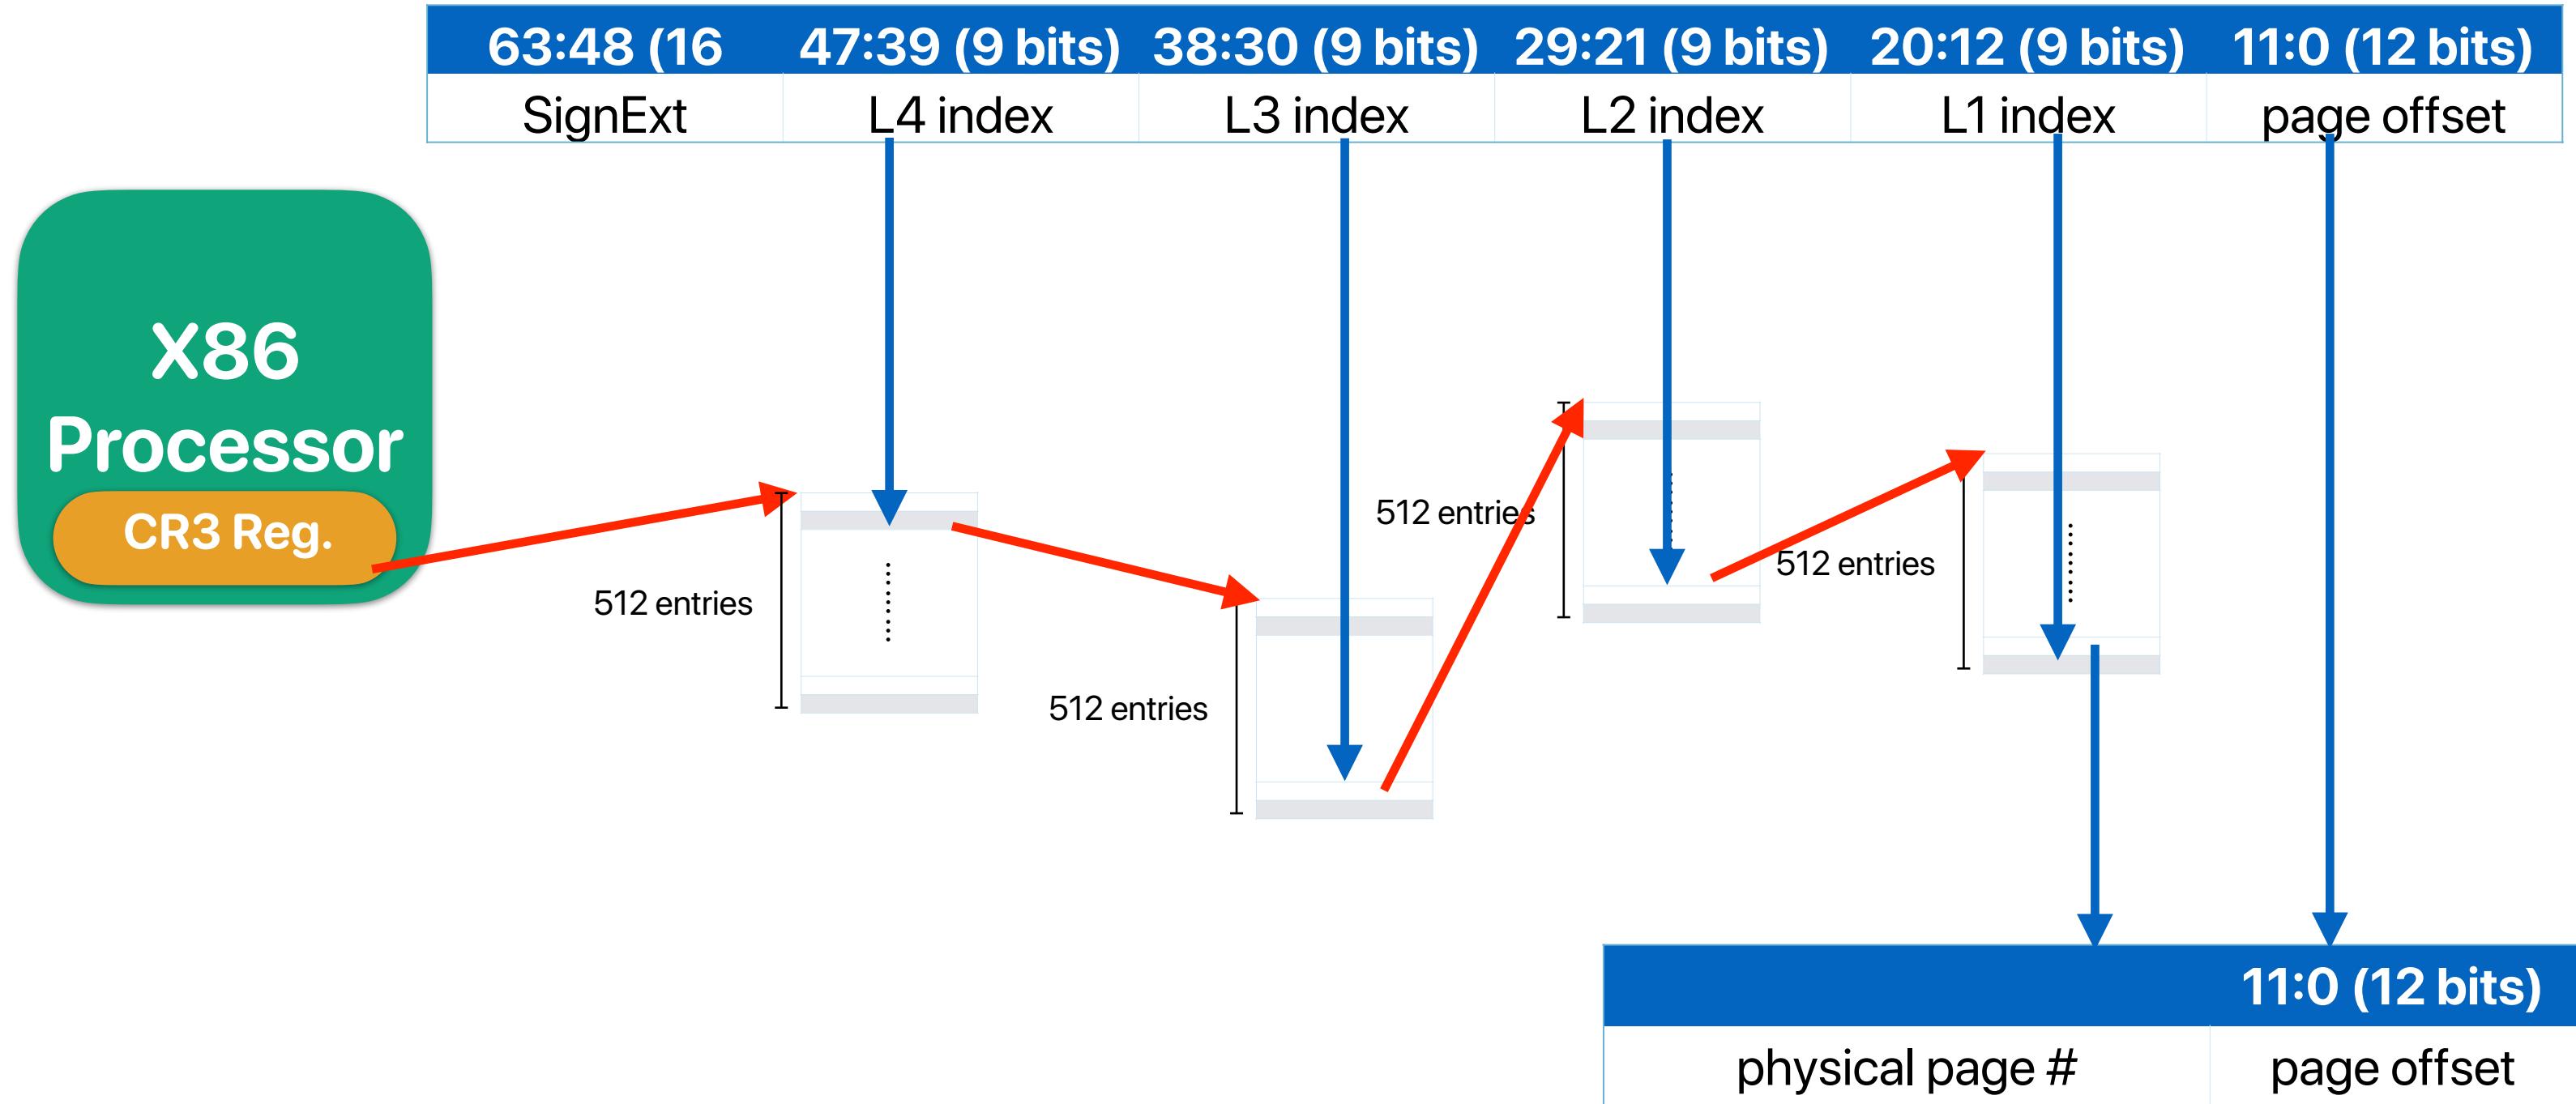
# Hierarchical Page Table

0x0

0xFFFFFFFFFFFFFFFFF



# Address translation in x86-64





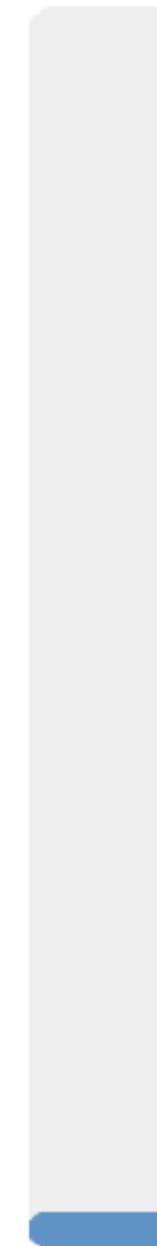
# When we have virtual memory...

- If an x86 processor supports virtual memory through the basic format of the page table as shown in the previous slide, how many memory accesses can a **mov** instruction that access data memory incur?
  - 2
  - 4
  - 6
  - 8
  - 10



 0

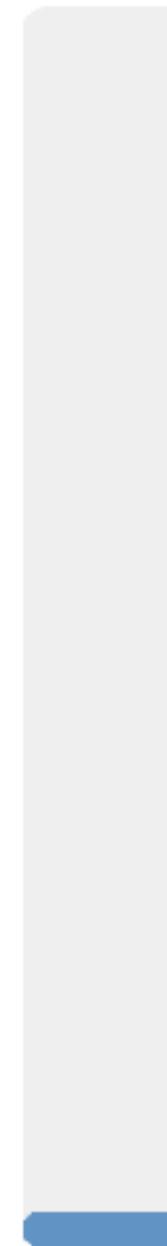
0



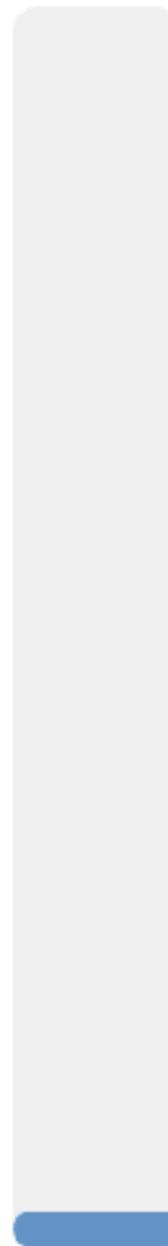
0



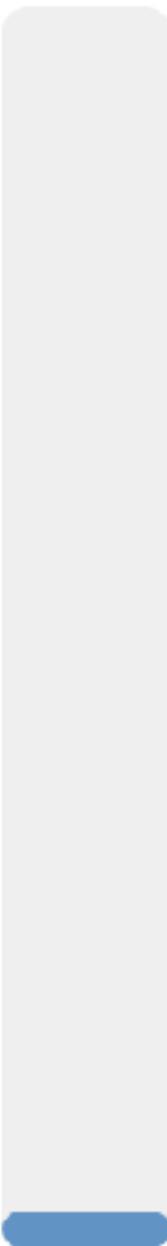
0



0



0



A

B

C

D

E



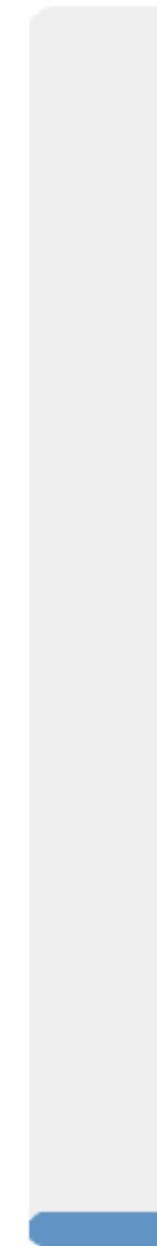
# When we have virtual memory...

- If an x86 processor supports virtual memory through the basic format of the page table as shown in the previous slide, how many memory accesses can a **mov** instruction that access data memory incur?
  - 2
  - 4
  - 6
  - 8
  - 10



 0

0



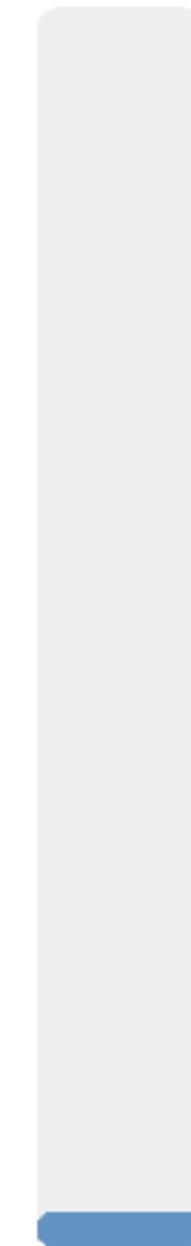
A

0



B

0



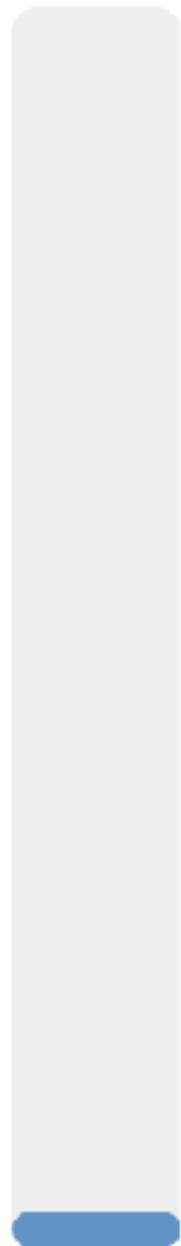
C

0



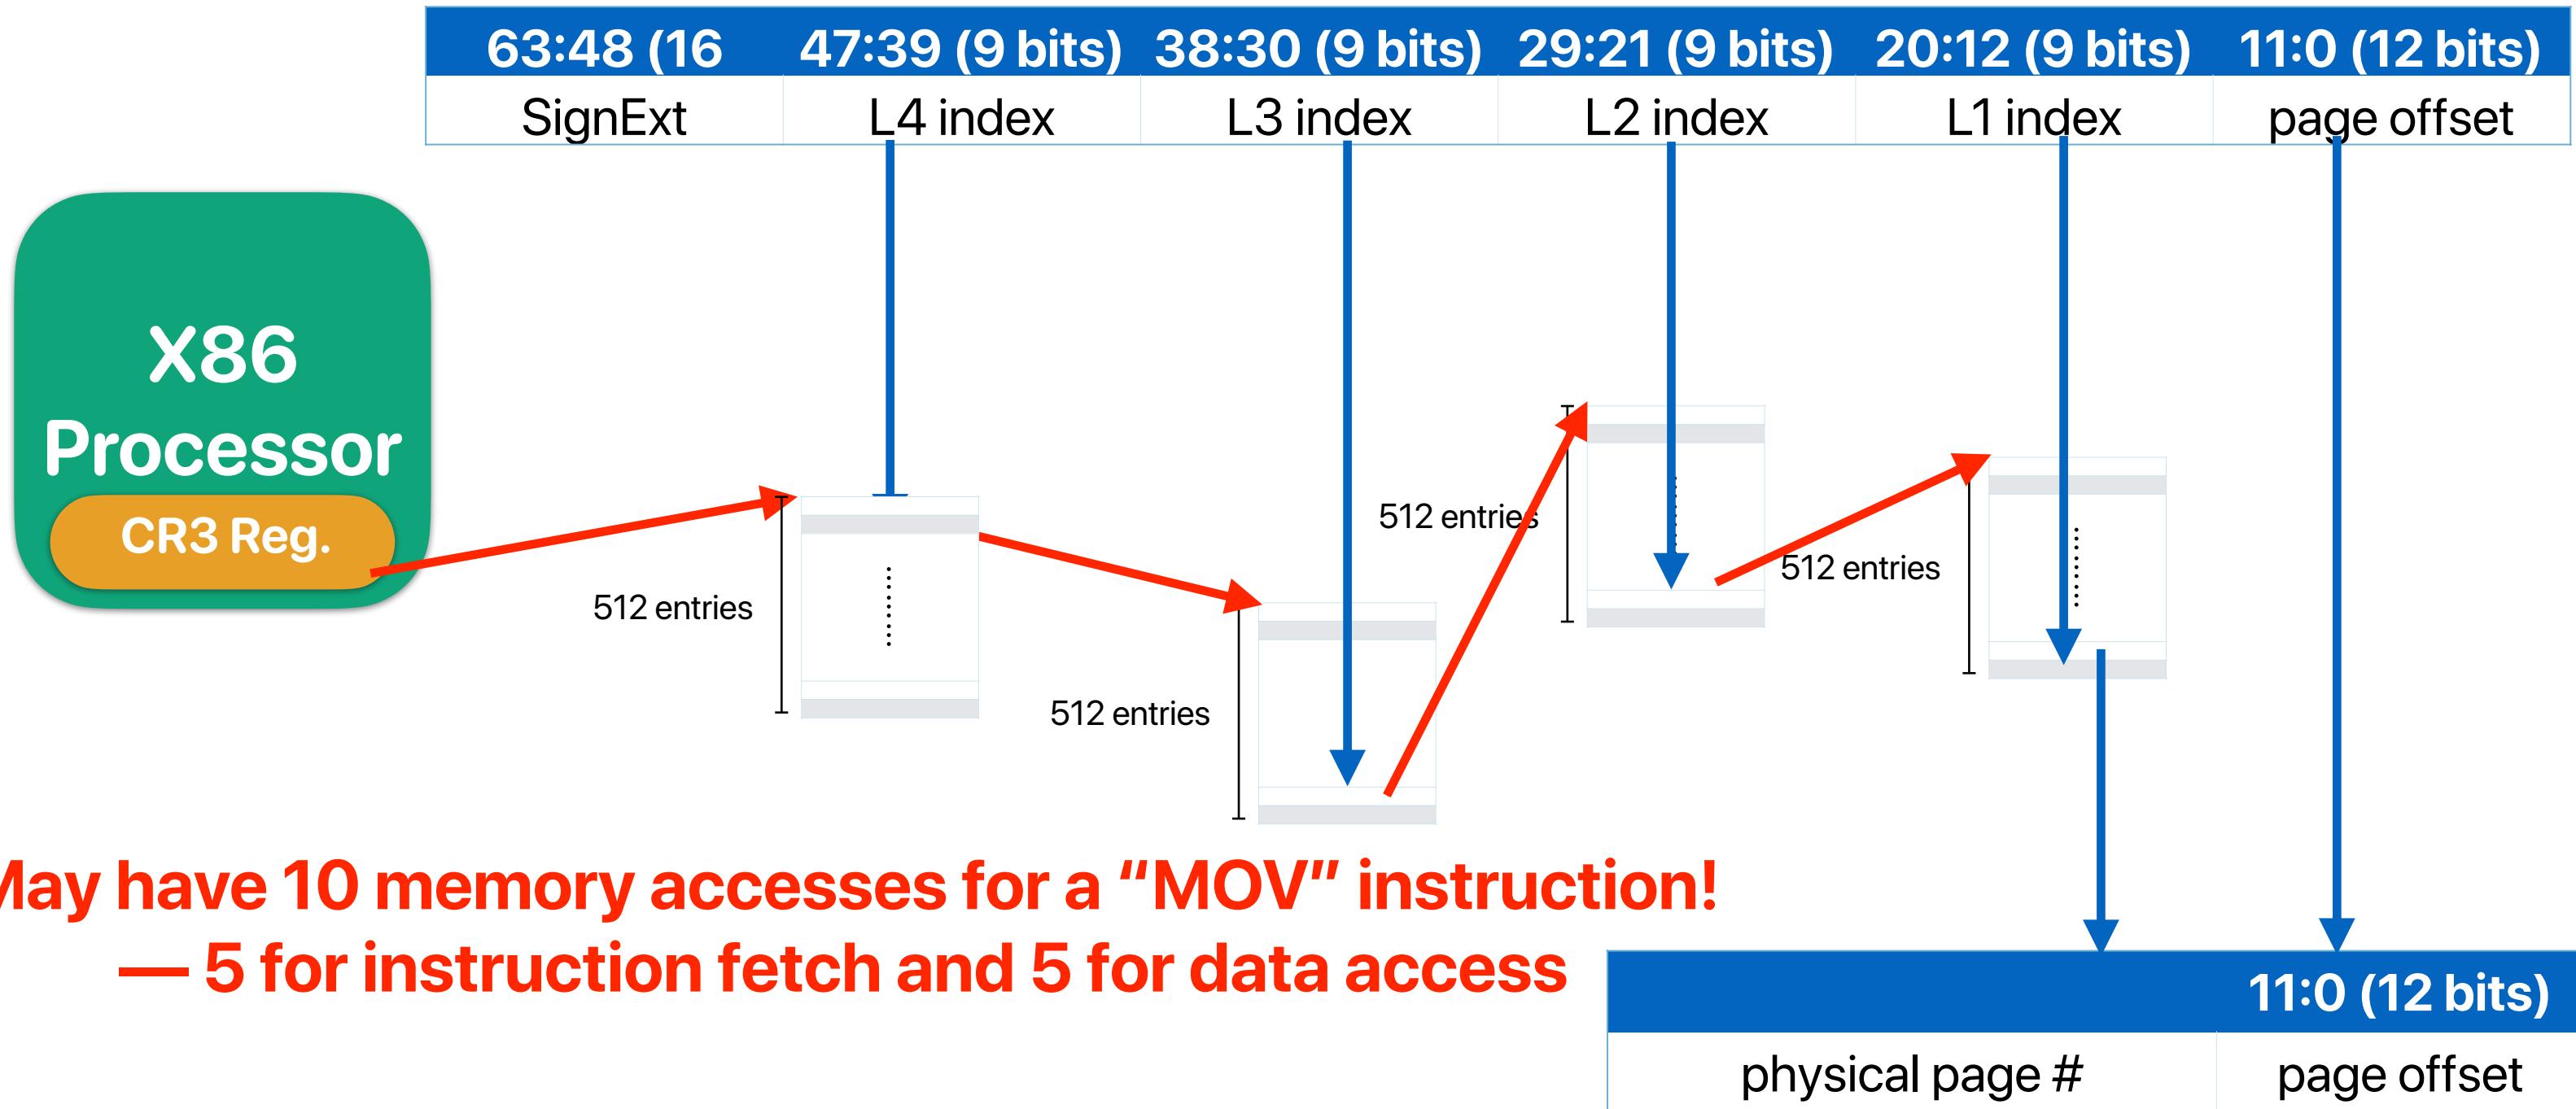
D

0



E

# Address translation in x86-64

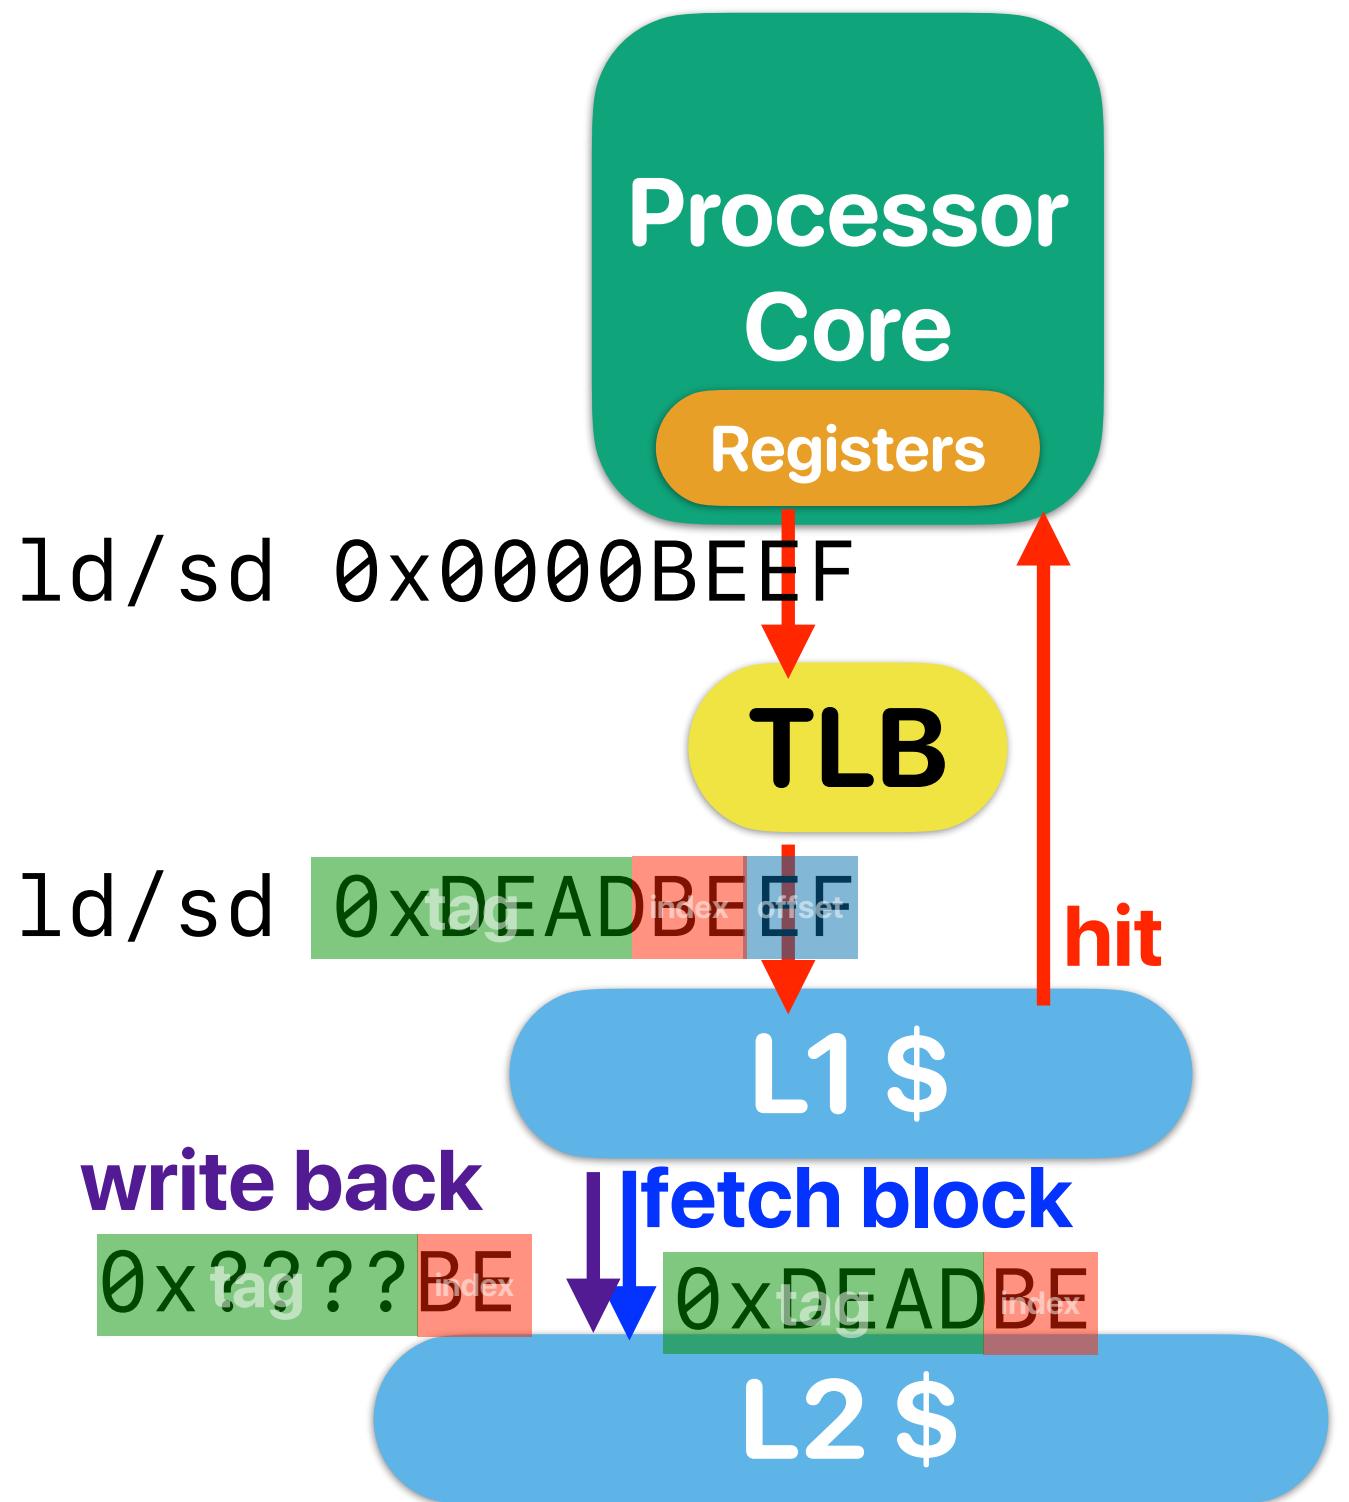


# When we have virtual memory...

- If an x86 processor supports virtual memory through the basic format of the page table as shown in the previous slide, how many memory accesses can a **mov** instruction that access data memory incur?
  - 2
  - 4
  - 6
  - 8
  - 10

# **Avoiding the address translation overhead**

# TLB: Translation Look-aside Buffer



- TLB — a small SRAM stores frequently used page table entries
- Good — A lot faster than having everything going to the DRAM
- Bad — Still on the critical path

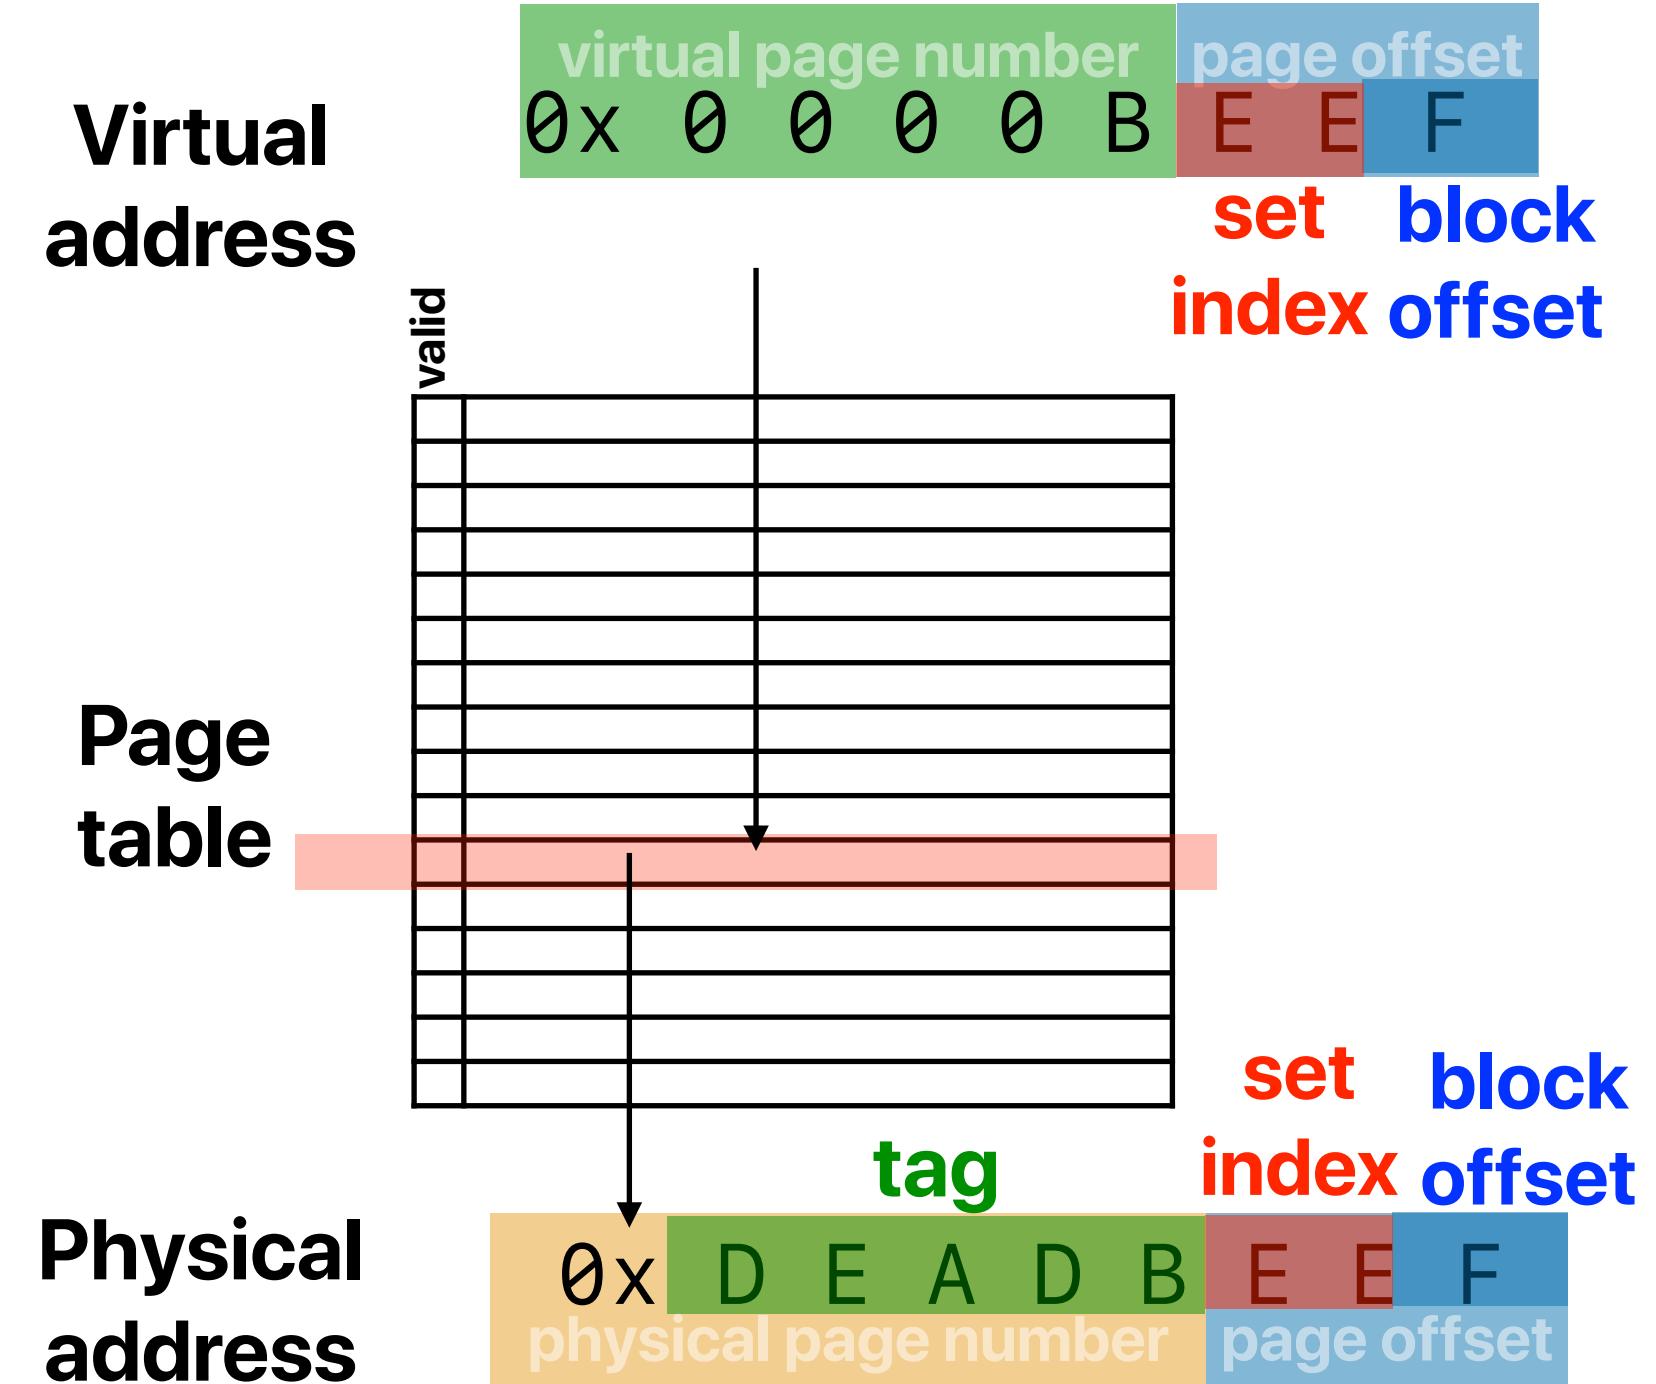
# TLB + Virtual cache

- L1 \$ accepts virtual address — you don't need to translate
- Good — you can access both TLB and L1-\$ at the same time and physical address is only needed if L1-\$ misses
- Bad — it doesn't work in practice
  - Many applications have the same virtual address but should be pointing different **physical addresses**
  - An application can have “aliasing virtual addresses” pointing to the same **physical address**

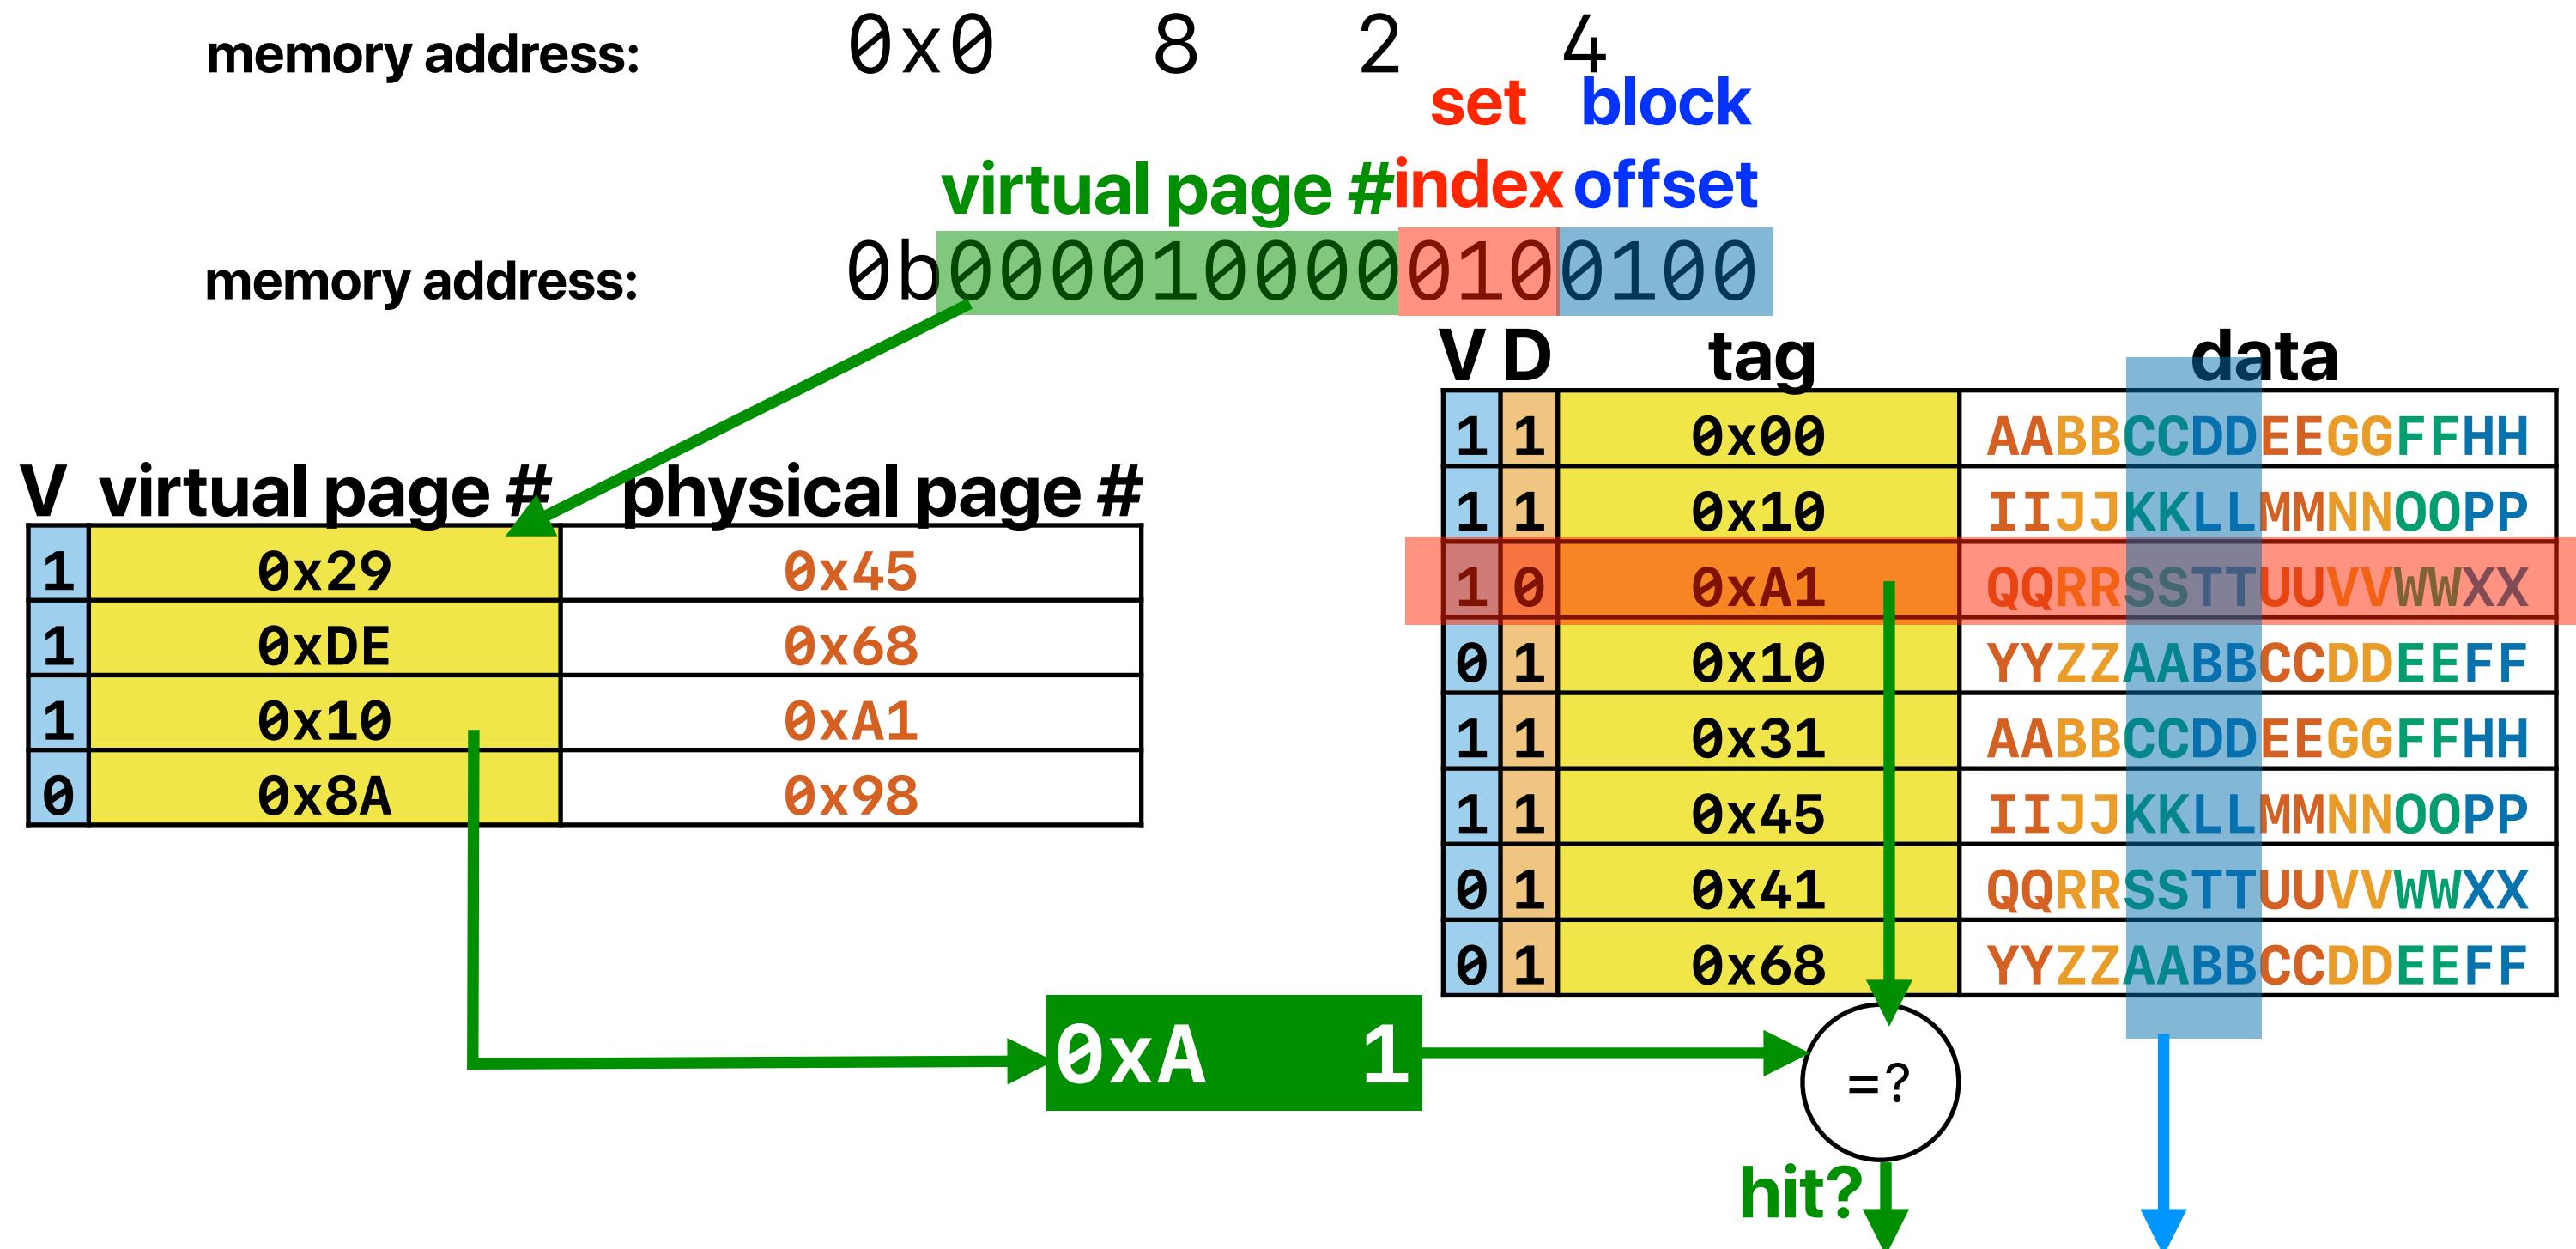


# Virtually indexed, physically tagged cache

- Can we find physical address directly in the virtual address
  - Not everything — but the page offset isn't changing!
- Can we indexing the cache using the "partial physical address"?
  - Yes — Just make set index + block set to be exactly the page offset



# Virtually indexed, physically tagged cache



# Virtually indexed, physically tagged cache

- If page size is 4KB —

$$\lg(B) + \lg(S) = \lg(4096) = 12$$

$$C = ABS$$

$$C = A \times 2^{12}$$

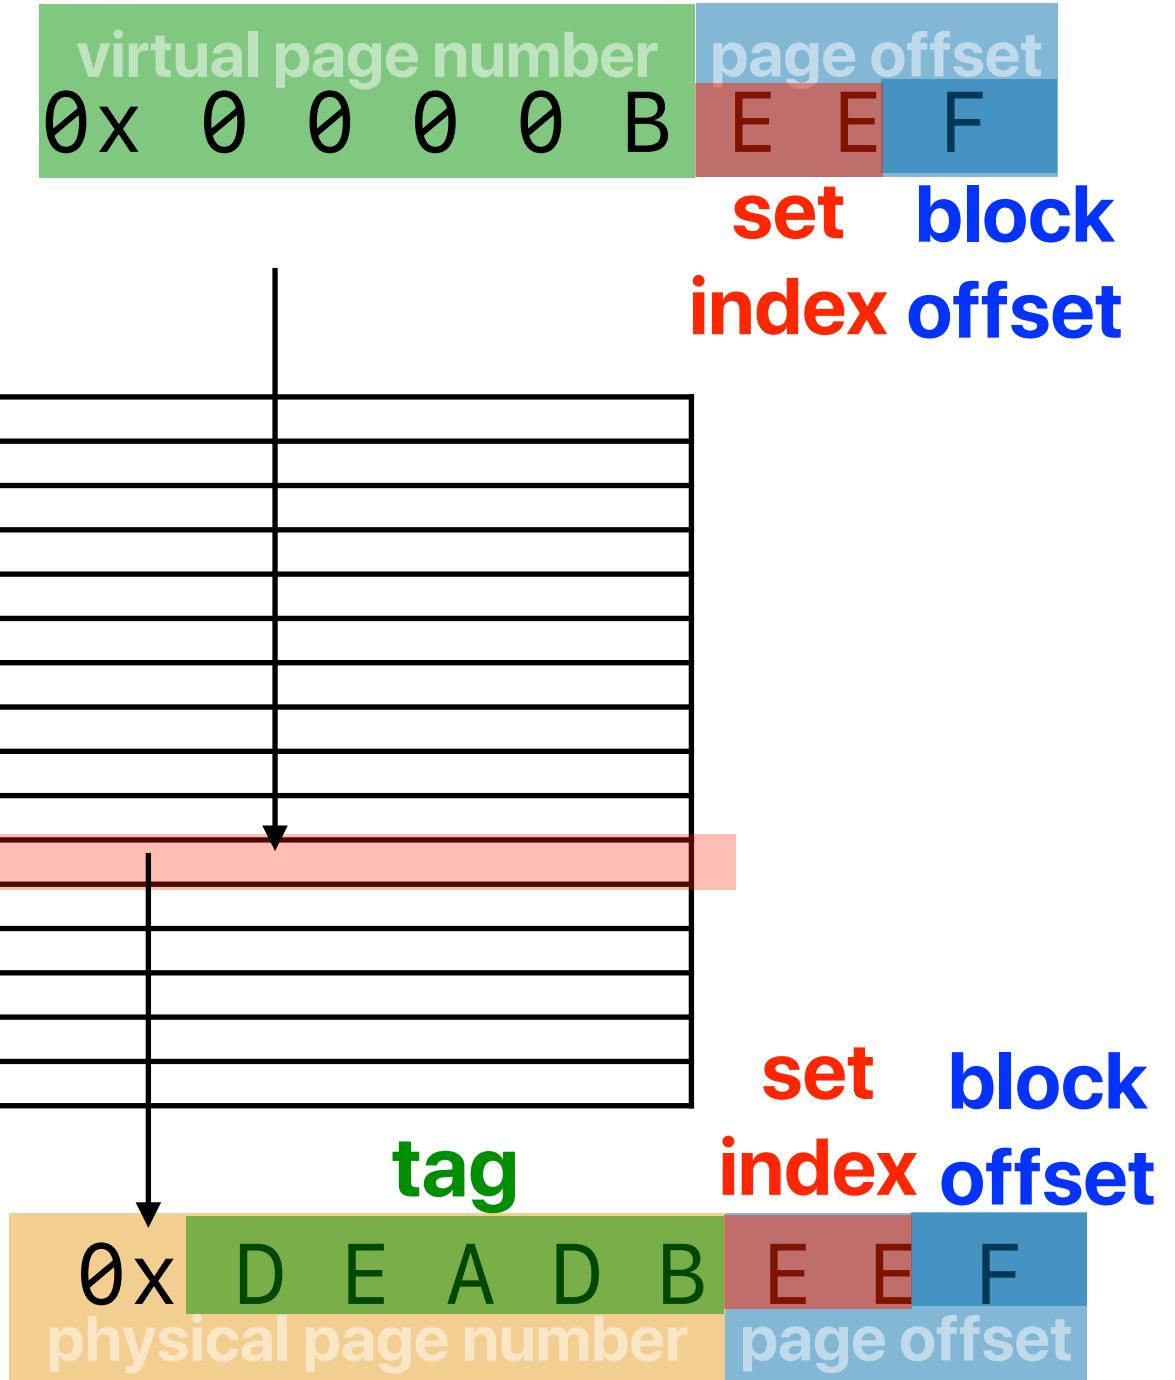
$$\text{if } A = 1$$

$$C = 4KB$$

Virtual address

Page table

Physical address



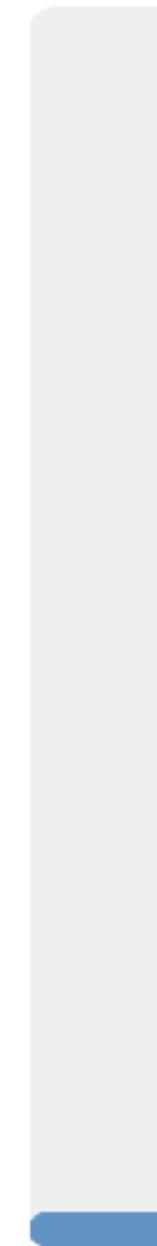


## Virtual indexed, physical tagged cache limits the cache size

- If you want to build a virtual indexed, physical tagged cache with 32KB capacity, which of the following configuration is possible? Assume the operating system use 4K pages.
  - A. 32B blocks, 2-way
  - B. 32B blocks, 4-way
  - C. 64B blocks, 4-way
  - D. 64B blocks, 8-way

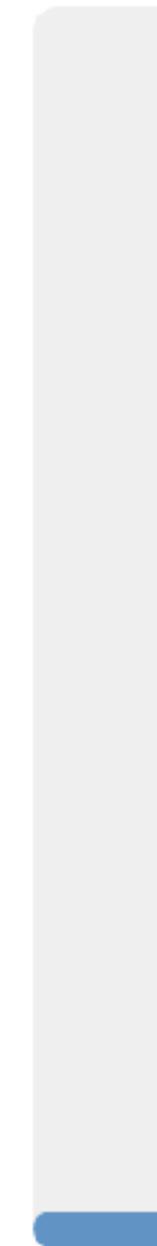
 0

0



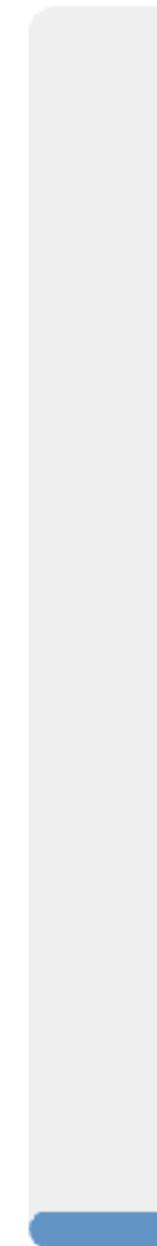
A

0



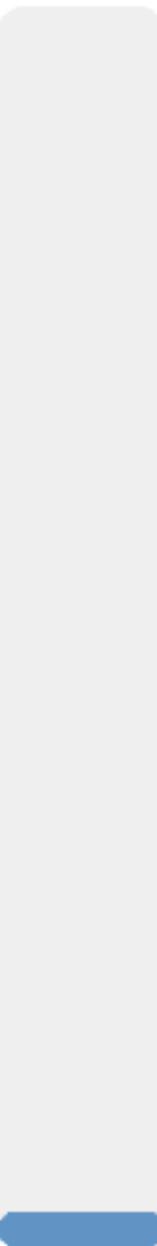
B

0



C

0



D

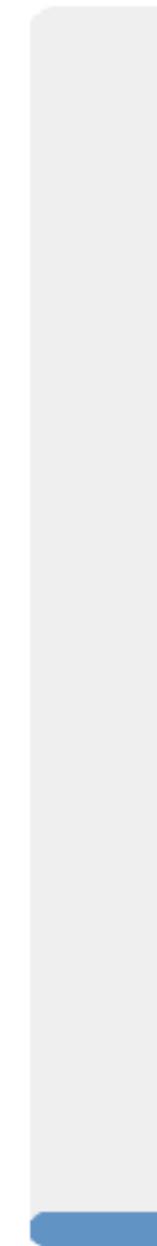


## Virtual indexed, physical tagged cache limits the cache size

- If you want to build a virtual indexed, physical tagged cache with 32KB capacity, which of the following configuration is possible? Assume the operating system use 4K pages.
  - A. 32B blocks, 2-way
  - B. 32B blocks, 4-way
  - C. 64B blocks, 4-way
  - D. 64B blocks, 8-way

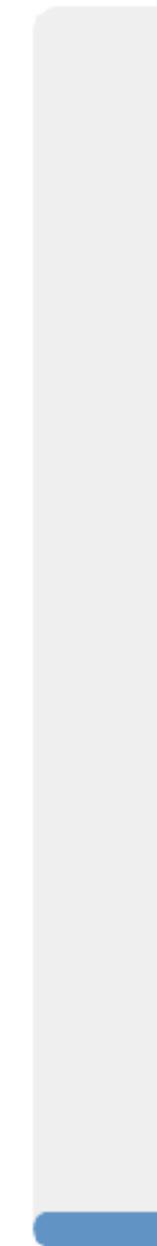
 0

0



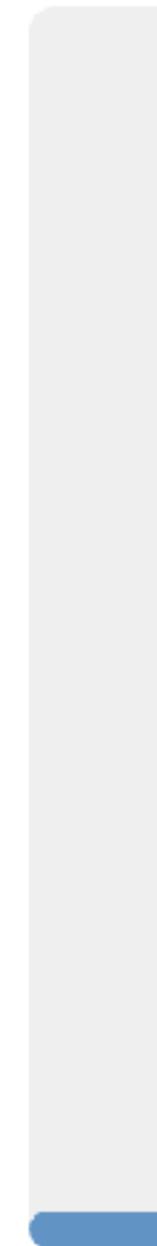
A

0



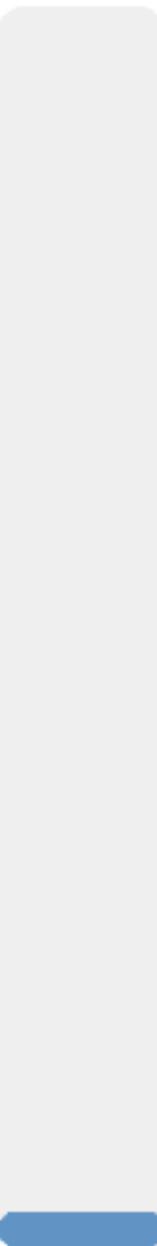
B

0



C

0



D

## Virtual indexed, physical tagged cache limits the cache size

- If you want to build a virtual indexed, physical tagged cache with 32KB capacity, which of the following configuration is possible? Assume the operating system use 4K pages.

- A. 32B blocks, 2-way
- B. 32B blocks, 4-way
- C. 64B blocks, 4-way
- D. 64B blocks, 8-way

$$\lg(B) + \lg(S) = \lg(4096) = 12$$

$$C = ABS$$

$$32KB = A \times 2^{12}$$

$$A = 8$$

Exactly how Core i7 9<sup>th</sup>  
generation configures its own  
cache



Start the presentation to see live content. Still no live content? Install the app or get help at [PollEv.com/app](https://PollEv.com/app)

Q & A



# Announcement

- Reading quiz due this Wednesday before the lecture
- Midterm tomorrow — 2p-3:20p, no exceptions
  - Be sure to watch the midterm review lecture
  - Sample midterm slides are already released with the last lecture — no solution should be provided as questions are mostly from previous lectures
  - In-person examine
    - Closed-book, closed notebook. No cheat sheet allowed
  - Online examine will have 1.5x questions compared to in-person ones
    - Since typing is faster — based on my teaching experience and scientific research
    - Learning LaTeX equation editing is part of the course — frequently used in scientific editing and becoming the standard in MS Office and Mac
    - It's an open-book/note test and self-proctored — you can preset some frequently used formula and we don't have many in this class
    - Discussion/communication with any other during the examine is prohibited

Typing is far faster than handwriting: the average American can type 40 words per minute, but can only handwrite 13 words per minute. And it's easier to organize, edit, and synthesize notes when they exist on a hard drive rather than on paper. Sep 7, 2021



The Daily Princetonian

<https://www.dailyprincetonian.com/article/2021/09/>

On handwriting - The Daily Princetonian

# Computer Science & Engineering

142

つづく

